



Microservices Deployment in Cloud Service Mesh

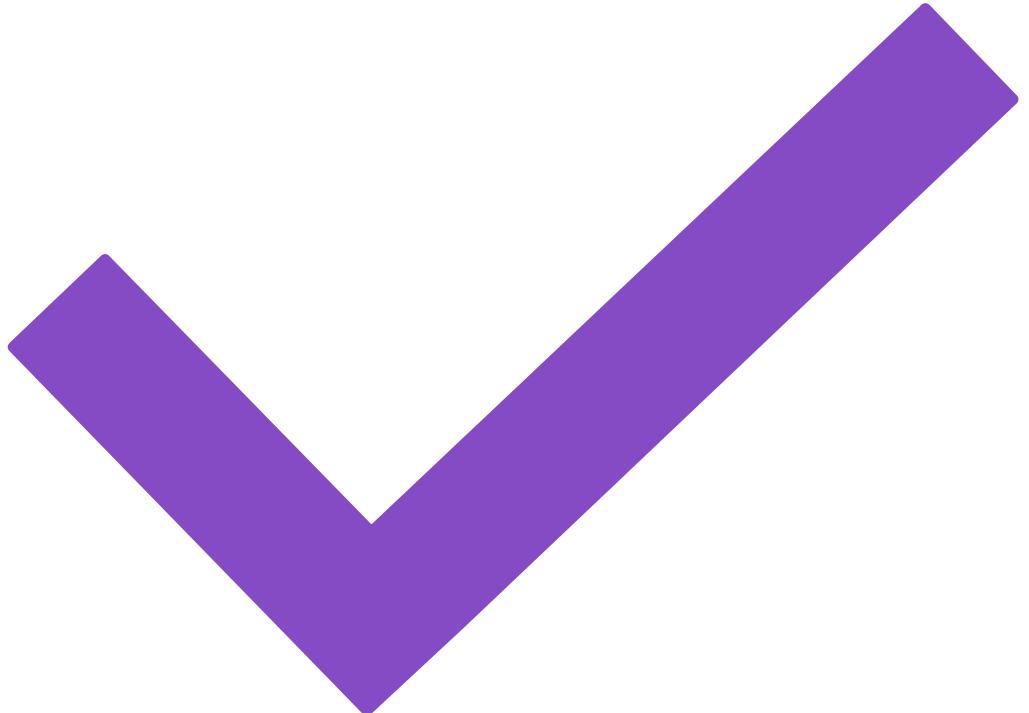
ECE 4150

Spring 2024

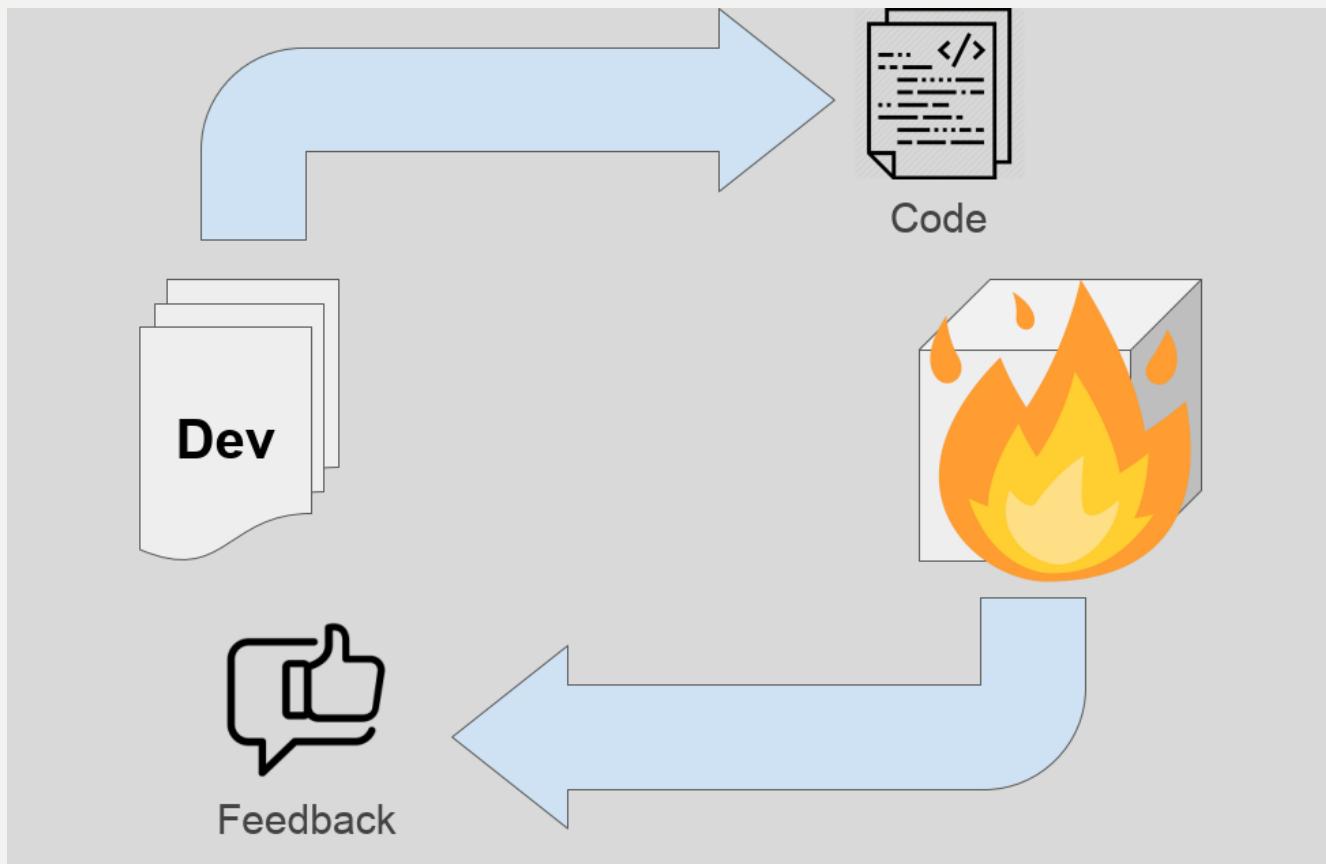
Madisetti

Outline

- **Problems with Microservices**
- **Kubernetes Model**
- **Istio Service Mesh**
- **Features that Istio Provides**
- **Examples**

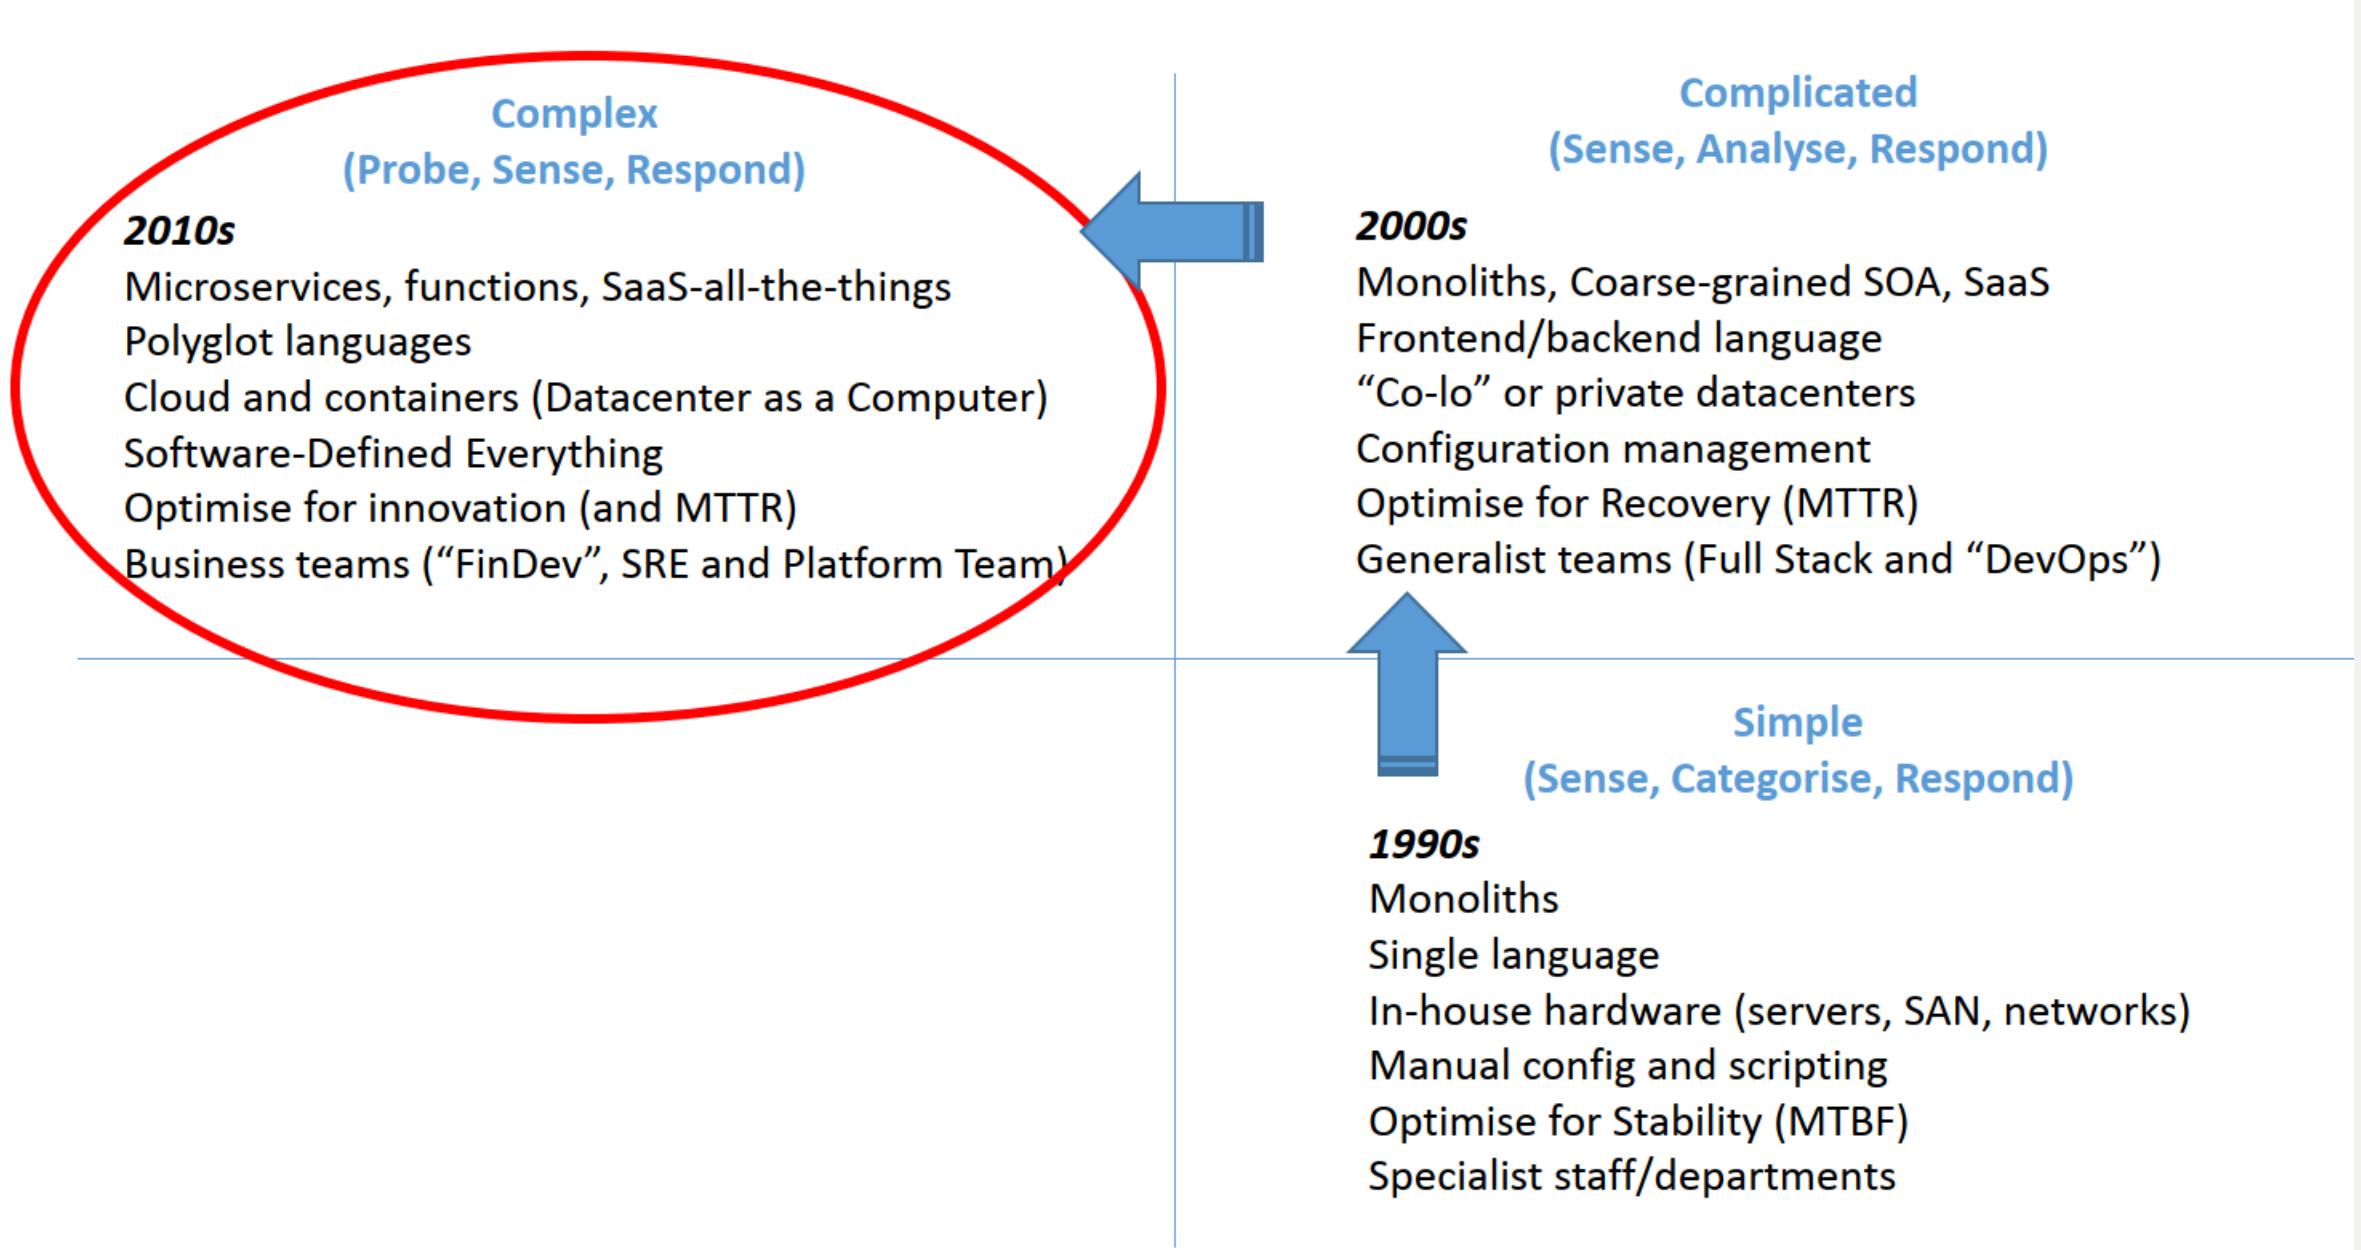


Ideal world vs practice

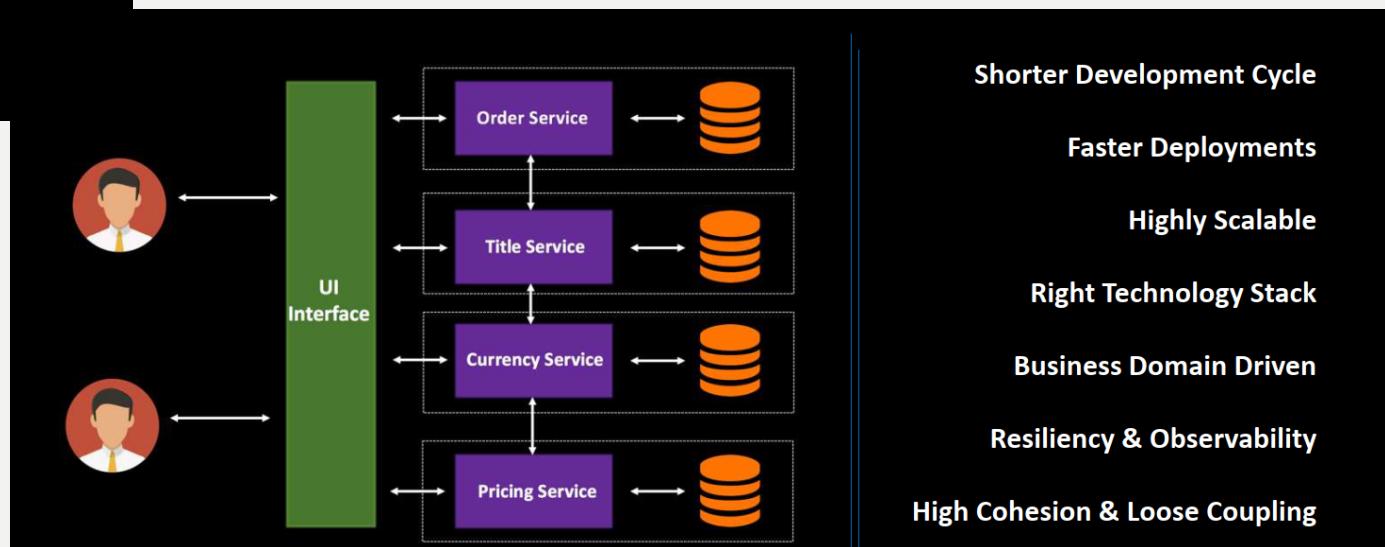
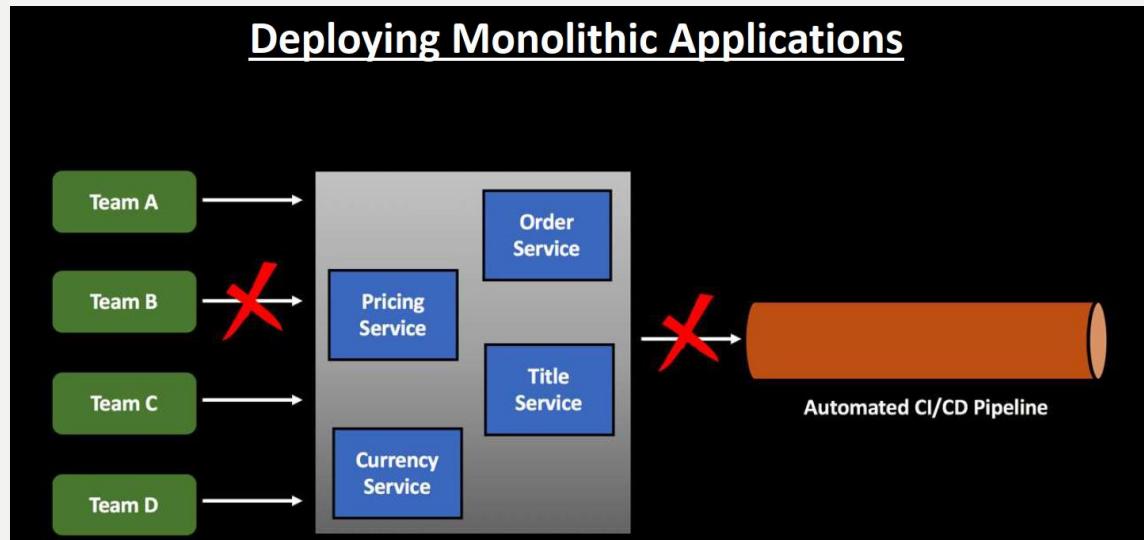


Fallacies of Distributed Computing

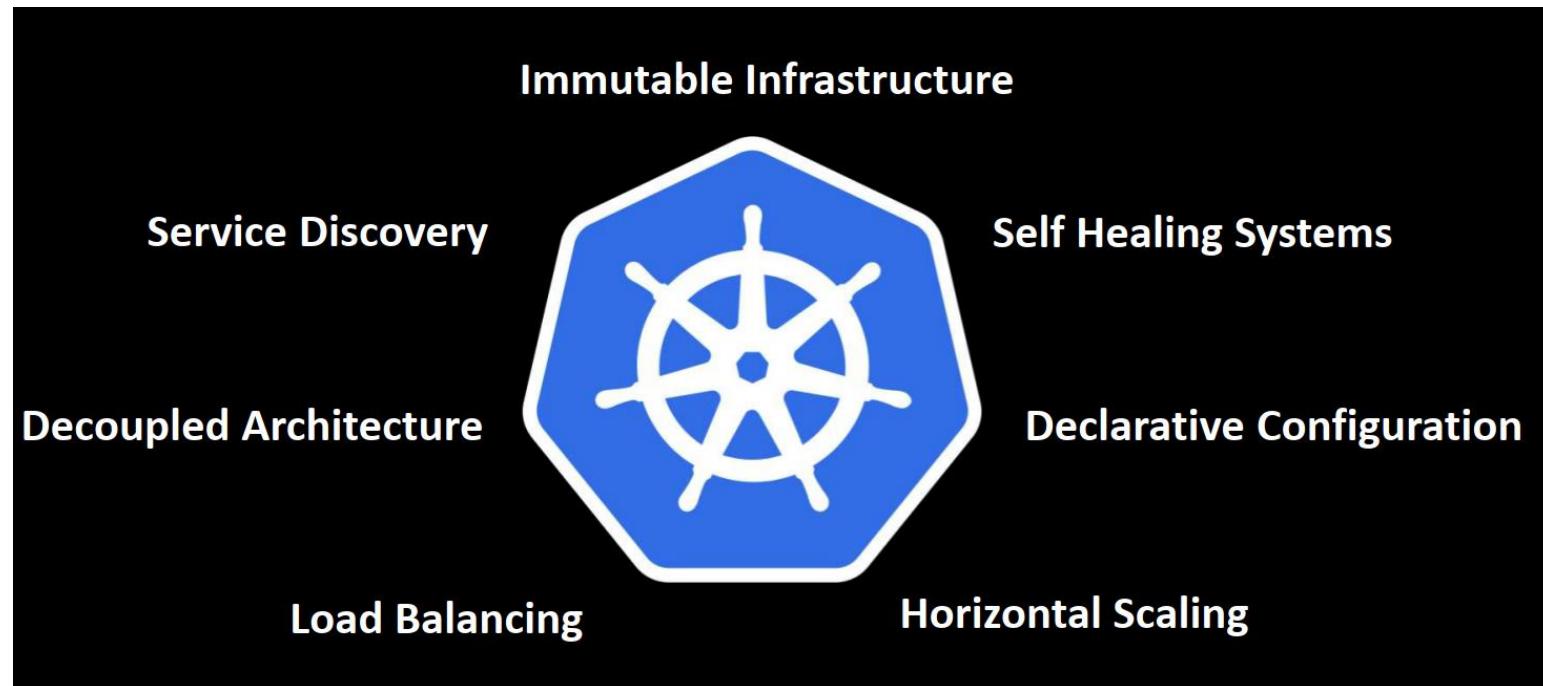
1. The network is reliable.
2. Latency is zero.
3. Bandwidth is infinite.
4. The network is secure.
5. Topology doesn't change.
6. There is one administrator.
7. Transport cost is zero.
8. The network is homogeneous.



Microservices



Kubernetes (K8S) provides the Infrastructure to Microservices



What is a *Tech Stack*?

A **tech stack** is a combination of tools, software, and frameworks with which you build your microservices.

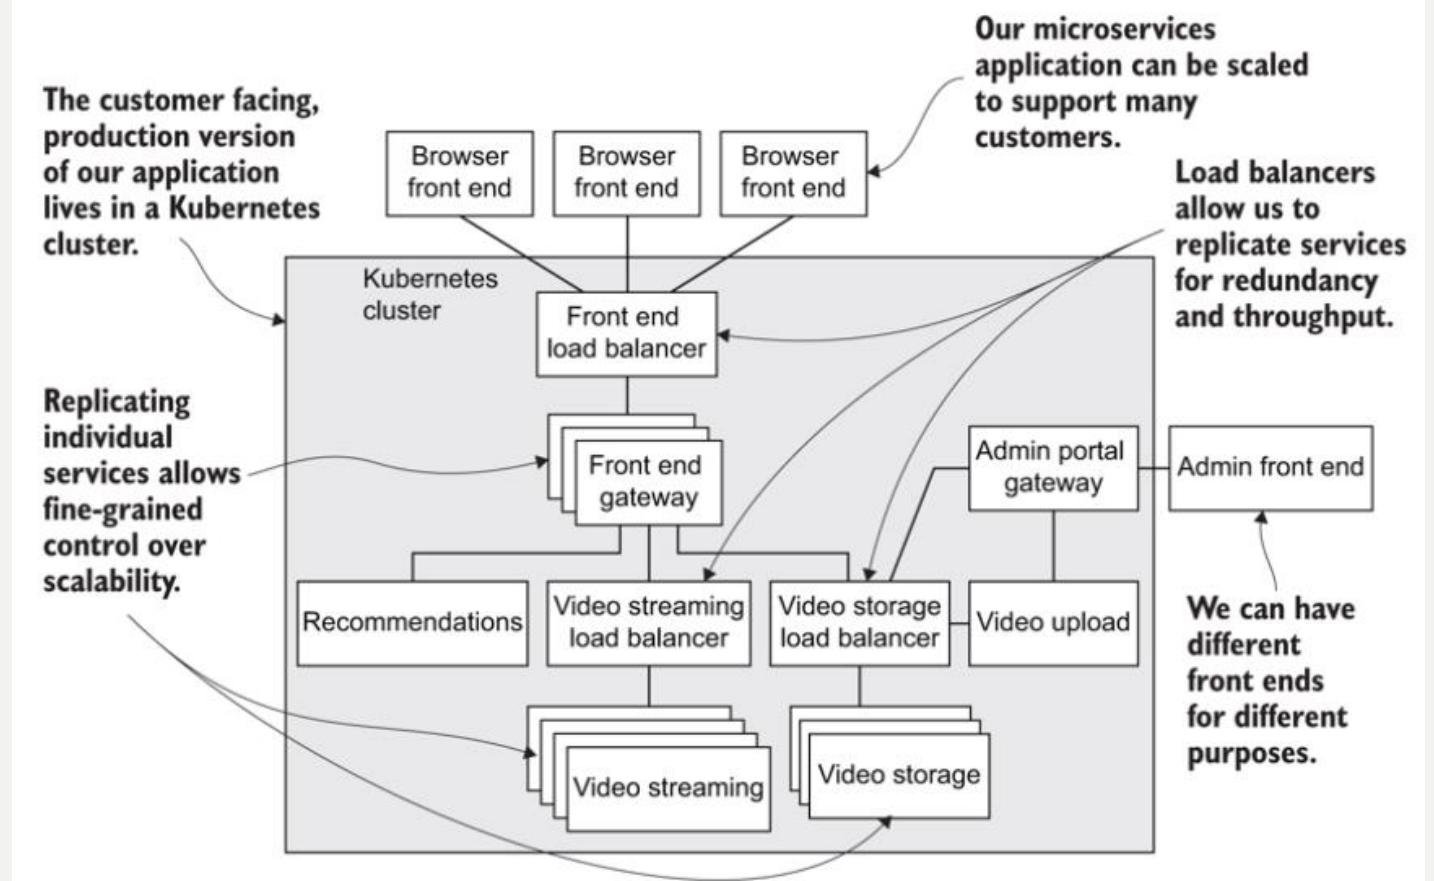
Examples of tech stacks: **MEAN** (MongoDB, Express, Angular, Node.js), **LAMP** (Linux, Apache, MySQL, PHP)... You can choose a different tech stack for each microservice in a microservices application.

Tools that are useful in Microservices

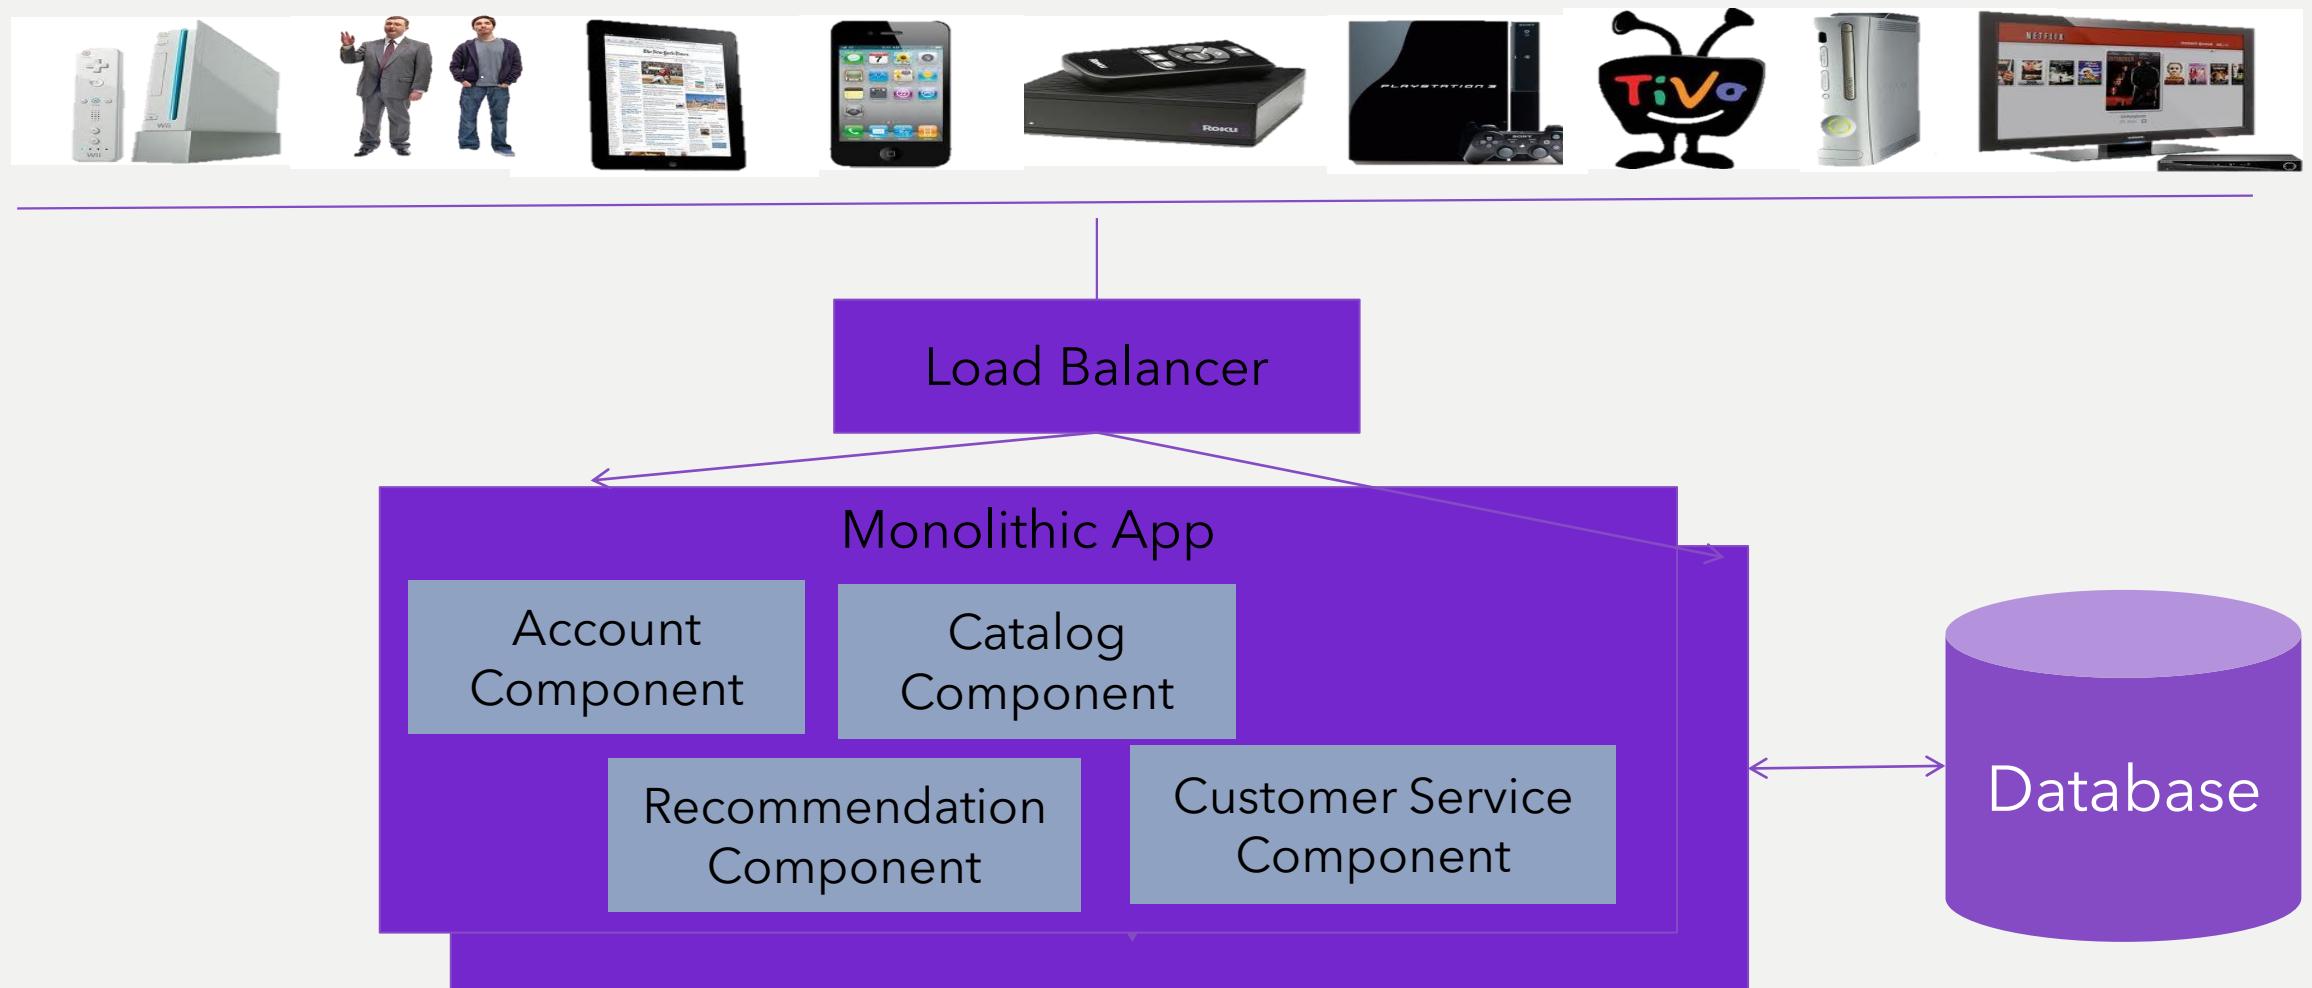
-  Docker - To package and deploy microservices
-  Docker Compose - To test microservices
-  Kubernetes (K8S) - To host application in the cloud
-  Terraform - To build our software factory production infrastructure in the cloud using code (**IaC** approach - a declarative approach).

Packaging for Production & Deployment

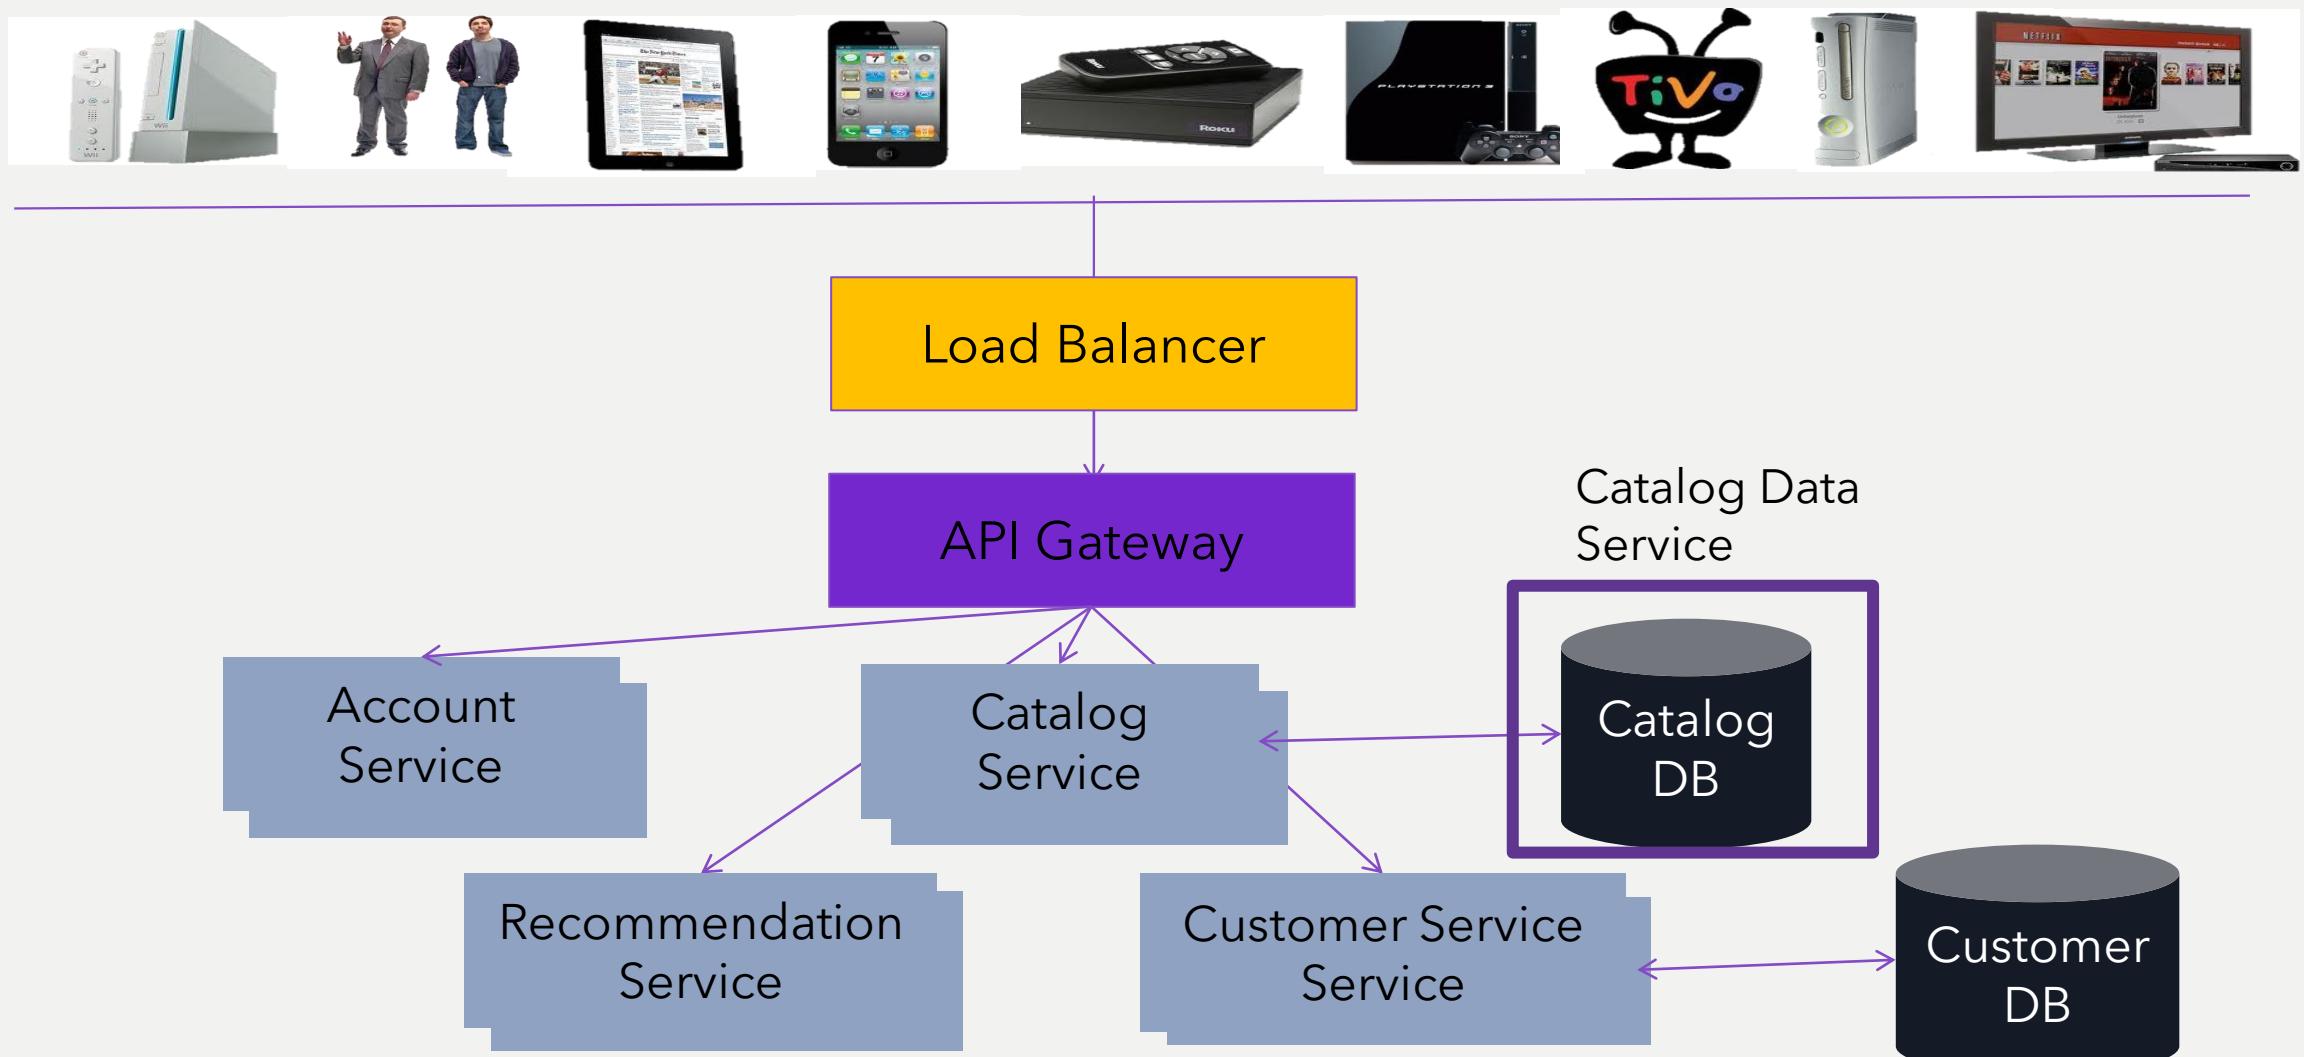
-  Docker – To package and deploy microservices
-  Docker Compose – To test microservices
-  Kubernetes (K8S) – To host application in the cloud
-  Terraform – To build our software factory production infrastructure in the cloud using code (*IaC* approach – a declarative approach).



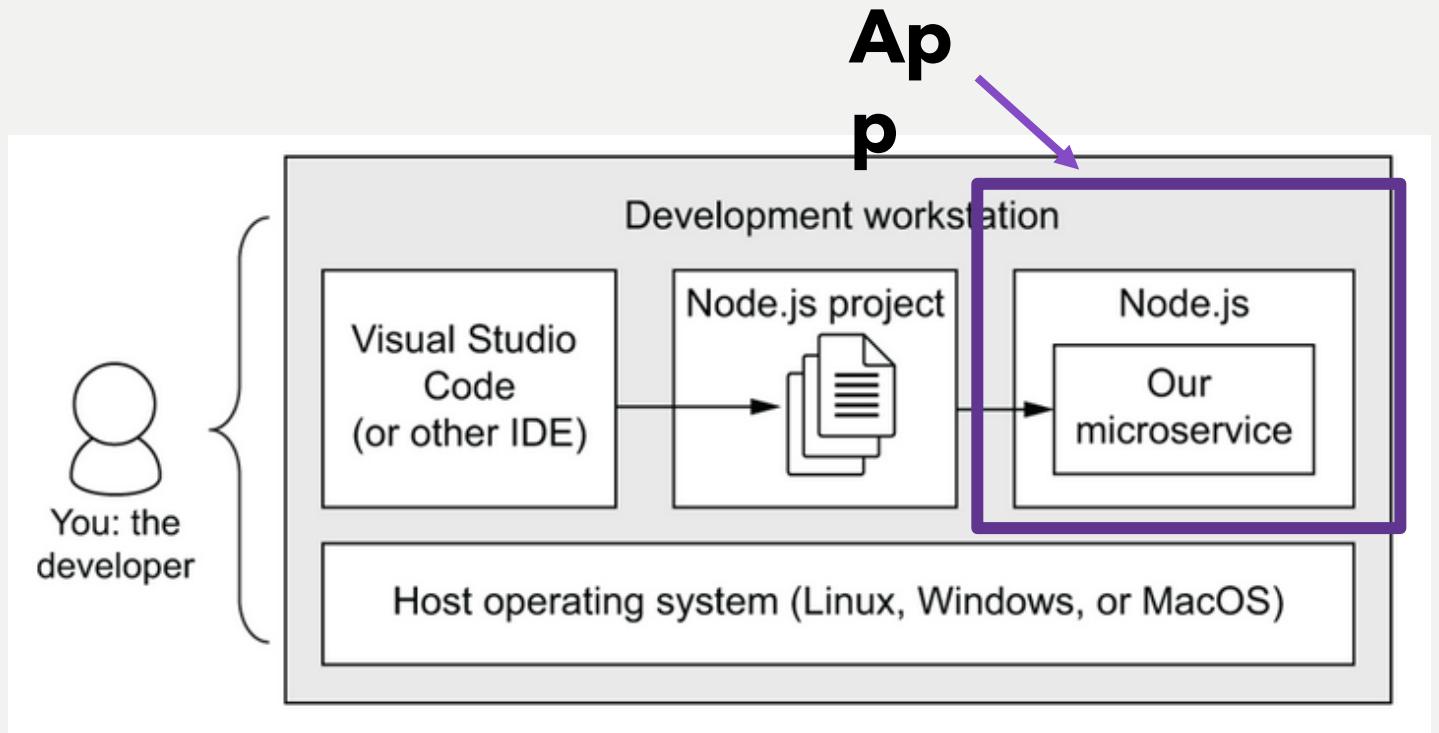
Monolithic Architecture Netflix (prior 2010)

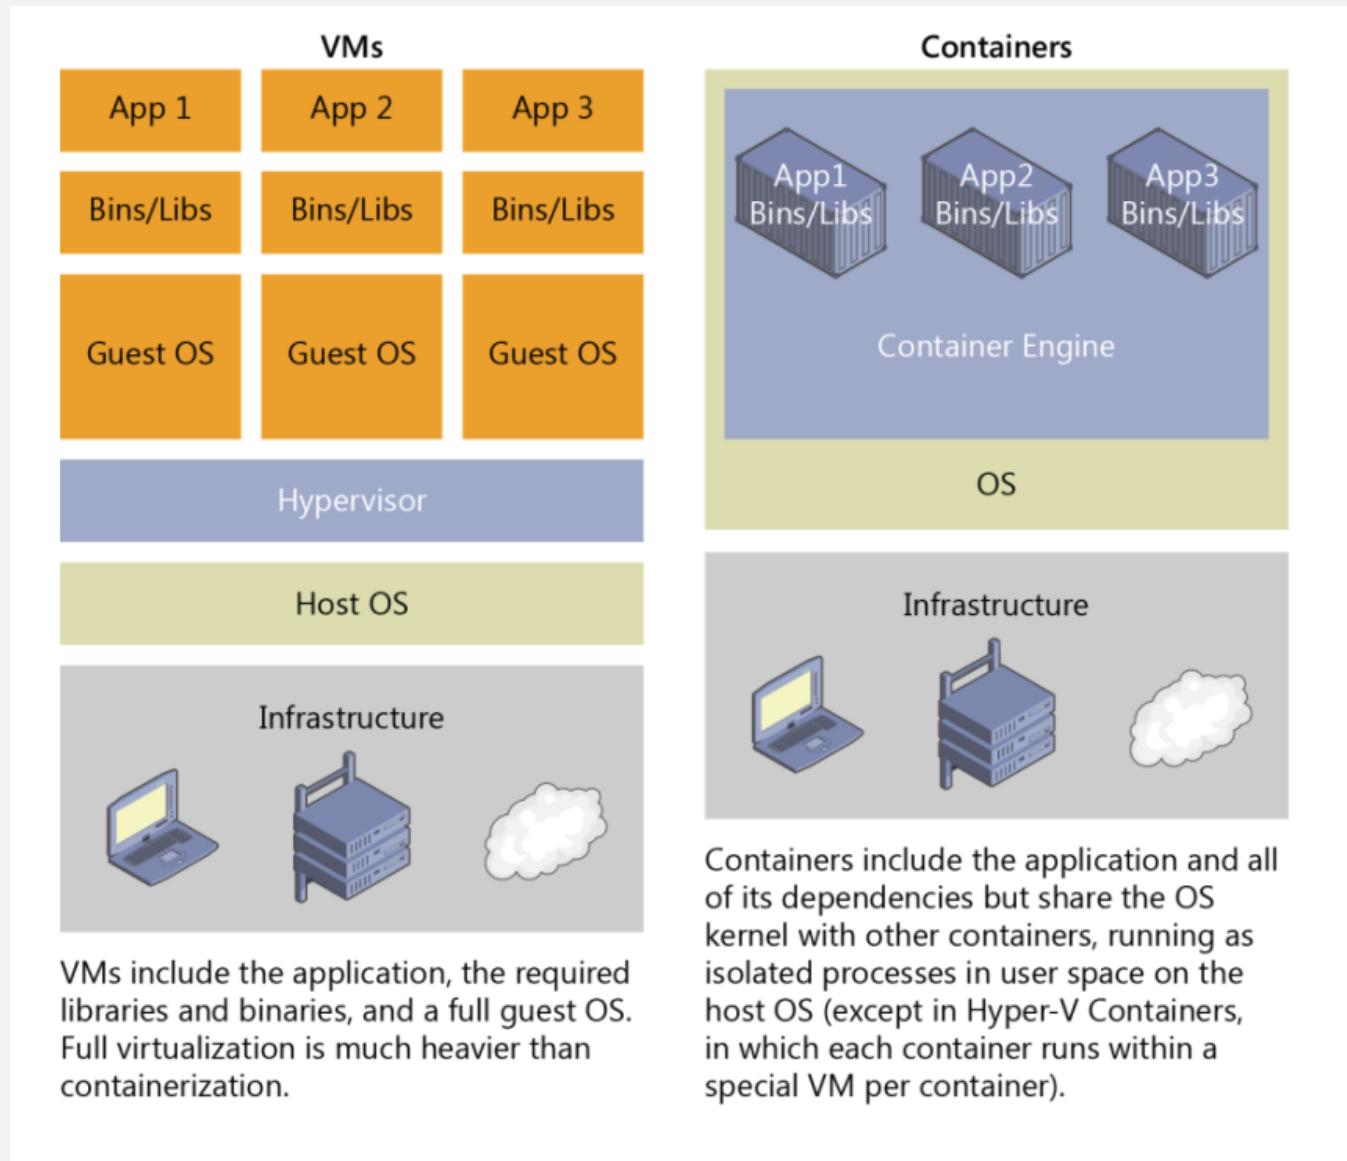


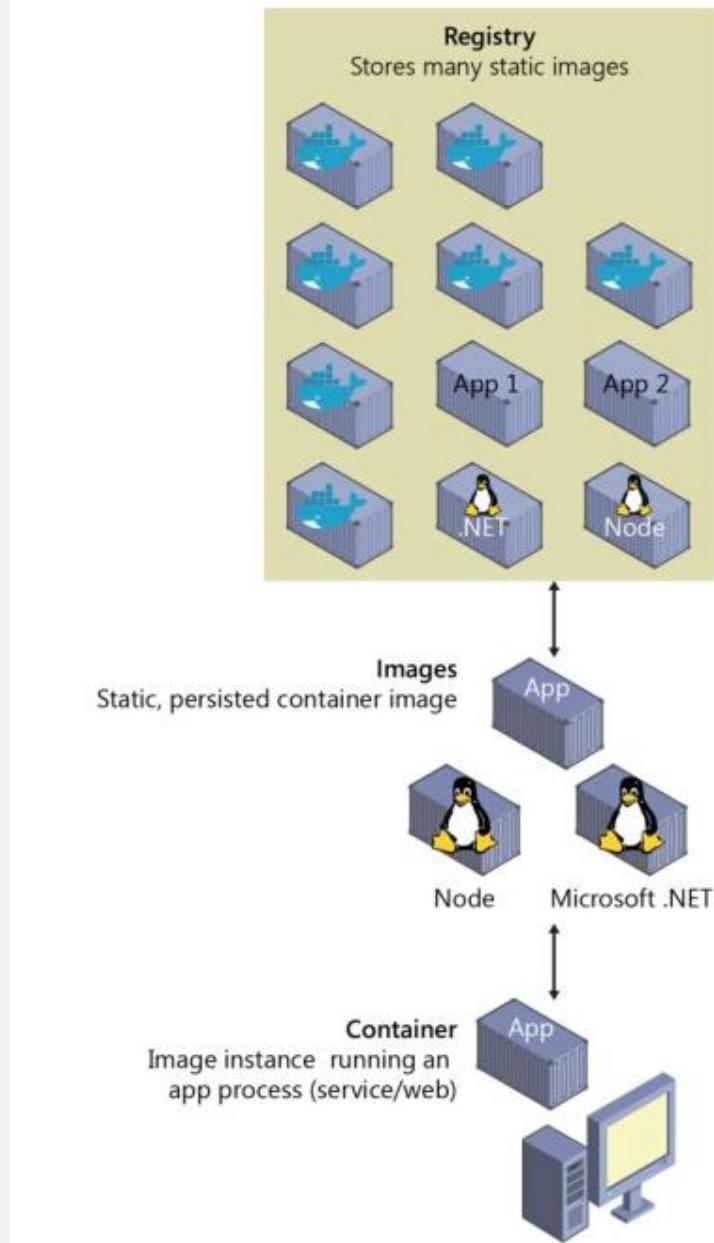
Microservices Architecture (2010+) on AWS

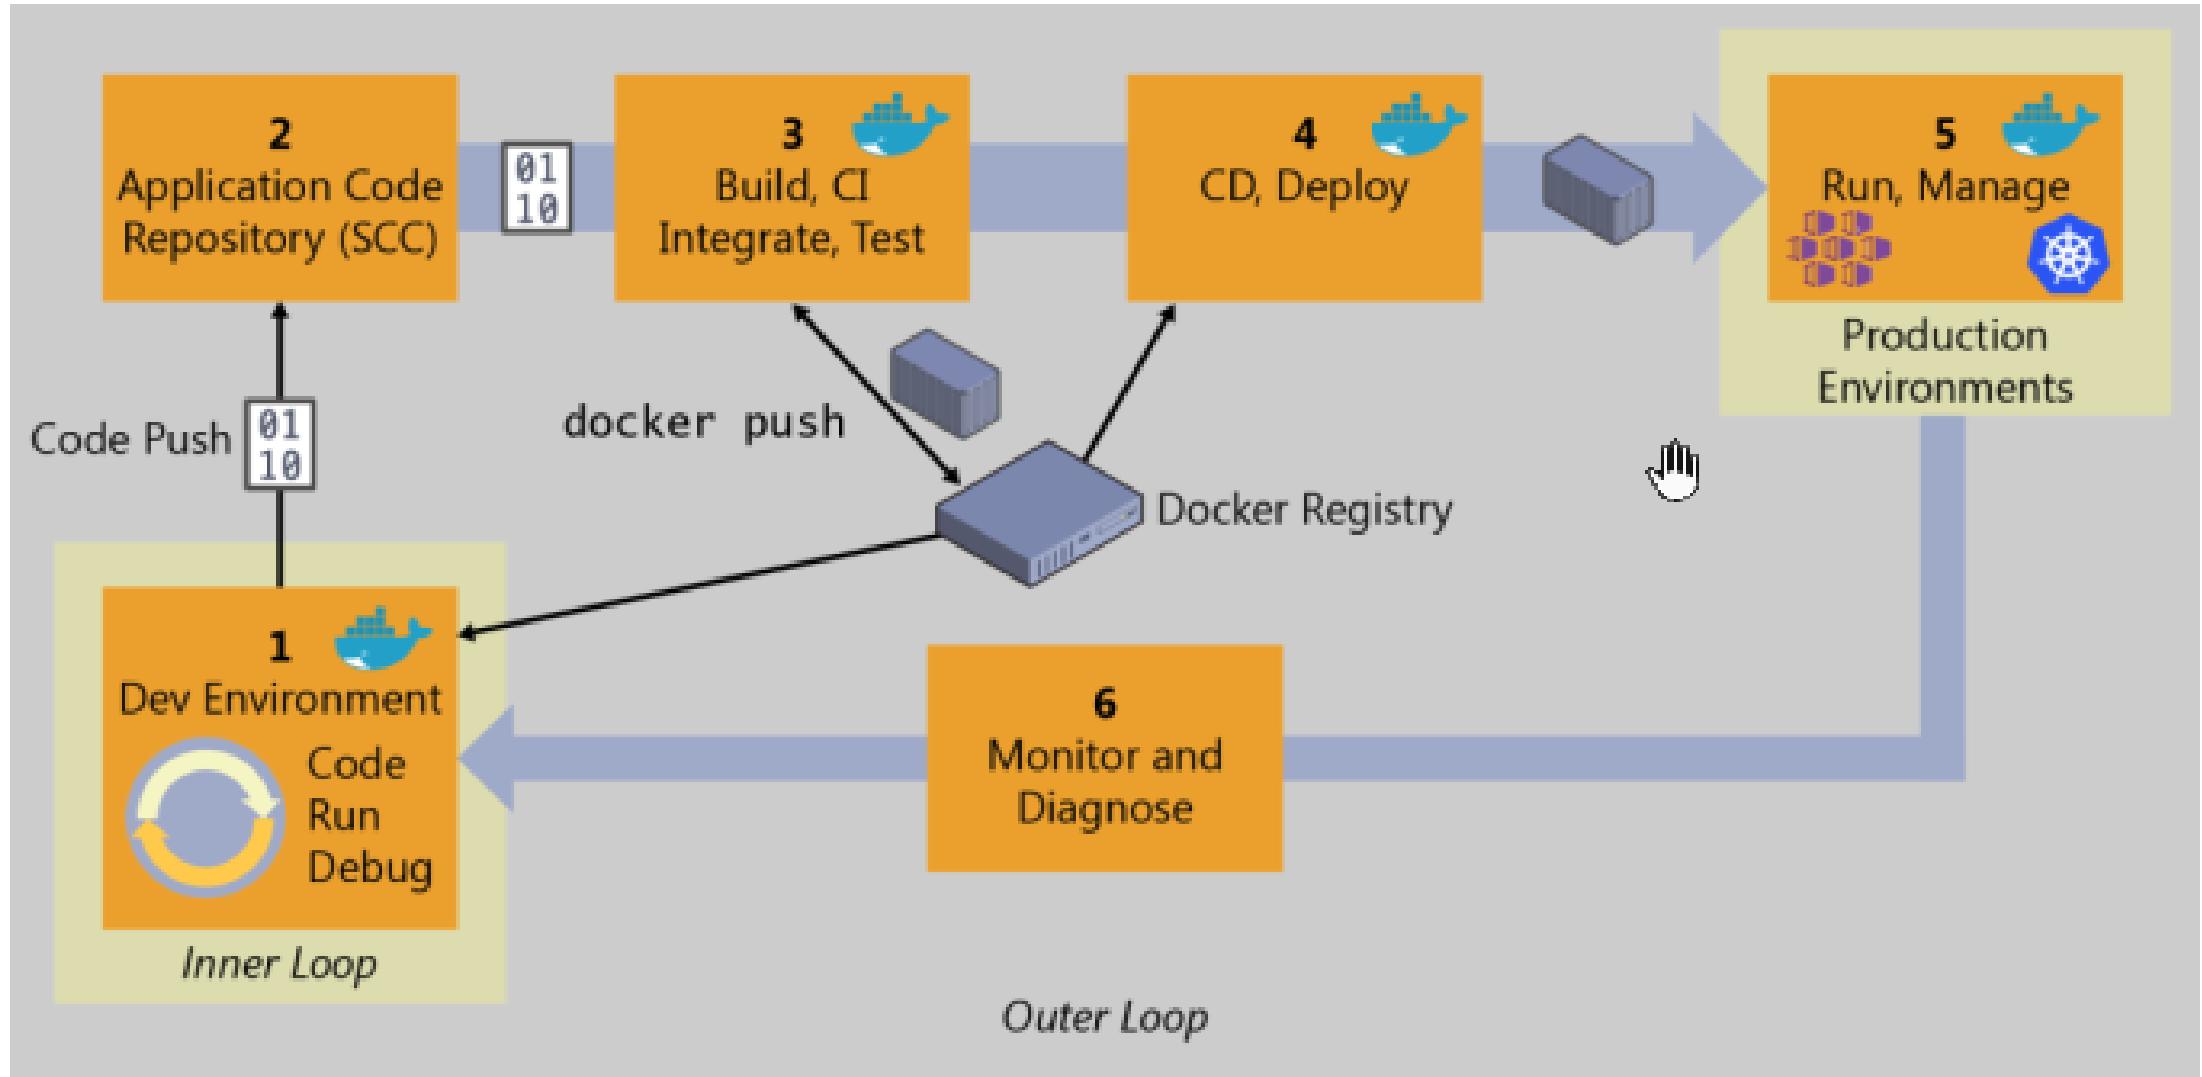


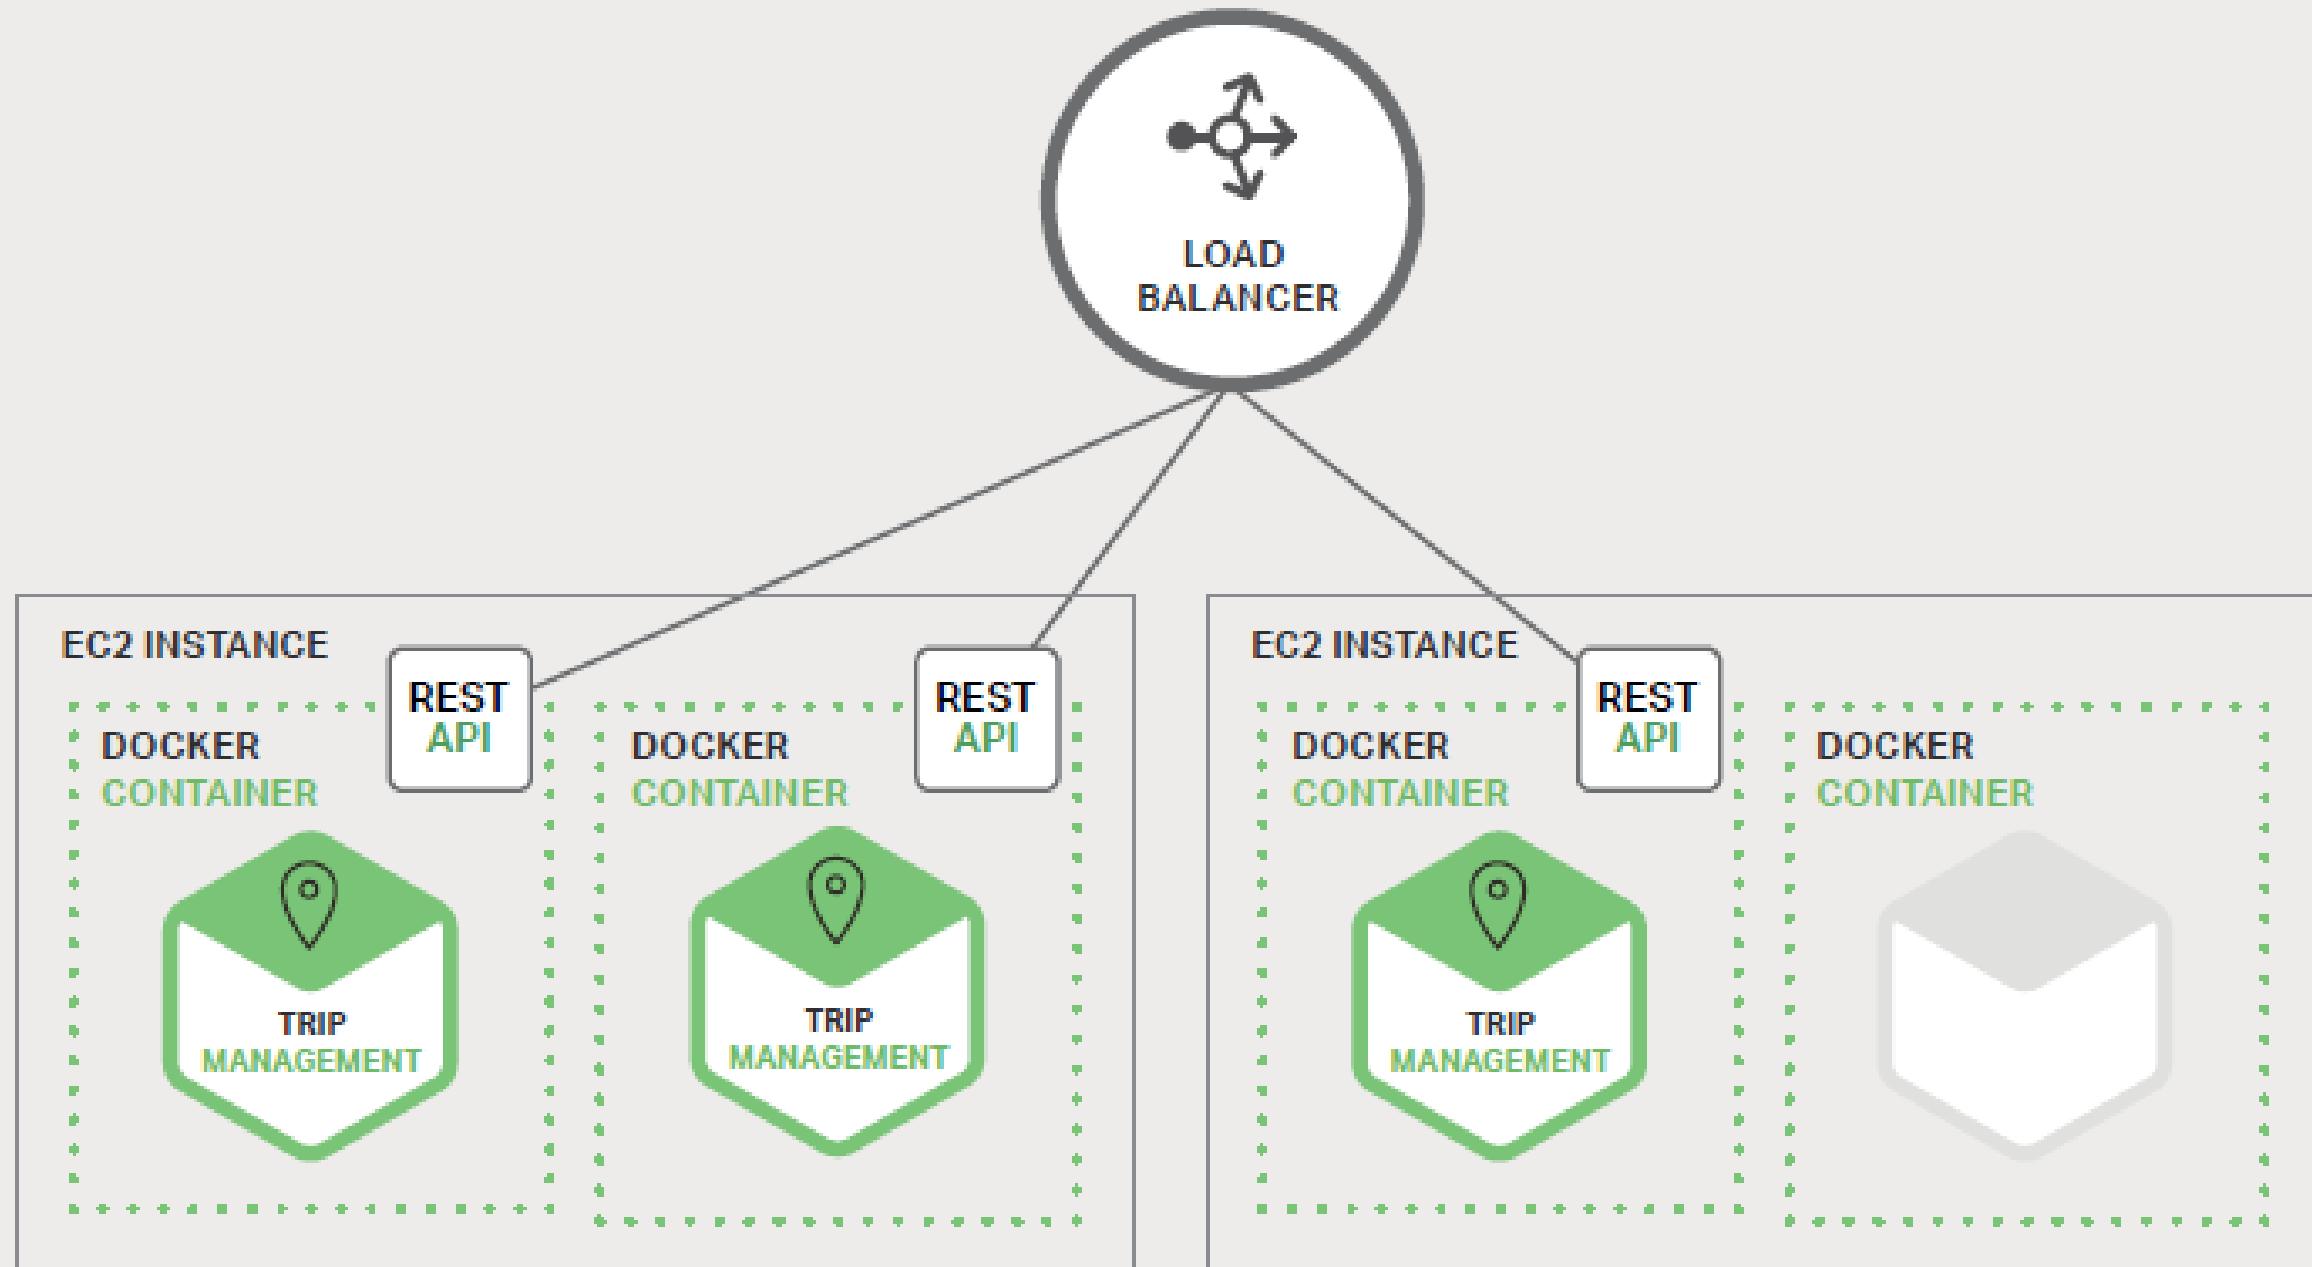
Example of Deploying a Single Service





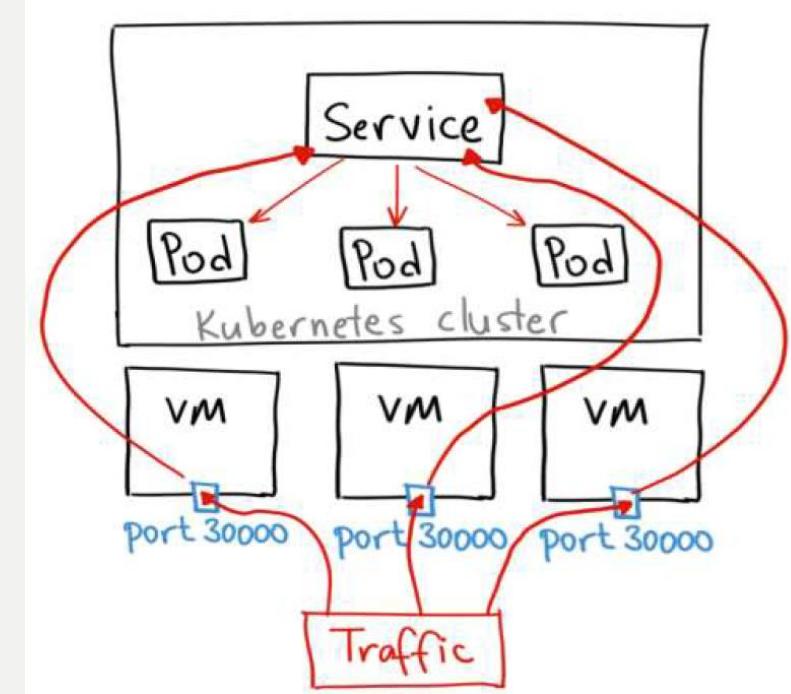
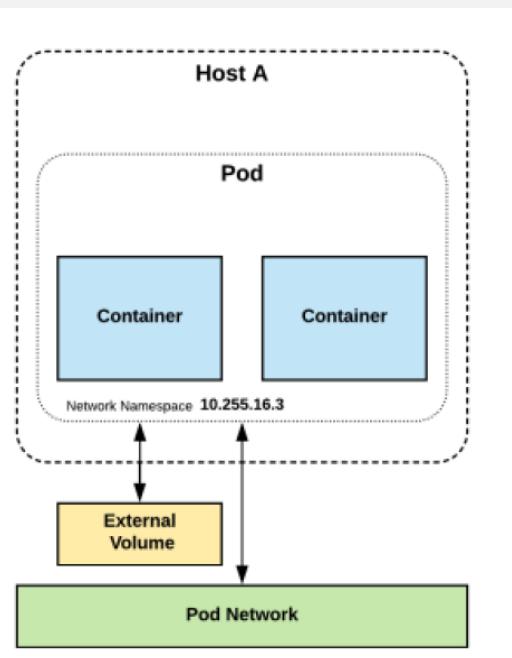






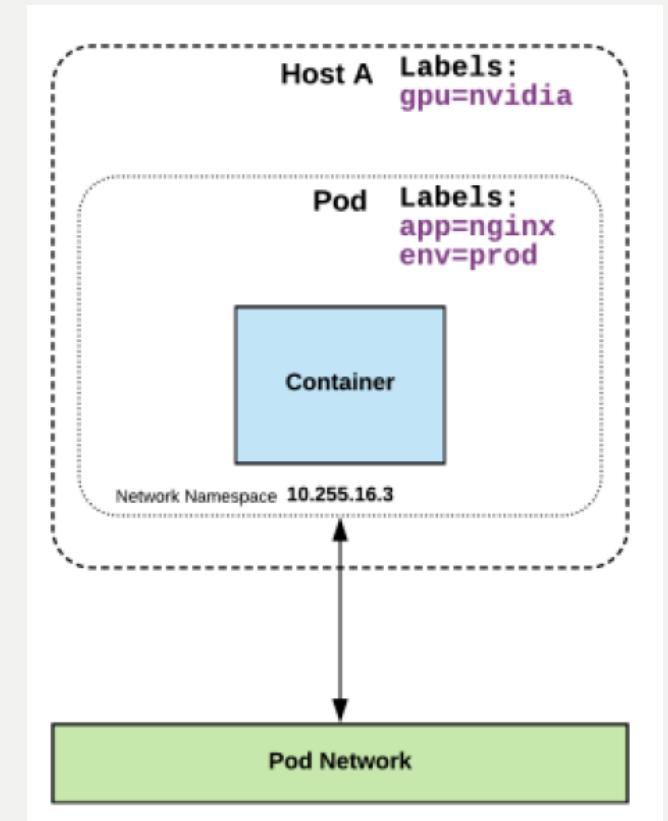
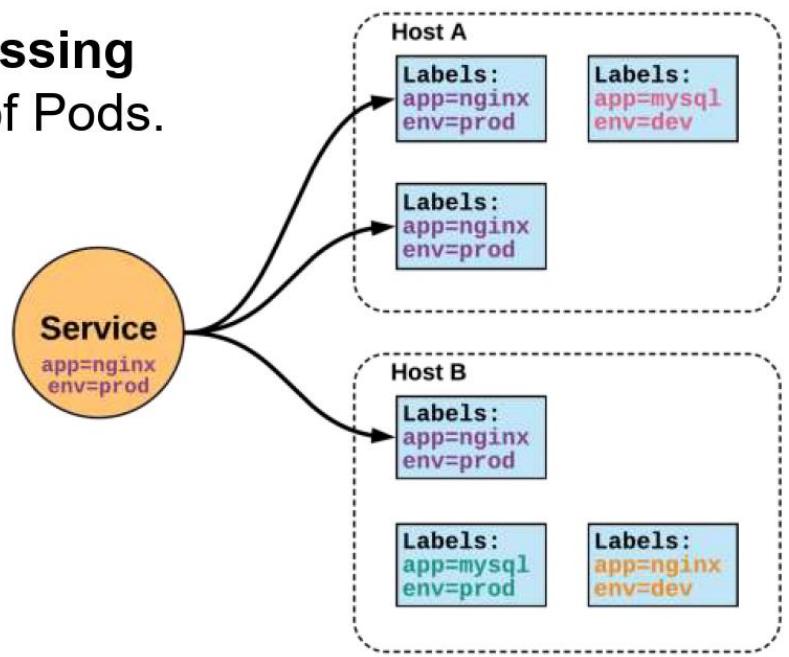
What are Pods?

- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- Pods are **one or MORE containers** that share volumes and namespace.
- They are also ephemeral!

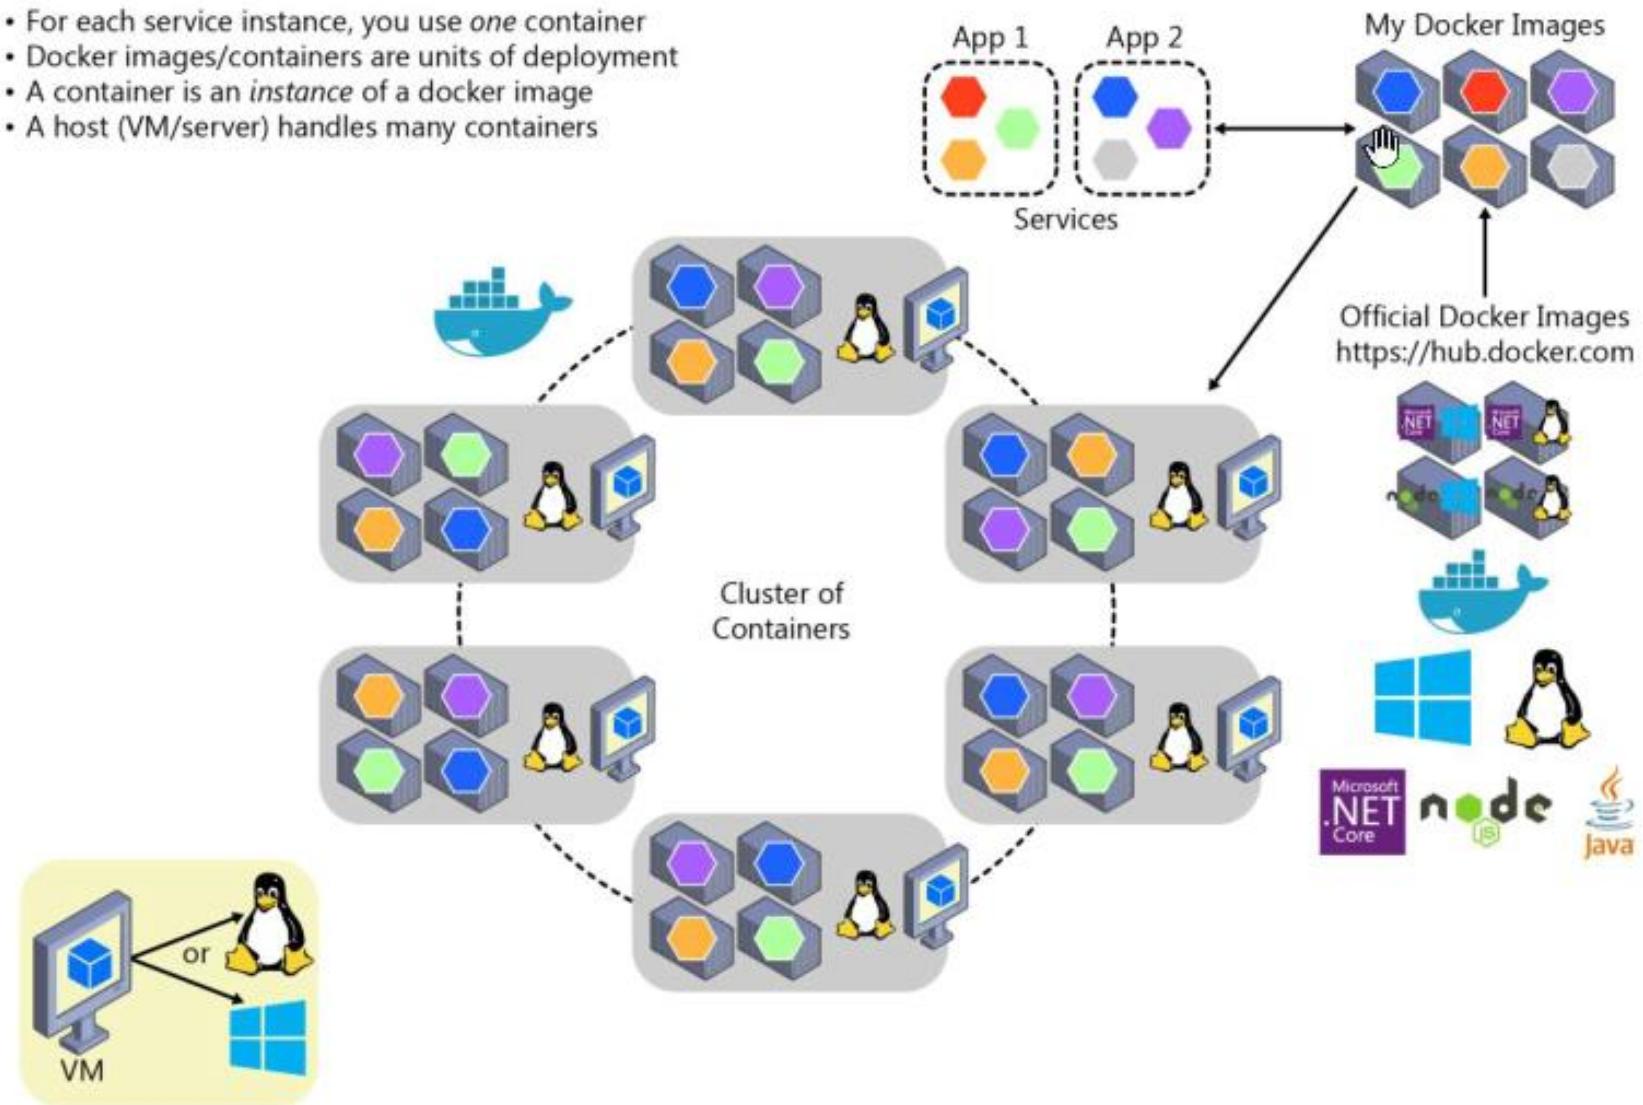


Services

- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
 - static cluster IP
 - static namespaced DNS name



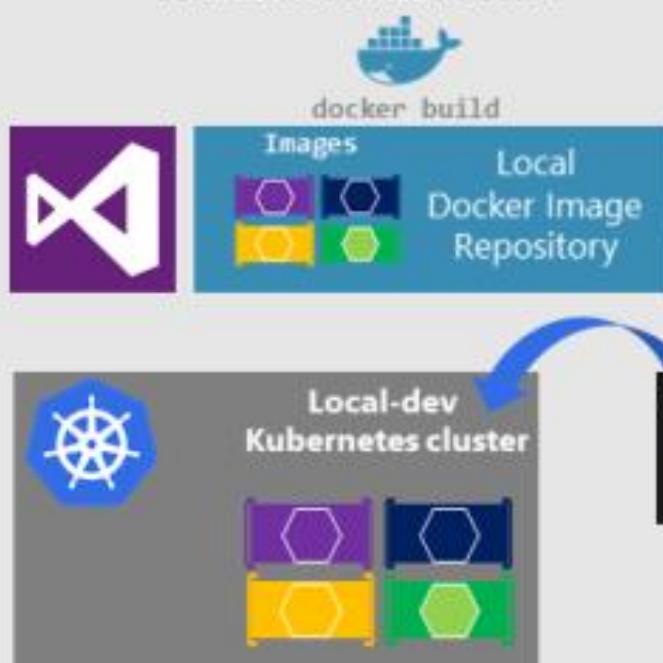
- For each service instance, you use *one* container
- Docker images/containers are units of deployment
- A container is an *instance* of a docker image
- A host (VM/server) handles many containers





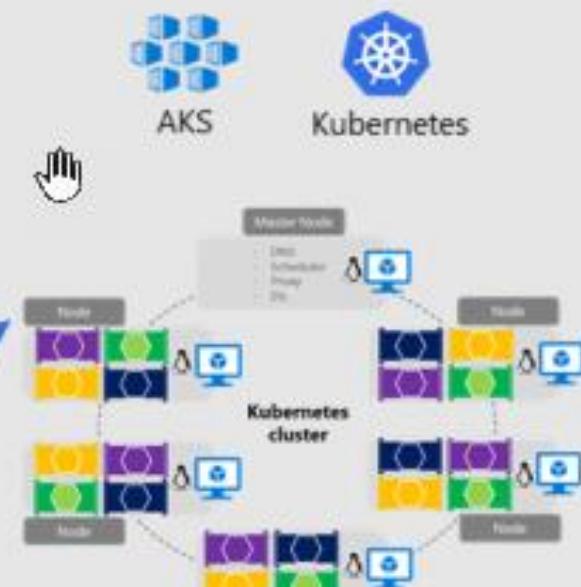
PC/Development Environment

Windows 10 or macOS
'Docker Desktop' with
local-dev Kubernetes cluster

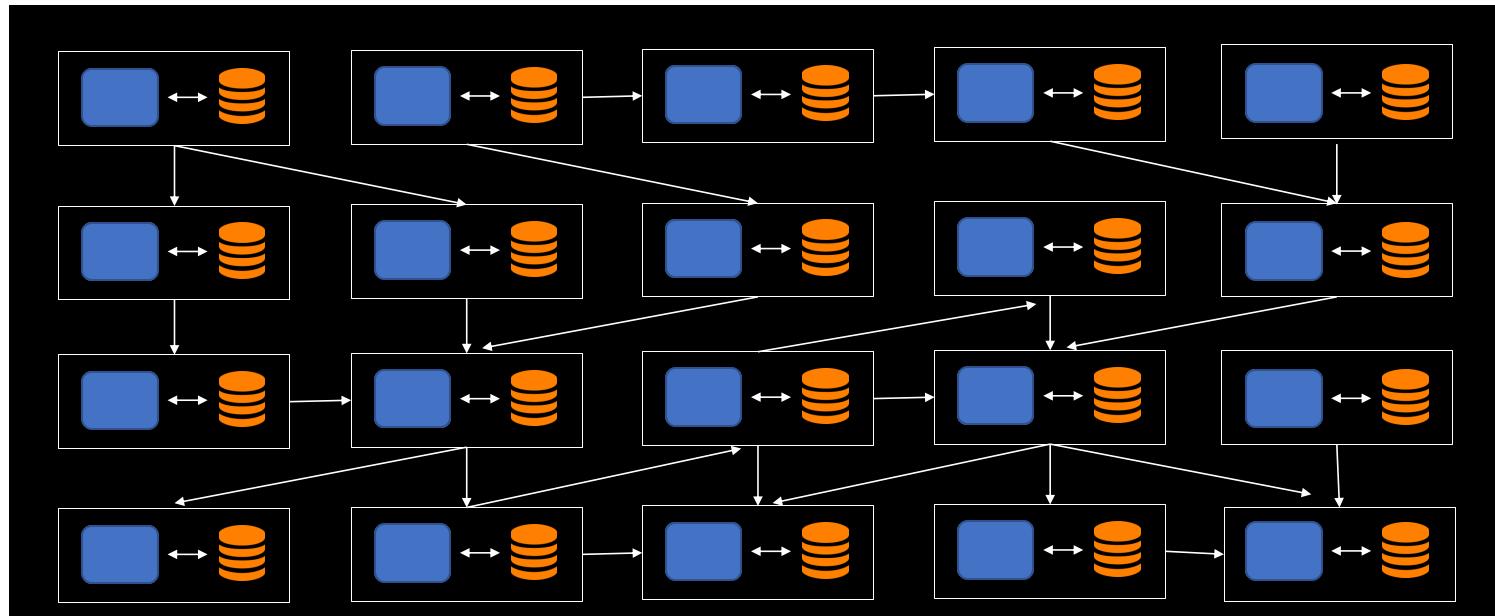


Azure Kubernetes Service (AKS)

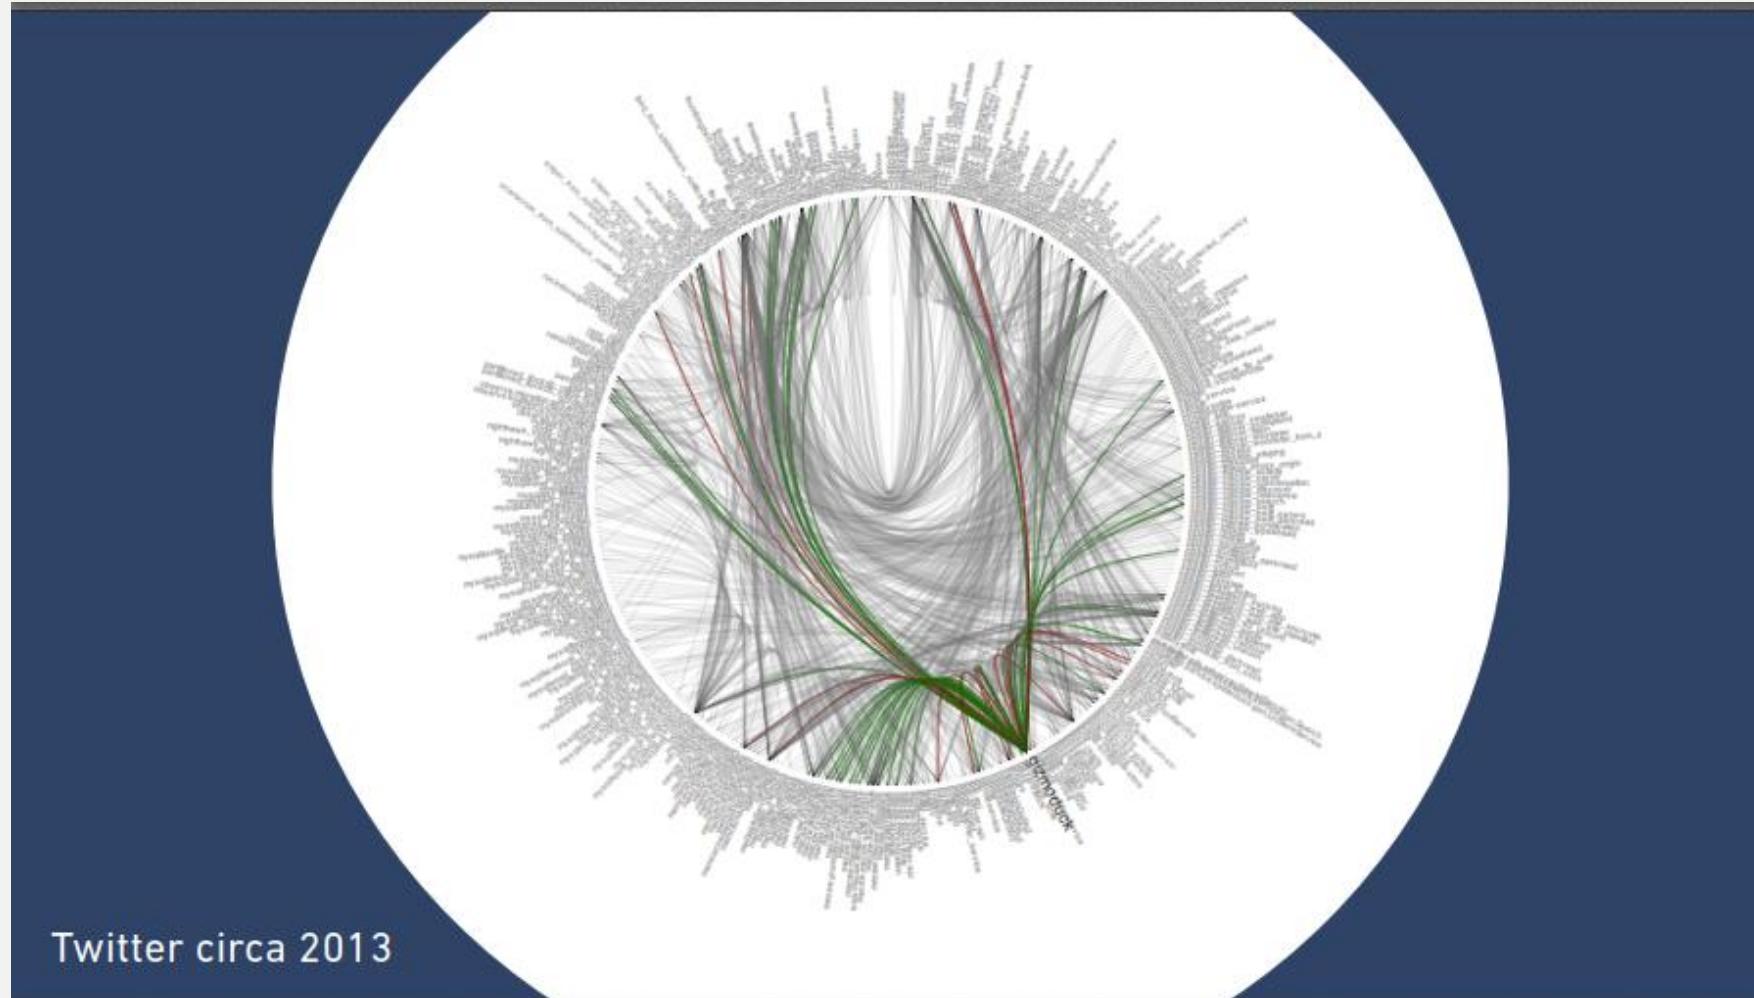
"Managed Kubernetes" for production



Microservices can quickly multiply in complexity

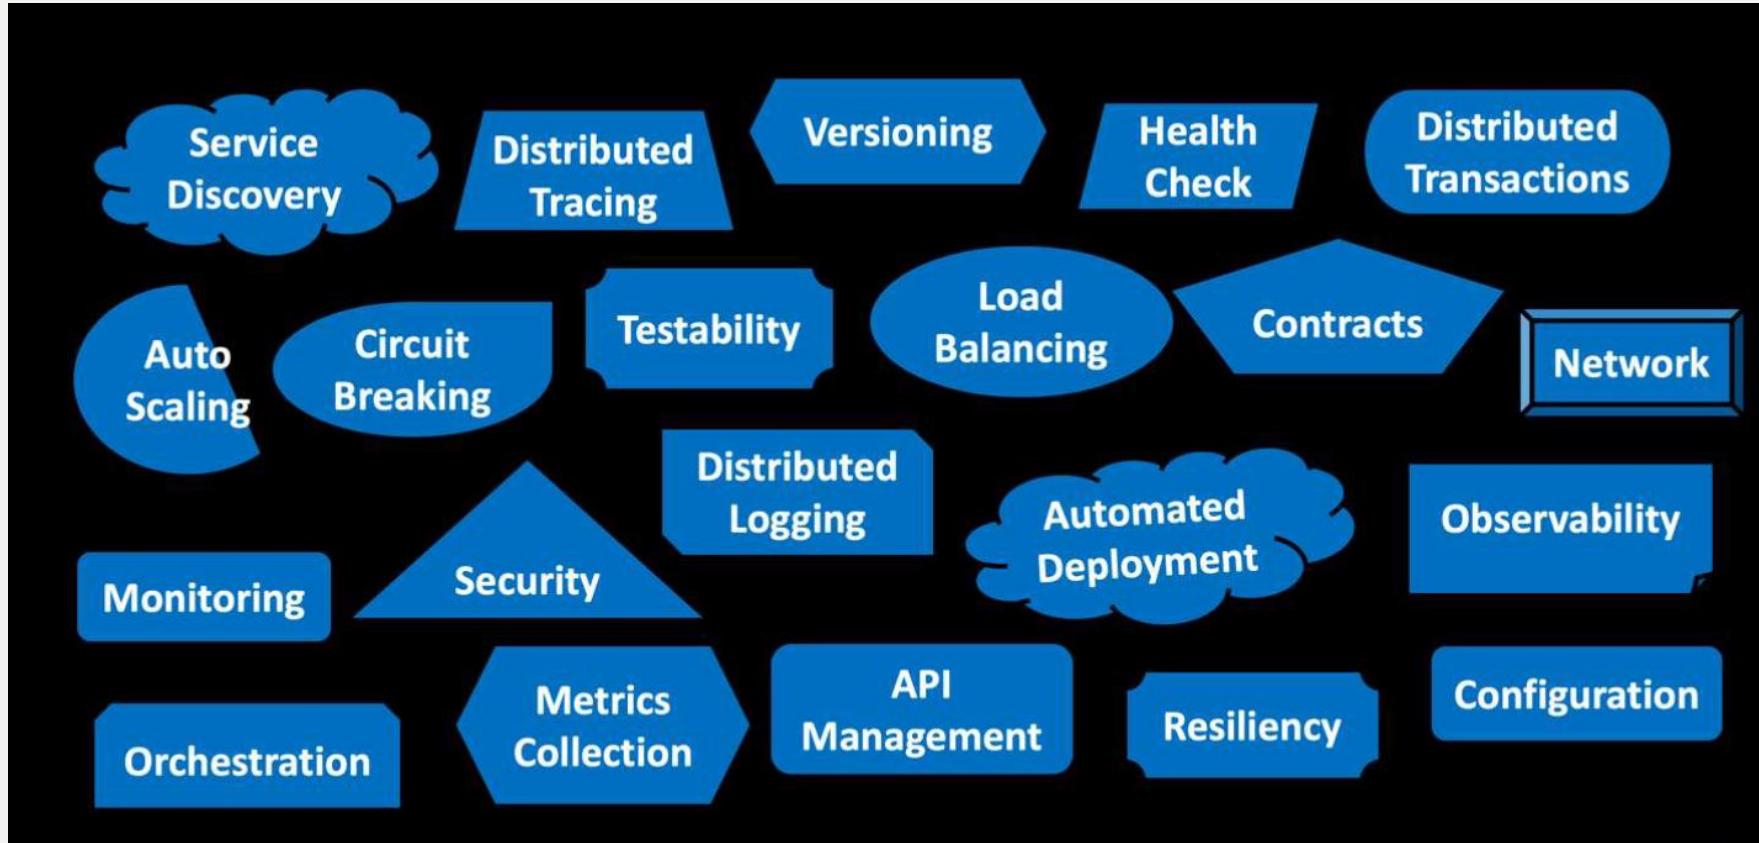


Microservices can lead to lot of complexity – 100k+ microservices in a large enterprise system



Open Problems to Solve in Microservices Architecture & Containers

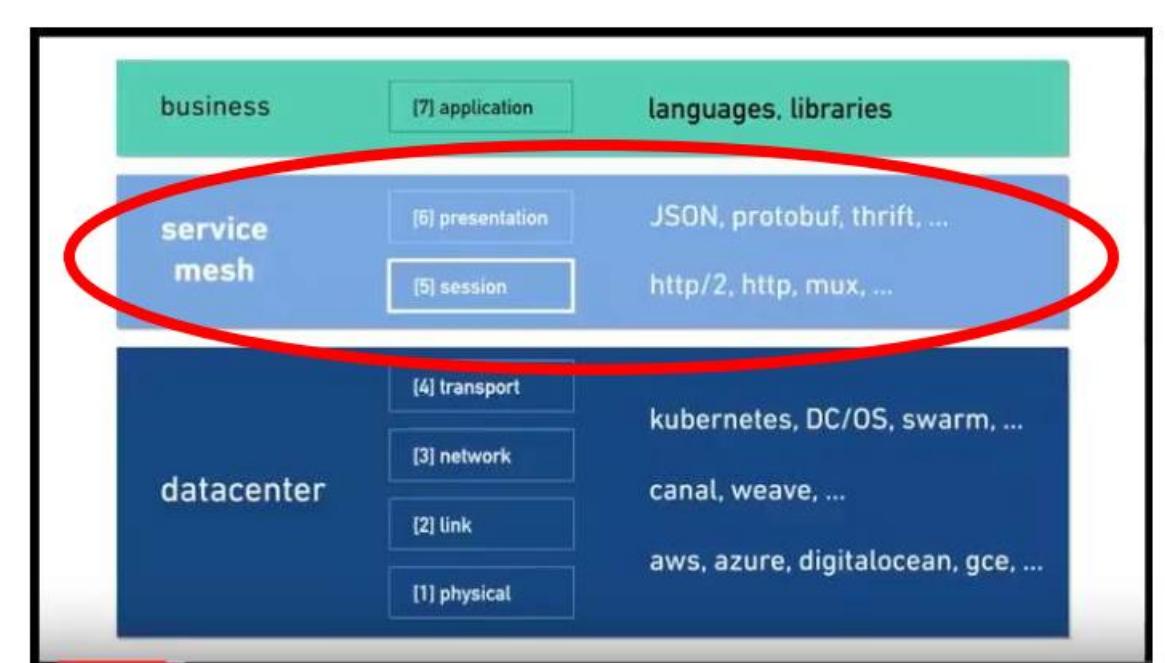
(Amazon, Google, Microsoft are actively working on this)



Service Mesh solved a lot of home-grown solutions

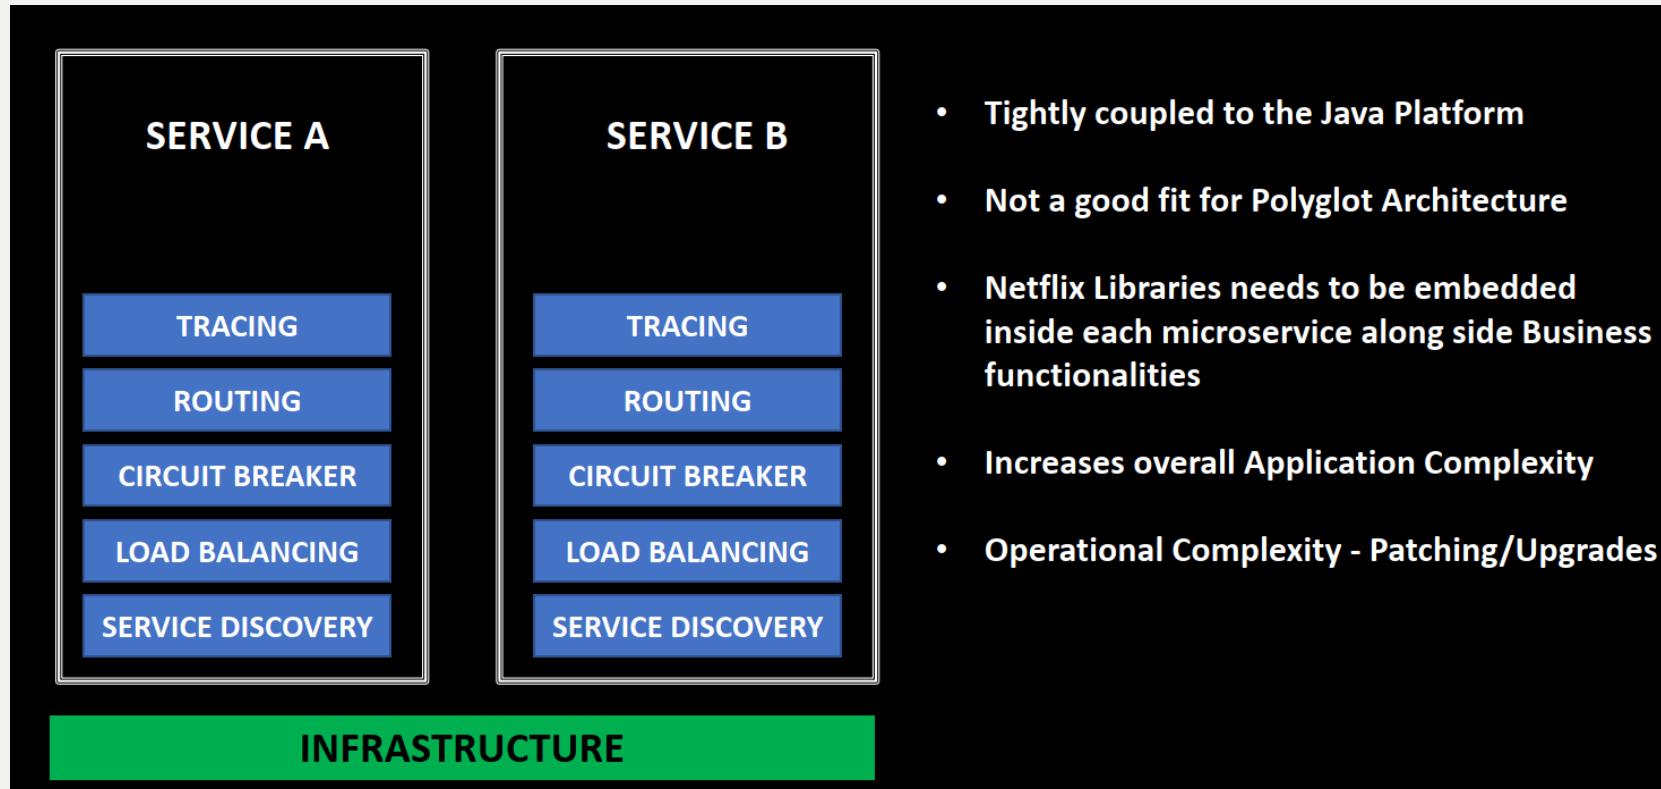
- Netflix
 - Karyon + HTTP/JSON or RxNetty RPC + Eureka + Hystrix + ...
- Twitter
 - Finagle + Thrift + ZooKeeper + Zipkin
- AirBnB
 - HTTP/JSON + SmartStack + ZooKeeper + Charon/Dyno
- Google
 - Stubby + GSLB + GFE + Dapper

Before Service Mesh

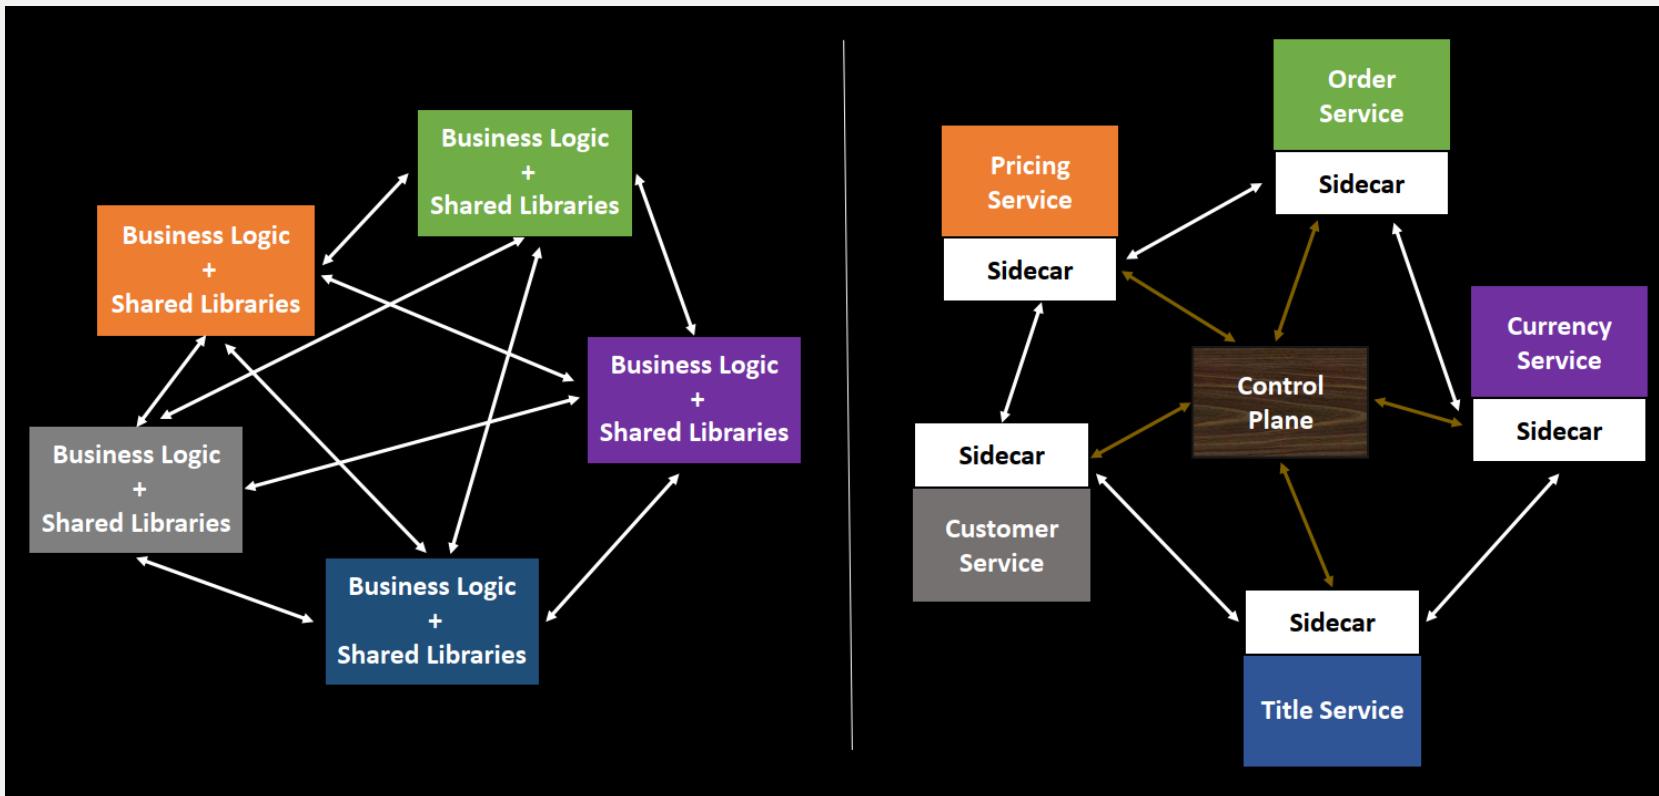


After Service Mesh

Why didn't “home-grown” Netflix approach work in general?



Service Mesh



Service Mesh - Benefits

- Normalises naming and adds logical routing
 - user-service -> AWS-us-east-1a/prod/users/v4
 - Adds traffic shaping and traffic shifting
 - Load balancing
 - Deploy control
 - Per-request routing (shadowing, fault injection, debug)
 - Adds baseline reliability
 - Health checks, timeouts/deadlines, circuit breaking, and retry (budgets)
-
- Increased security
 - Transparent mutual TLS
 - Policies (service Access Control Lists - ACL)
 - Observability
 - Top-line metrics like request volume, success rates and latencies
 - Distributed tracing
 - Sane defaults (to protect the system)
 - Options to tune



What is a Service Mesh - Generic Definition

- A Service Mesh is the connective tissue between your *microservices* that adds the capabilities of traffic control, service discovery, load balancing, resilience, observability, security,

Origins of ISTIO

Collaboration
between
Google IBM and
Lyft

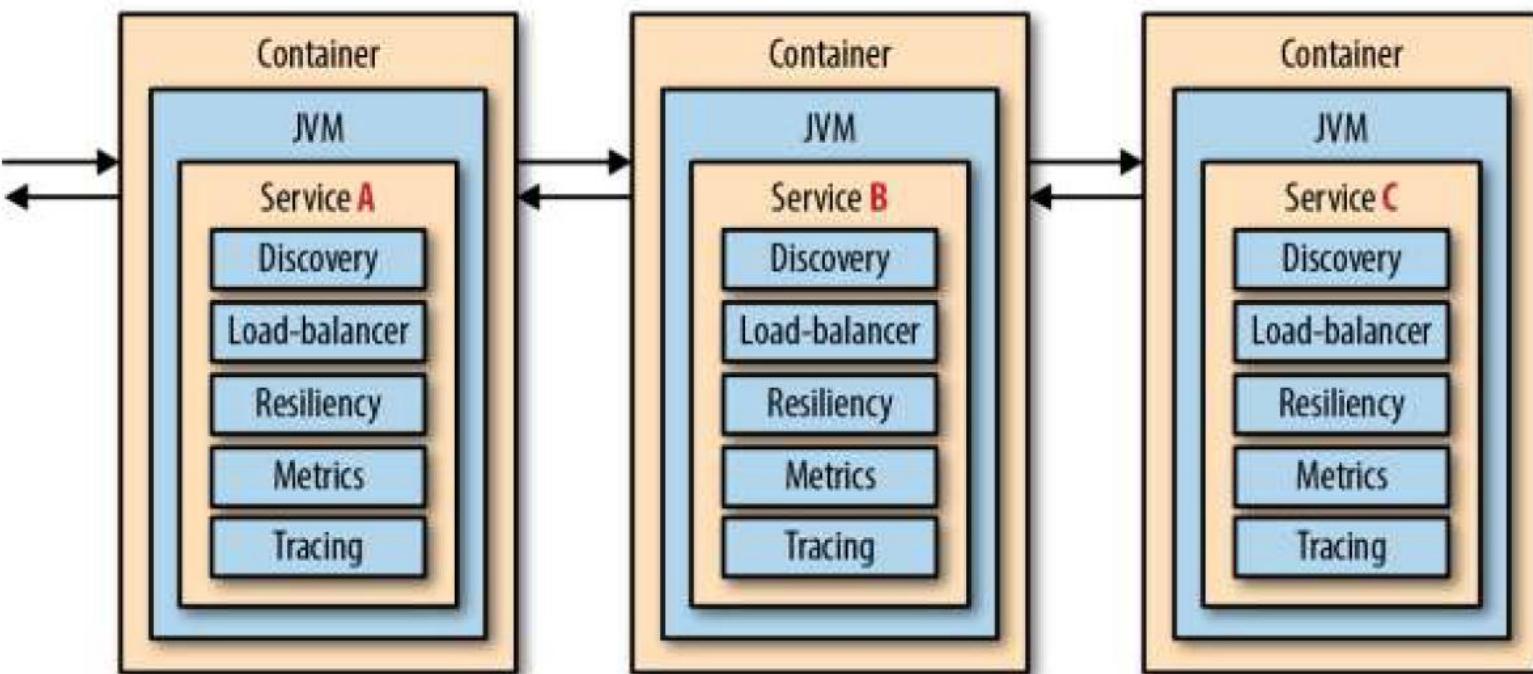
Open-source
service mesh for
microservices

How did Service Mesh Help Solve these Open Problems?

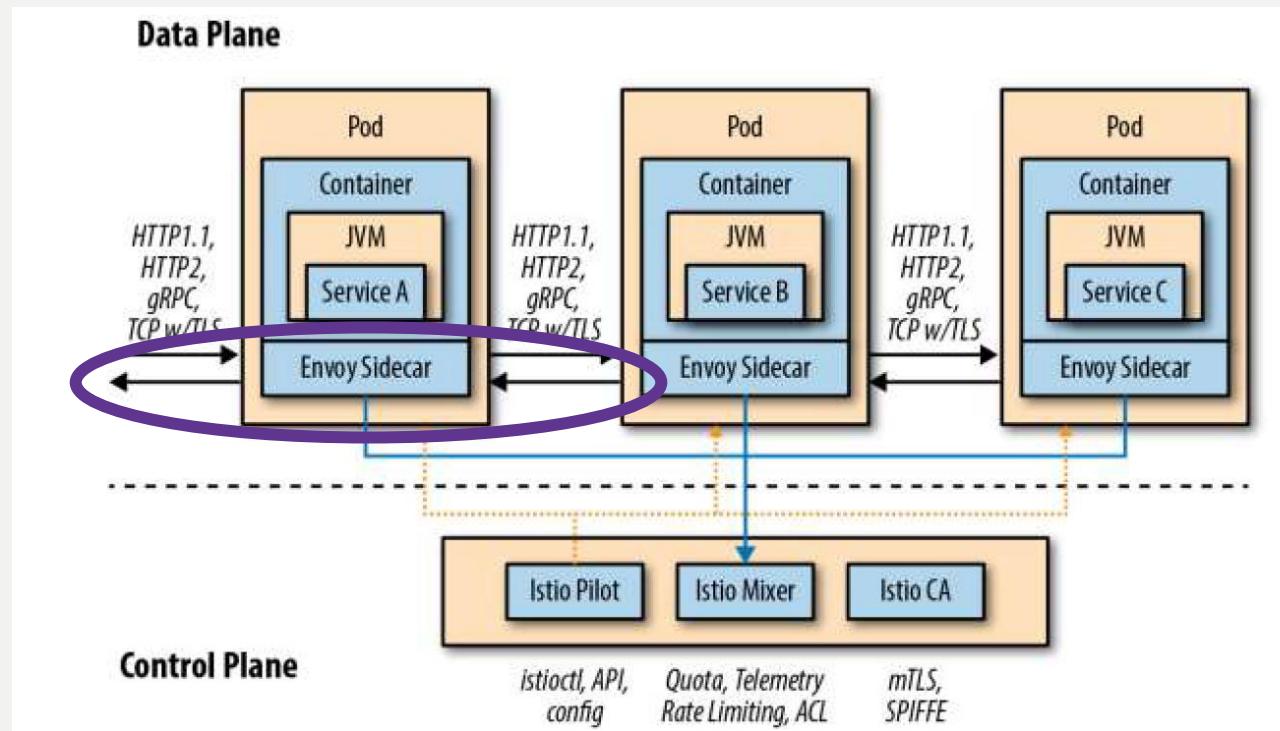
Real World v. Practical World

Fallacy	Solutions
The network is reliable	Automatic Retries, Message Queues
Latency is zero	Caching Strategy, Bulk Requests, Deploy in AZs near client
Bandwidth is infinite	Throttling Policy, Small payloads with Microservices
The network is secure	Network Firewalls, Encryption, Certificates, Authentication
Topology does not change	No hardcoding IP, Service Discovery Tools
There is one administrator	DevOps Culture eliminates Bus Factor
Transport cost is zero	Standardized protocols like JSON, Cost Calculation
The network is homogenous	Circuit Breaker, Retry and Timeout Design Pattern

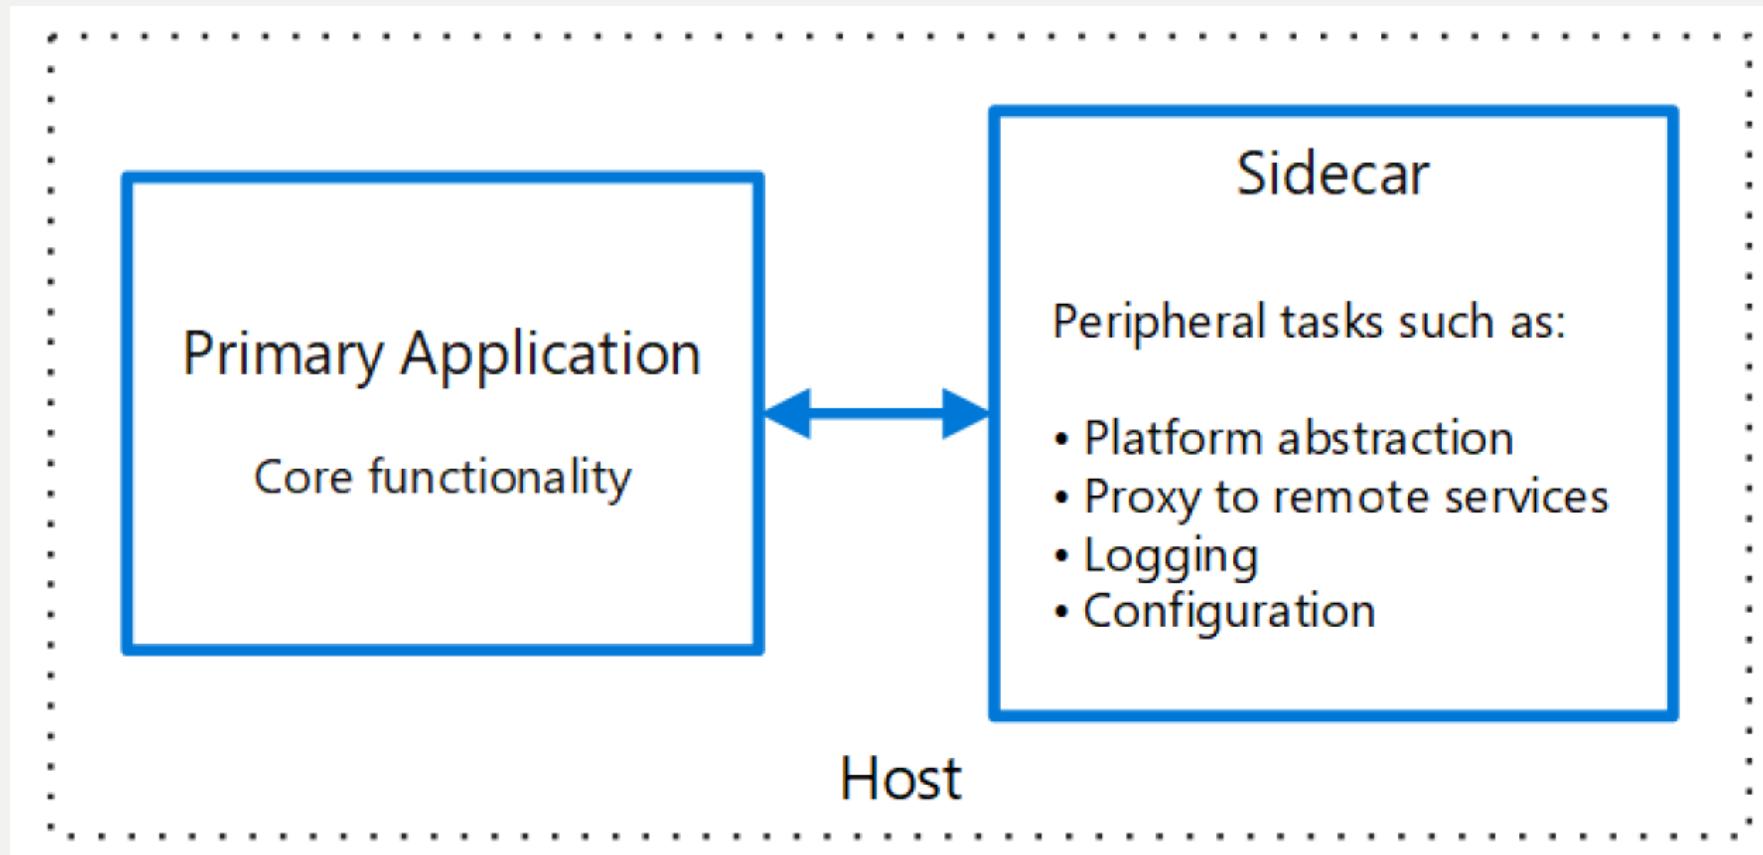
Microservices before ISTIO



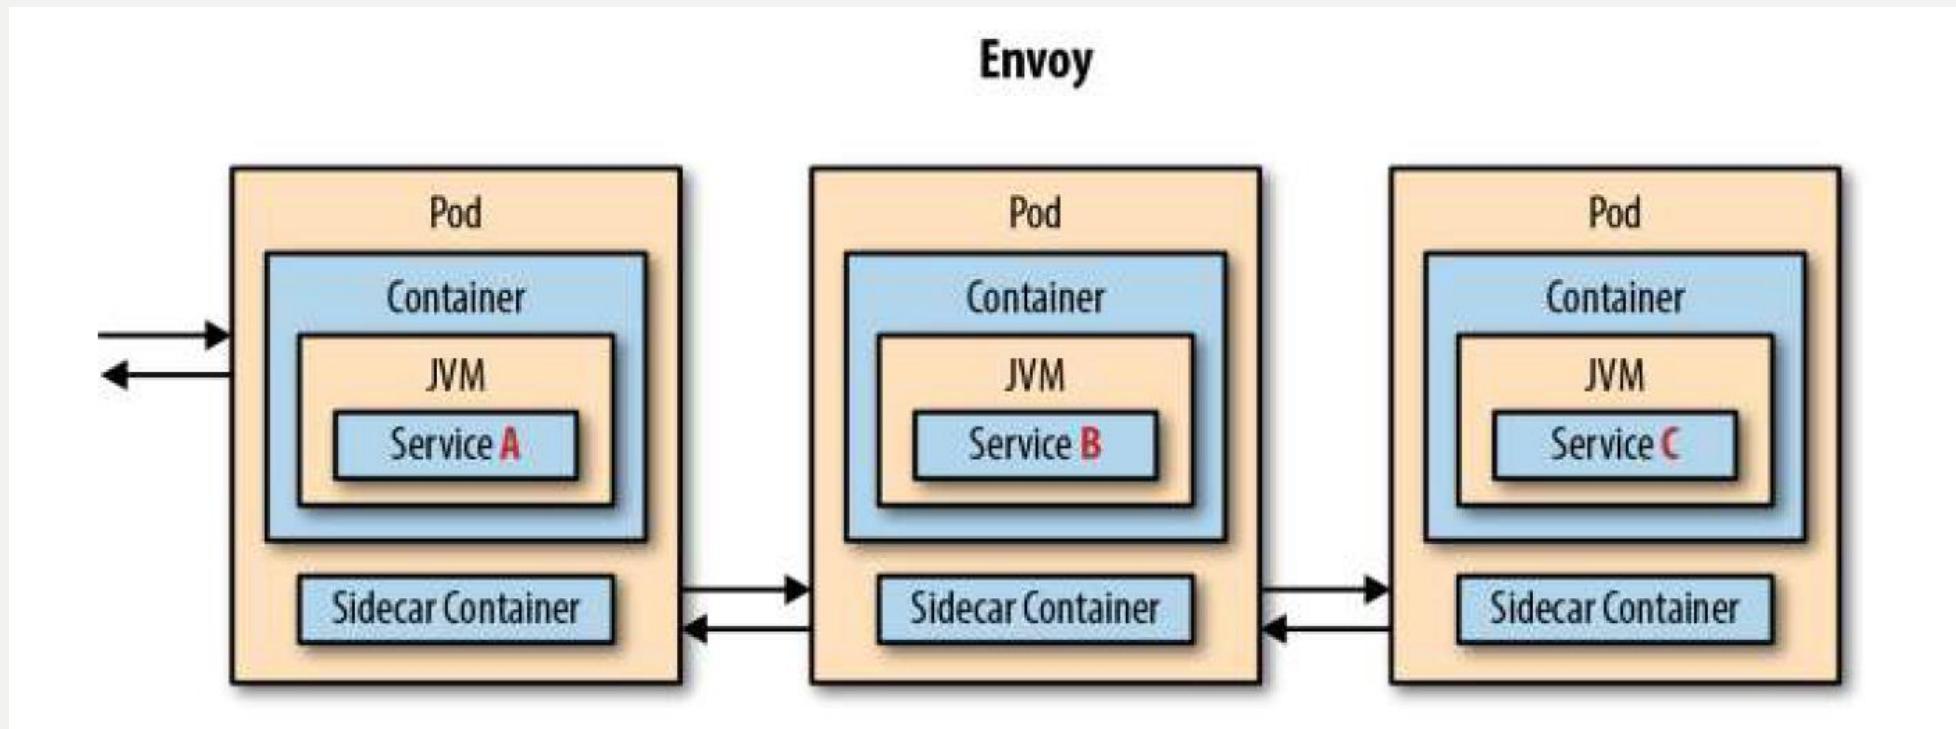
Microservices after ISTIO



What is the Sidecar Pattern?



Envoy is a Sidecar

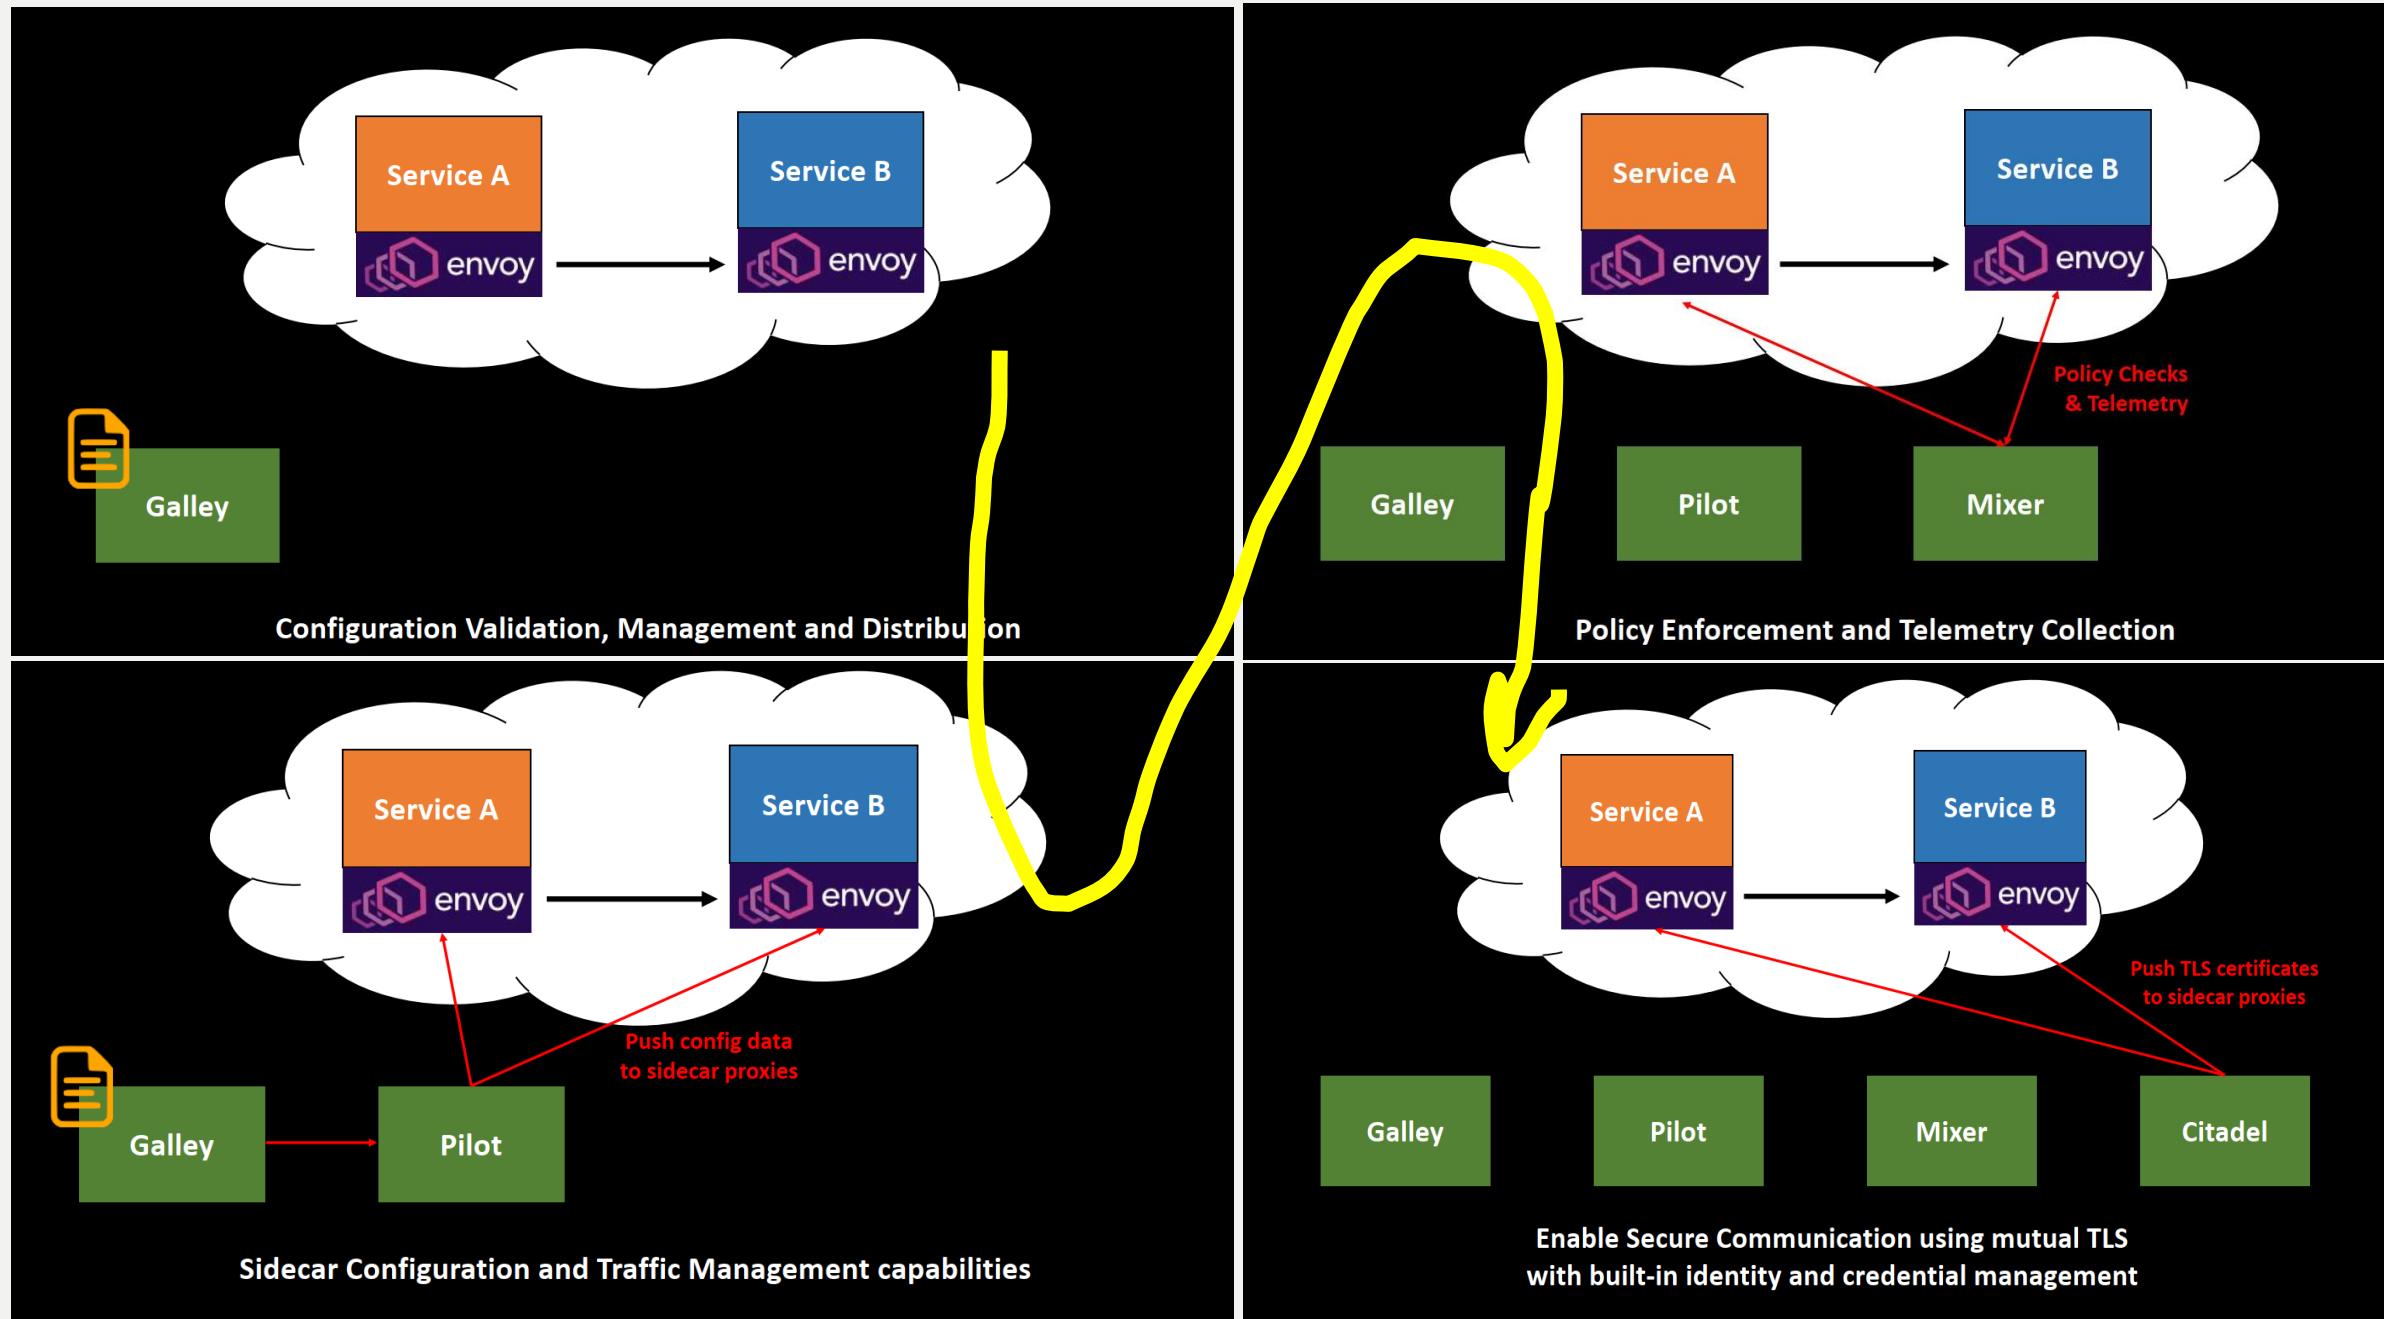


What are the three primary Istio services?

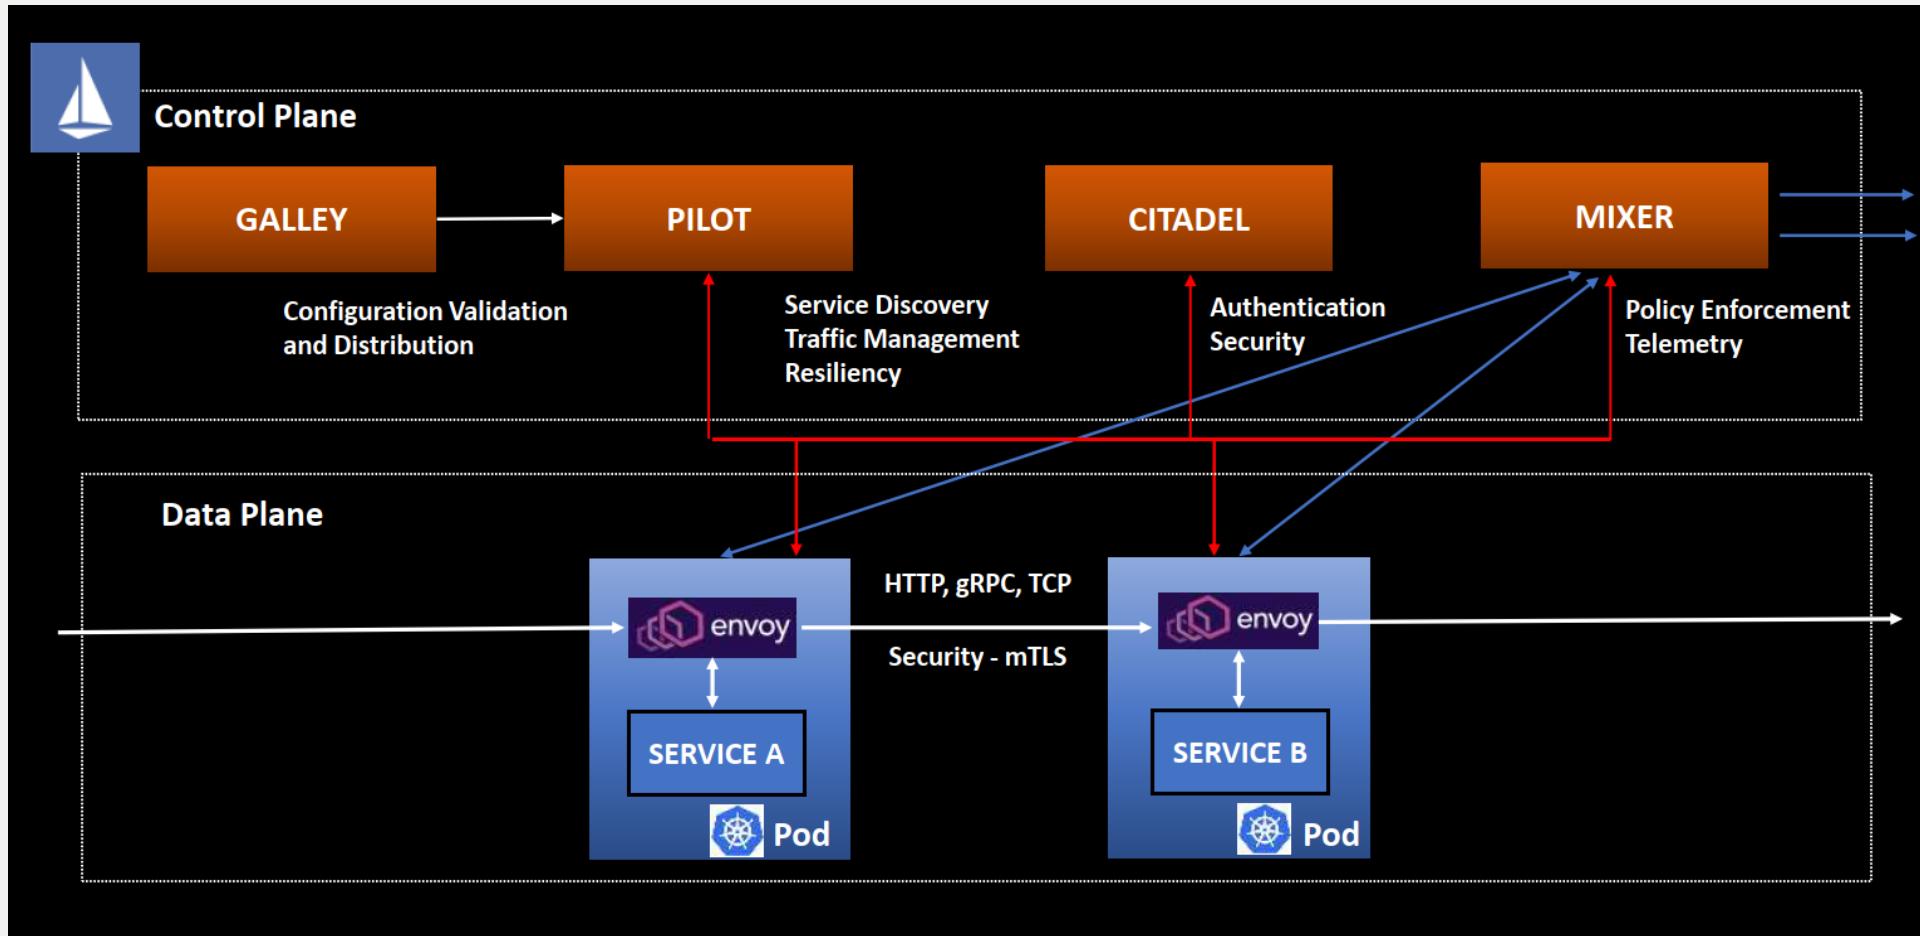
Pilot

Mixer

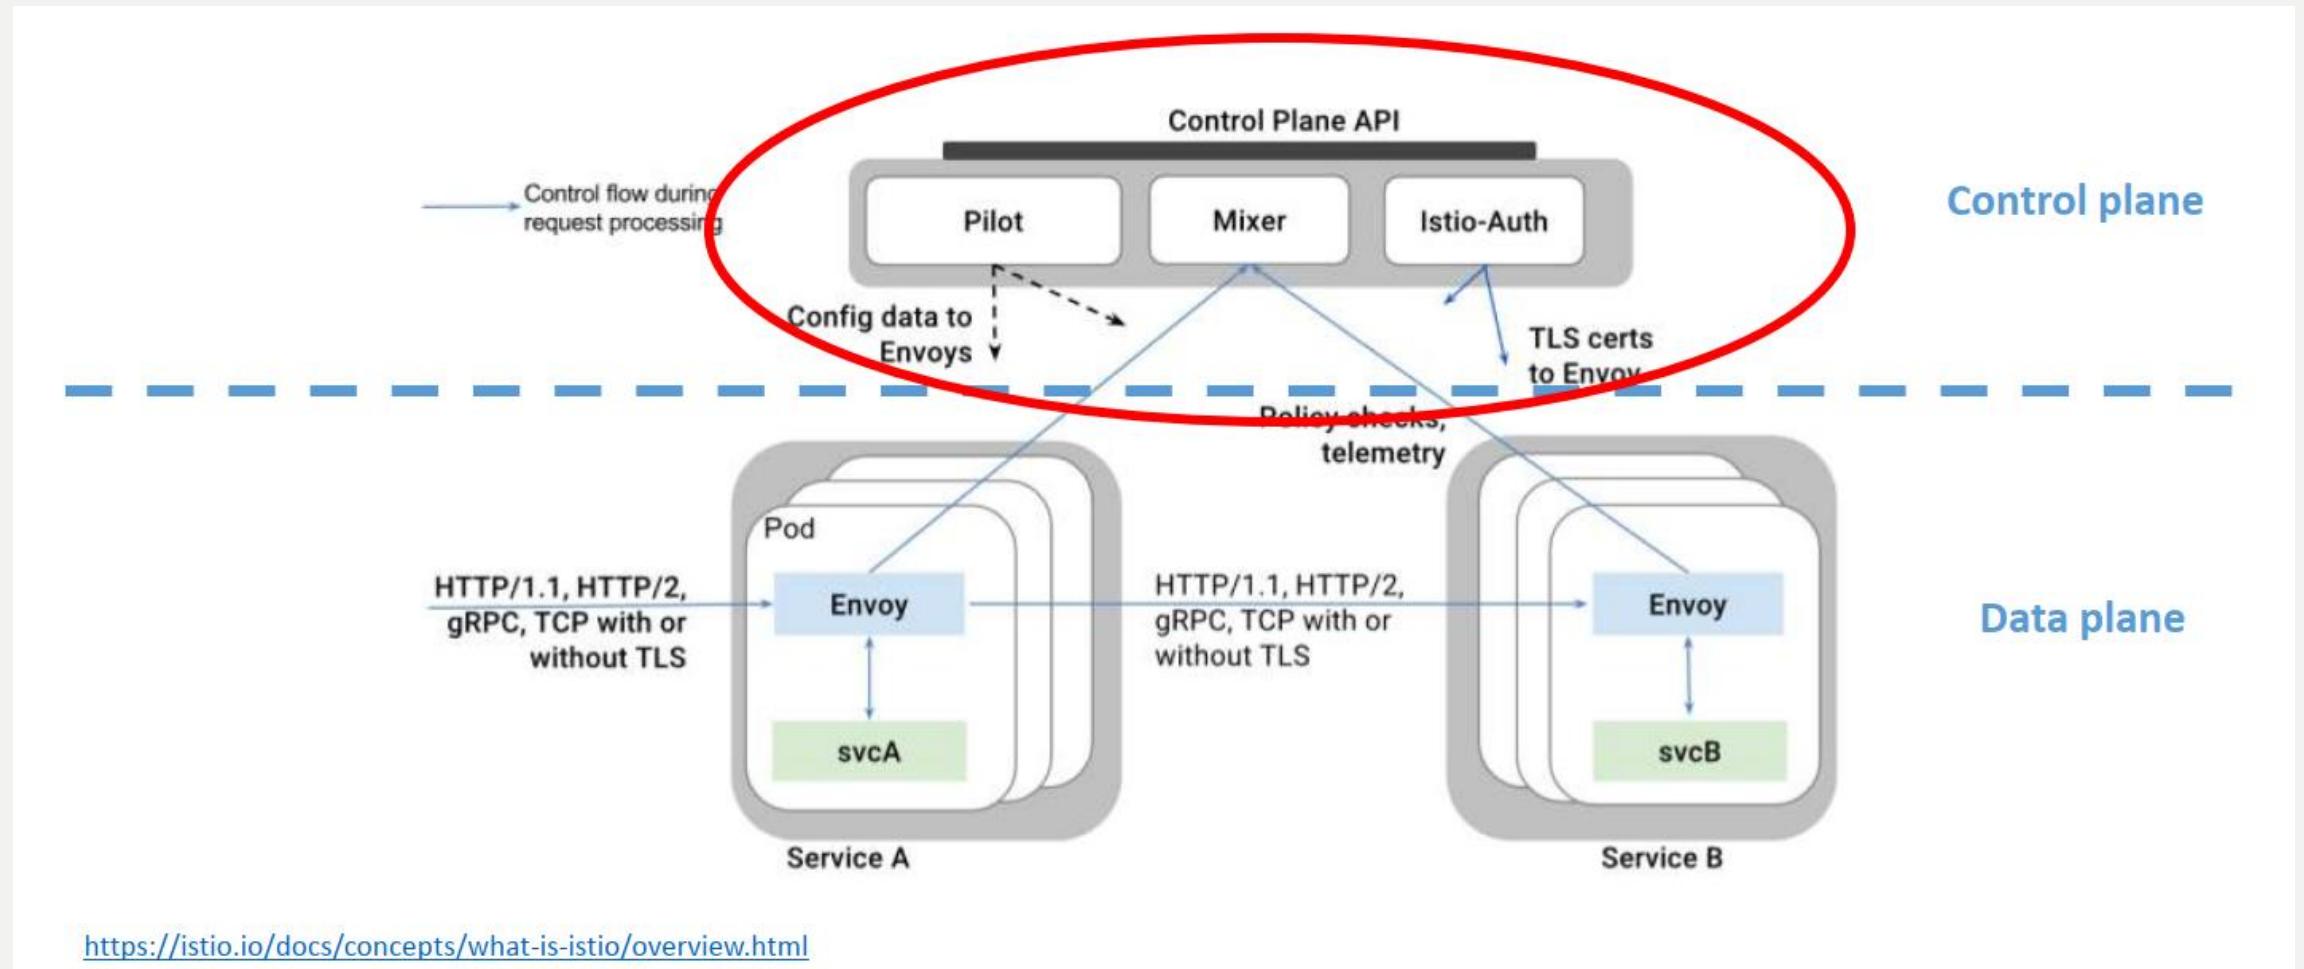
Citadel



Istio Components



Control – How they fit the control plane





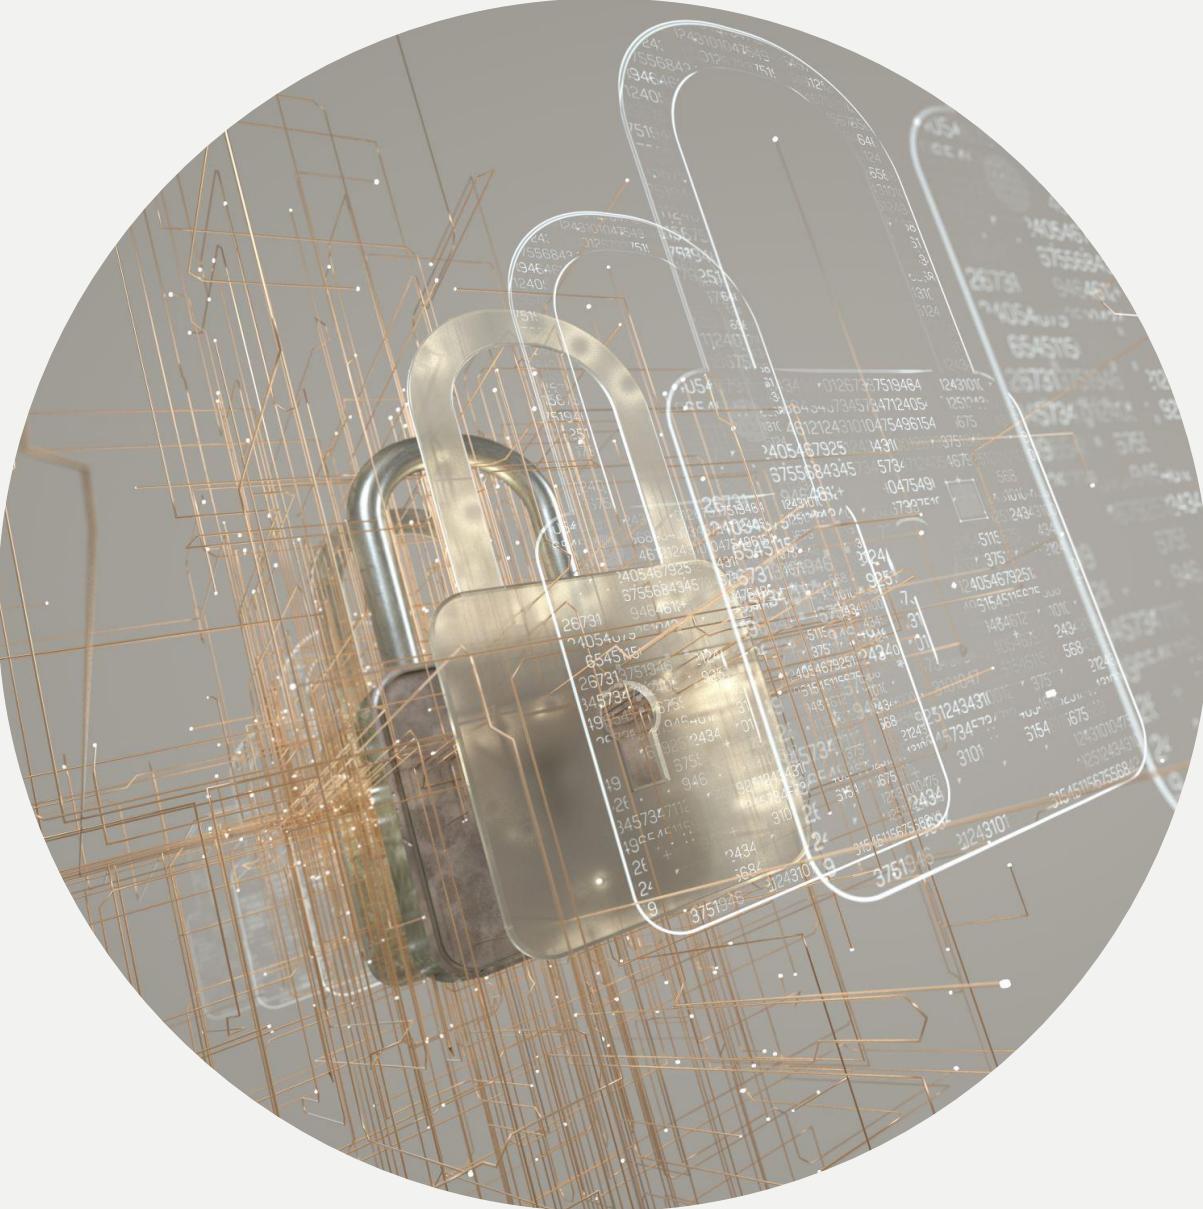
Pilot

- Pilot provides
 - Service Discovery
 - Fine-grained request distribution, retries, timeouts.



Mixer

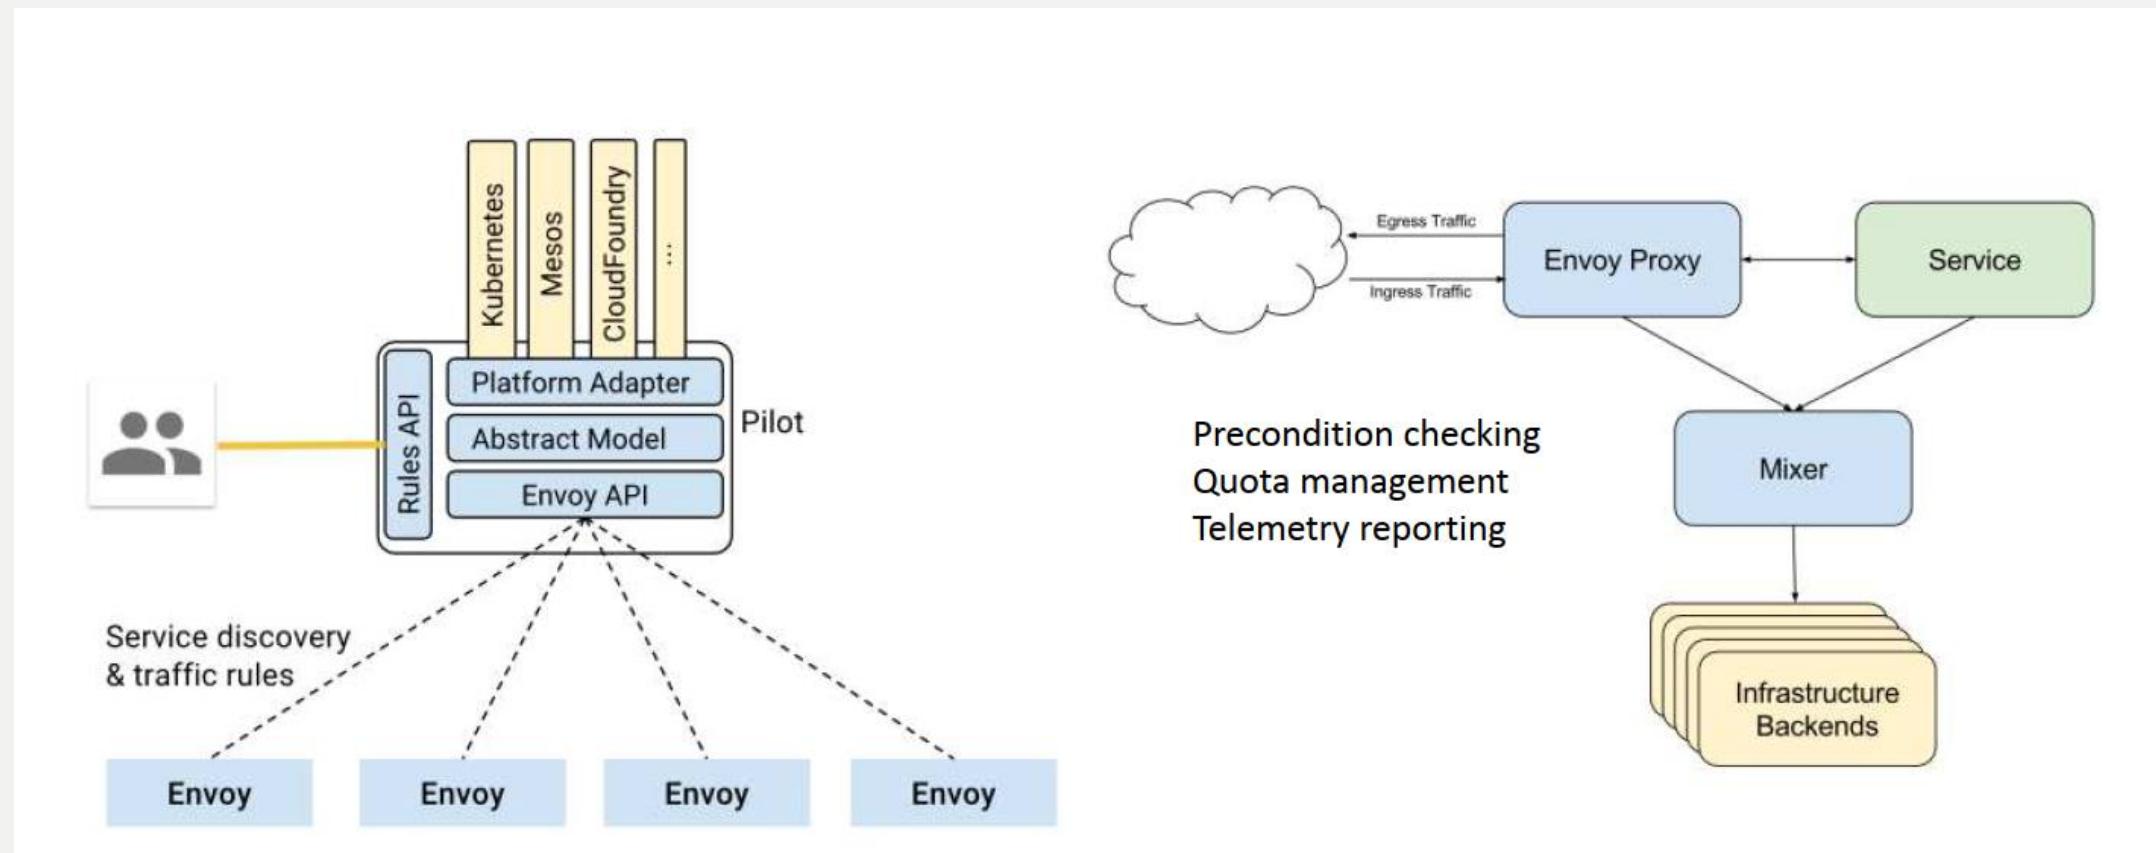
- Mixer allows you to:
 - Create policies and rules
 - Rate-Limiting rules
 - Custom metrics
 - Pluggable architecture



Citadel

- Responsible for secure operations
 - Certificate Signing
 - Certificate Issuance
 - mTLS security between services through transparent encryption

Pilot, Mixer and Citadel



What features does Service Mesh provide?

Traffic Management

Policy Enforcement

Platform Support

Observability

Security

Telemetry

Traffic Control with Istio - Canary and Dark Launches

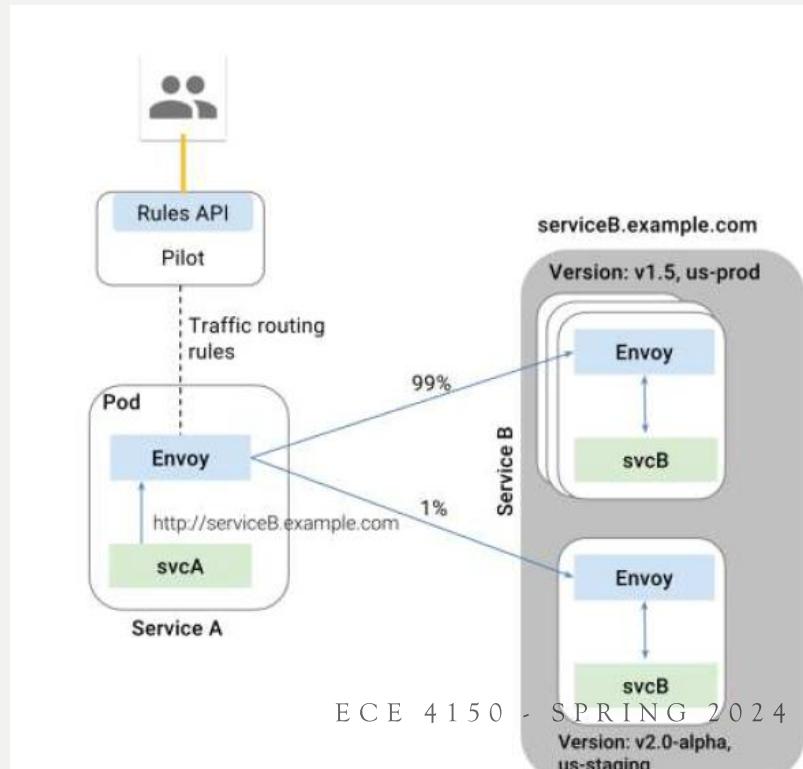
```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: recommendation
  namespace: tutorial
spec:
  host: recommendation
  subsets:
    - labels:
        version: v1
      name: version-v1
    - labels:
        version: v2
      name: version-v2
```

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
    http:
      - route:
          - destination:
              host: recommendation
              subset: version-v1
              weight: 100
```

Original Launch

Traffic Control: Canary Release for API v2

- Allows for phased releases with smaller amounts of load



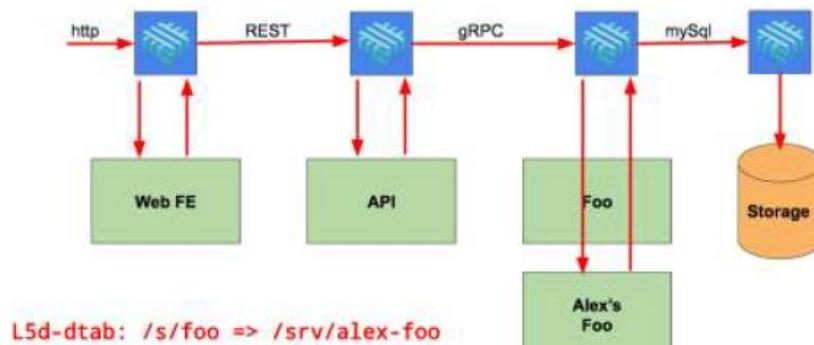
```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
  http:
    - route:
        - destination:
            host: recommendation
            subset: version-v1
            weight: 90
        - destination:
            host: recommendation
            subset: version-v2
            weight: 10
```

Traffic Control: Dark Launch

- Create duplicate or mirror traffic that duplicates traffic to a new version of your microservice to see how it behaves compared to the live application pod

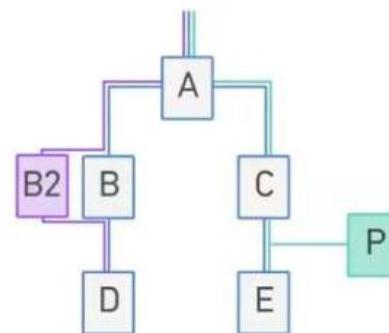
```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
    http:
      - route:
          - destination:
              host: recommendation
              subset: version-v1
          mirror:
            host: recommendation
            subset: version-v2
```

Traffic Routing: Shadow, Fault Inject, Debug



Per-request routing: staging

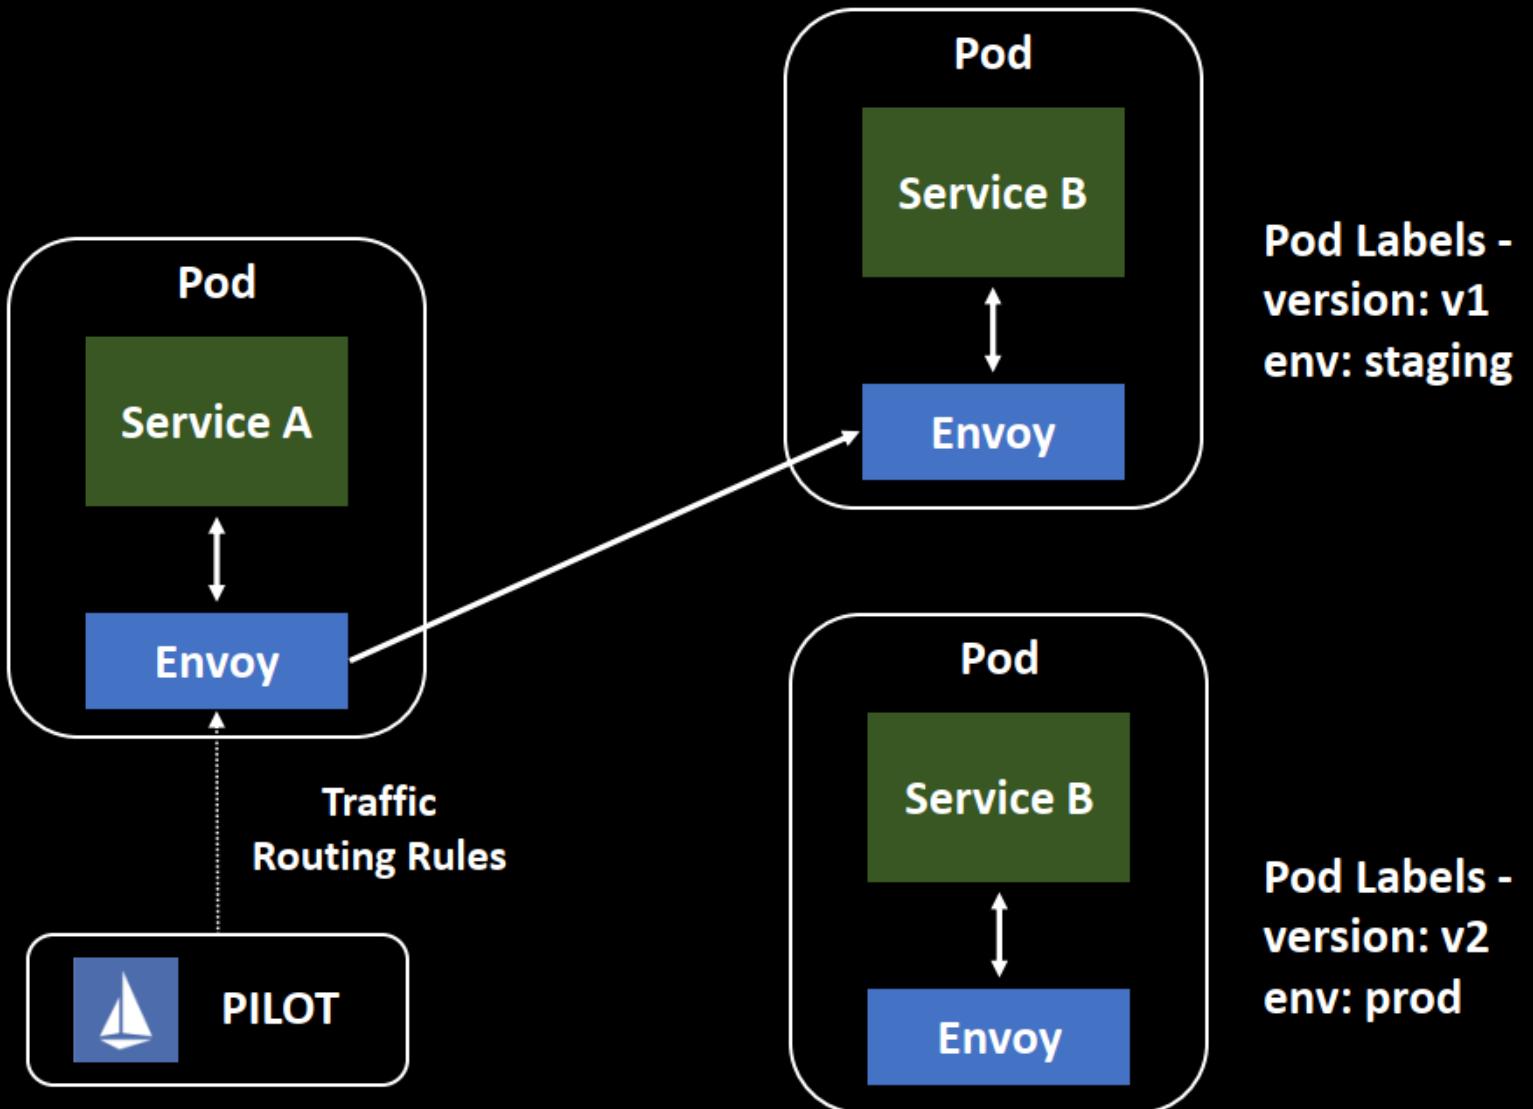
GET / HTTP/1.1
Host: mysite.com
l5d-dtab: /svc/B => /svc/B2



Traffic Routing

Route all traffic to v1 of ServiceB

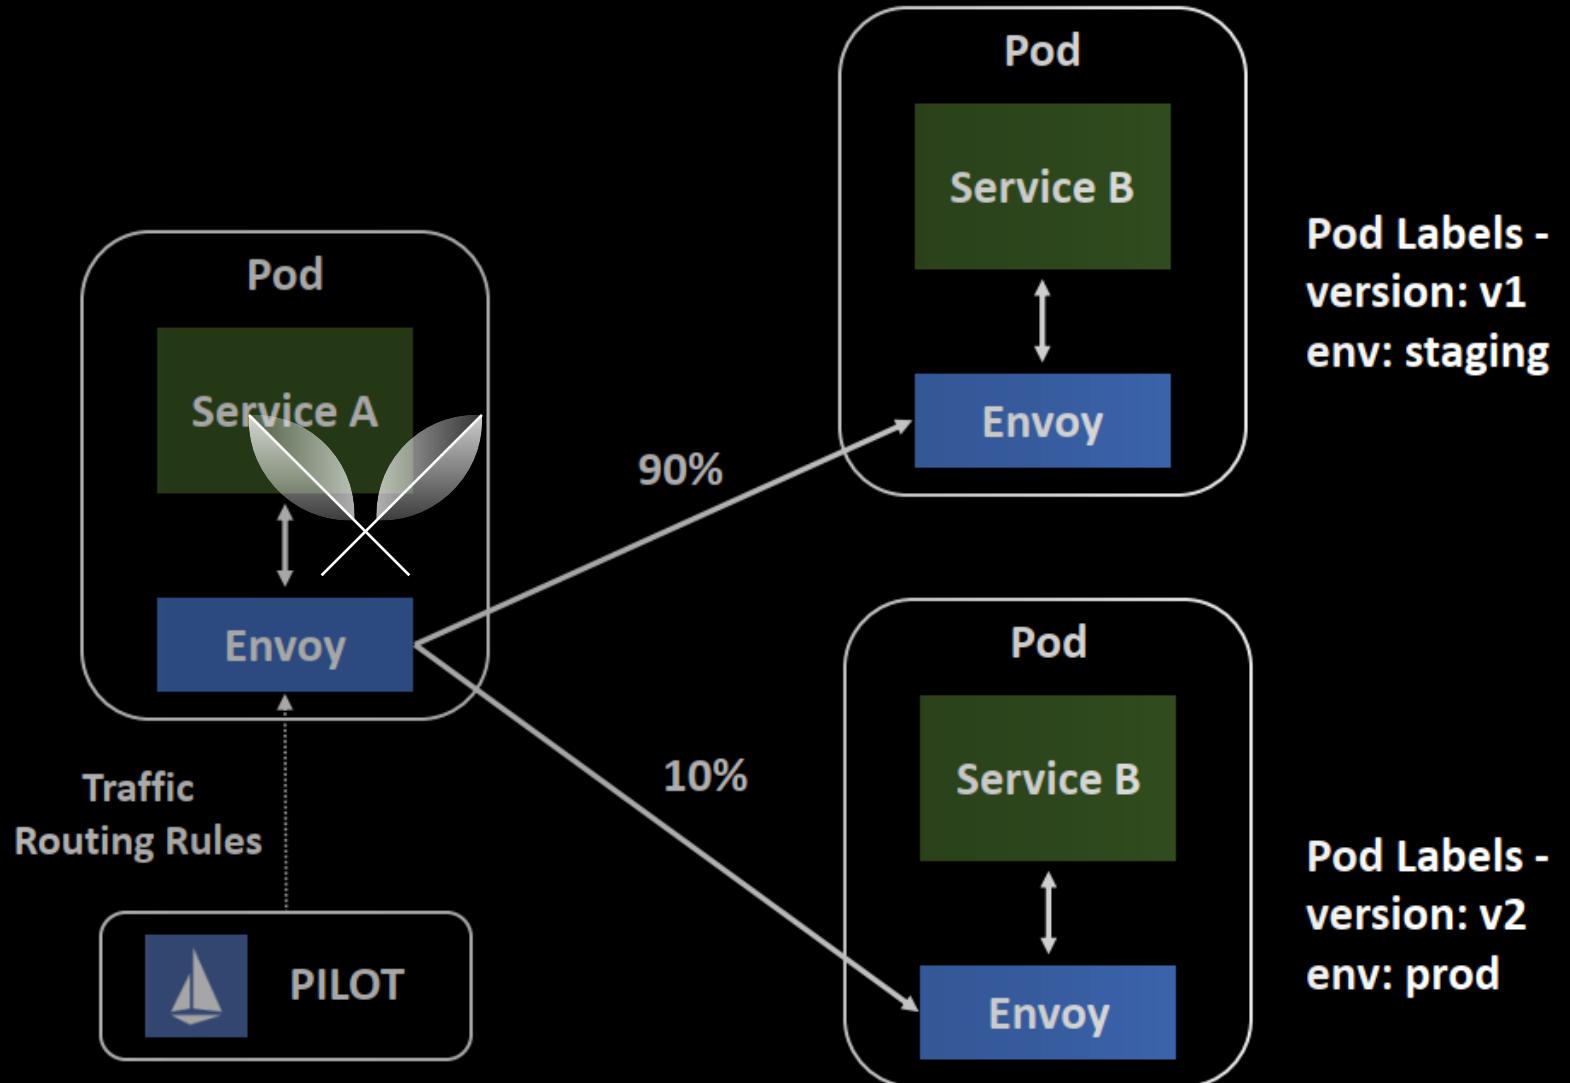
```
kind: virtualService
metadata:
  name: serviceB
spec:
  hosts:
  - serviceB
  http:
  - route:
    - destination:
        host: serviceB
        subset: v1
```



Canary Deployment

Percentage based Traffic Split

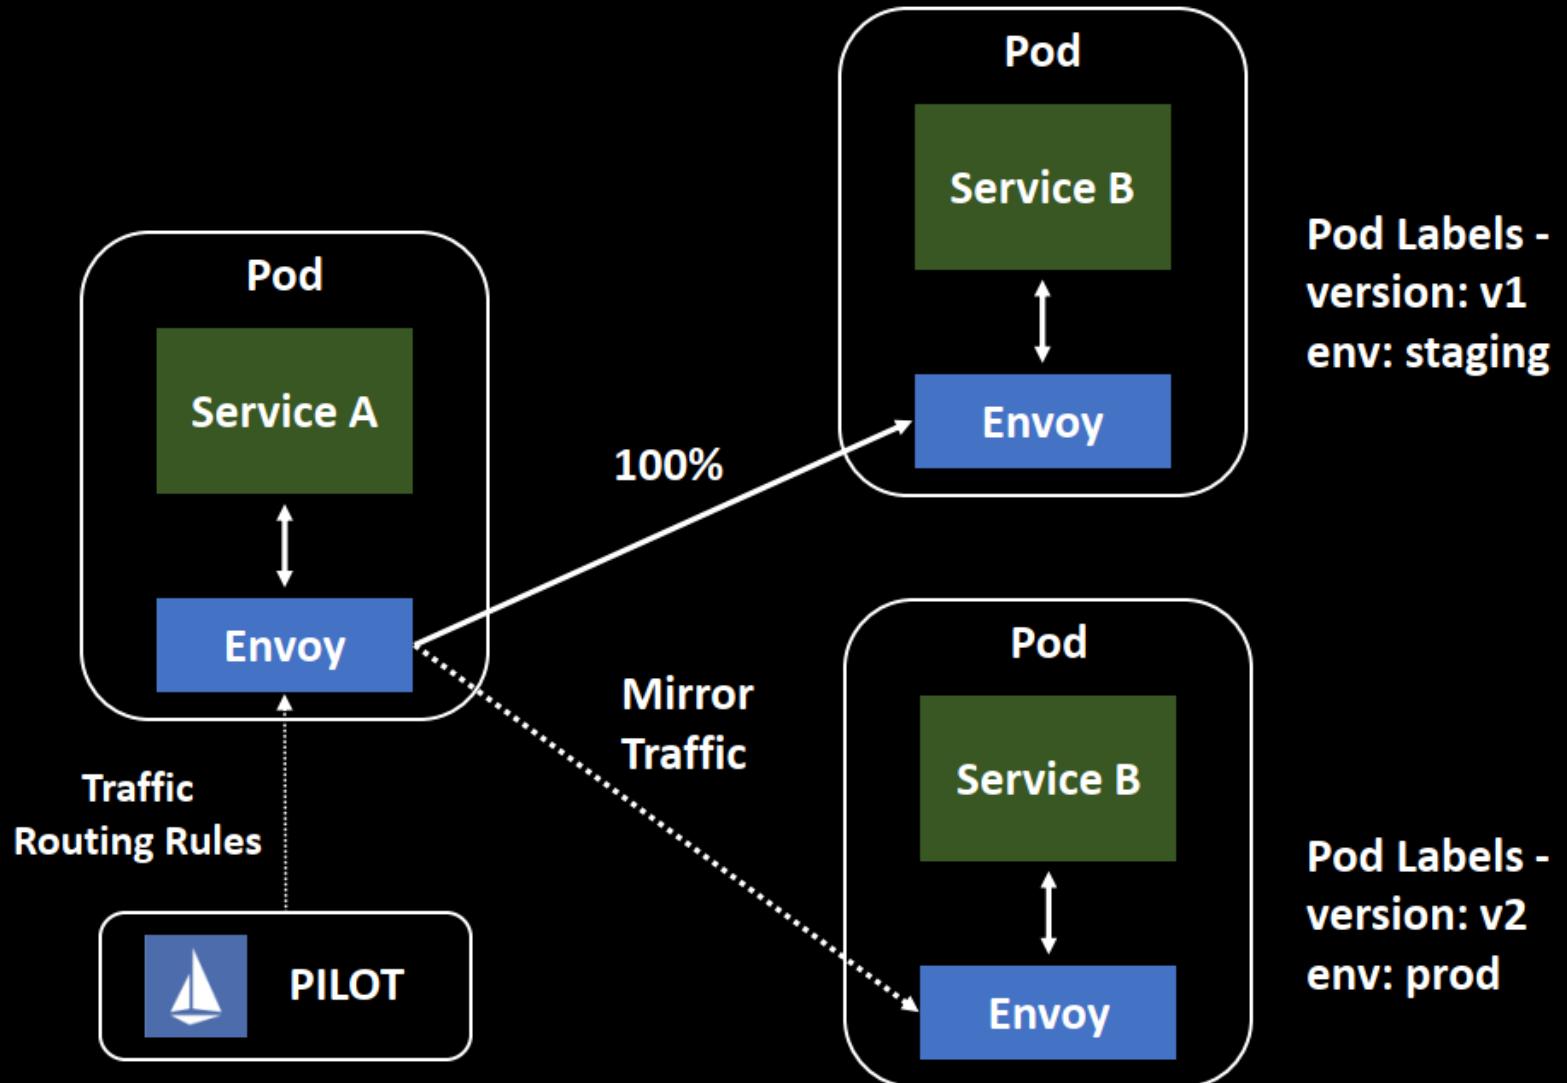
```
kind: virtualService
metadata:
  name: serviceB
spec:
  hosts:
    - serviceB
  http:
    - route:
        - destination:
            host: serviceB
            subset: v1
            weight: 90
        - destination:
            host: serviceB
            subset: v2
            weight: 10
```



Dark Launches

Traffic Mirroring

```
kind: virtualService
metadata:
  name: serviceB
spec:
  hosts:
    - serviceB
  http:
    - route:
        - destination:
            host: serviceB
            subset: v1
            weight: 100
        mirror:
            host: serviceB
            subset: v2
```



Service Resiliency with Istio: Load balancing Rules

- Options

Round Robin - Even distributes load across endpoints

Random - Evenly distributes load without any order

Least_Conn - Picks to two random hosts and determines which of the two has fewer outstanding requests to that endpoint

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: recommendation
  namespace: tutorial
spec:
  host: recommendation
  trafficPolicy:
    loadBalancer:
      simple: RANDOM
```

Service Resiliency: Timeout

- Calls to reduce unpredictable behavior.
- Did the service fail?

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  creationTimestamp: null
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
  http:
    - route:
        - destination:
            host: recommendation
            timeout: 1.000s
```

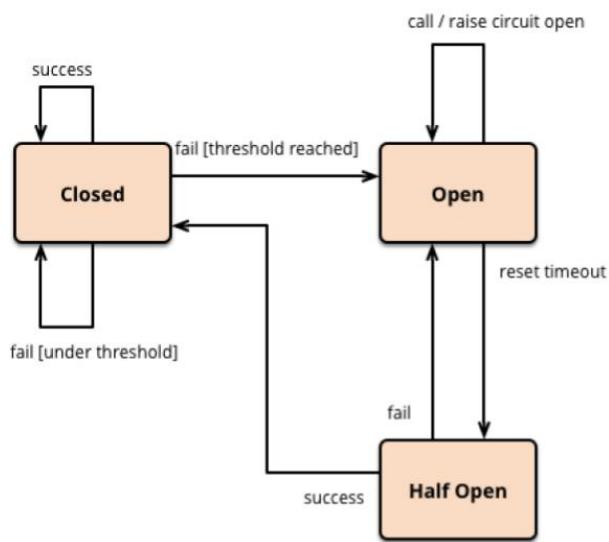
Service Resiliency: Retry

Retry capabilities allow more
resilient services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
    http:
      - route:
          - destination:
              host: recommendation
            retries:
              attempts: 3
              perTryTimeout: 2s
```

Circuit Breaker Rule

– Service Resiliency



```
1 destination: serviceB.example.cluster.local policy:  
2 - tags:  
3   version: v1  
4   circuitBreaker:  
5     simpleCb:  
6       maxConnections: 100  
7       httpMaxRequests: 1000  
8       httpMaxRequestsPerConnection: 10  
9       httpConsecutiveErrors: 7  
10      sleepWindow: 15m  
11      httpDetectionInterval: 5m
```

```
1 - protocol: http  
2   client:  
3     failureAccrual:  
4       kind: io.l5d.successRate  
5       successRate: 0.9  
6       requests: 20  
7       backoff:  
8         kind: constant  
9         ms: 10000
```

Service Resiliency: Multiple Concurrent Requests

- Circuit Breaker Rule

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  creationTimestamp: null
  name: recommendation
  namespace: tutorial
spec:
  host: recommendation
  subsets:
    - name: version-v1
      labels:
        version: v1
    - name: version-v2
      labels:
        version: v2
      trafficPolicy:
        connectionPool:
          http:
            http1MaxPendingRequests: 1
            maxRequestsPerConnection: 1
          tcp:
            maxConnections: 1
        outlierDetection:
          baseEjectionTime: 120.000s
          consecutiveErrors: 1
          interval: 1.000s
          maxEjectionPercent: 100
```

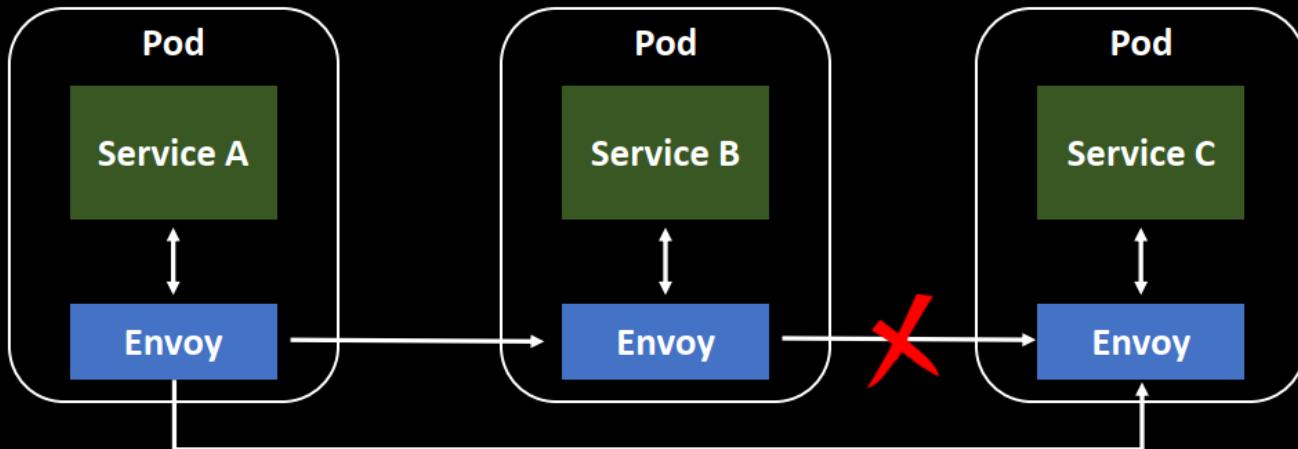
Service Resiliency: Introduce delays to stress test

- Introduces 7 seconds of delay into 50% of responses from the recommendation

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  creationTimestamp: null
  name: recommendation
  namespace: tutorial
spec:
  hosts:
    - recommendation
  http:
    - fault:
        delay:
          fixedDelay: 7.000s
          percent: 50
      route:
        - destination:
            host: recommendation
            subset: app-recommendation
```

Circuit Breaker

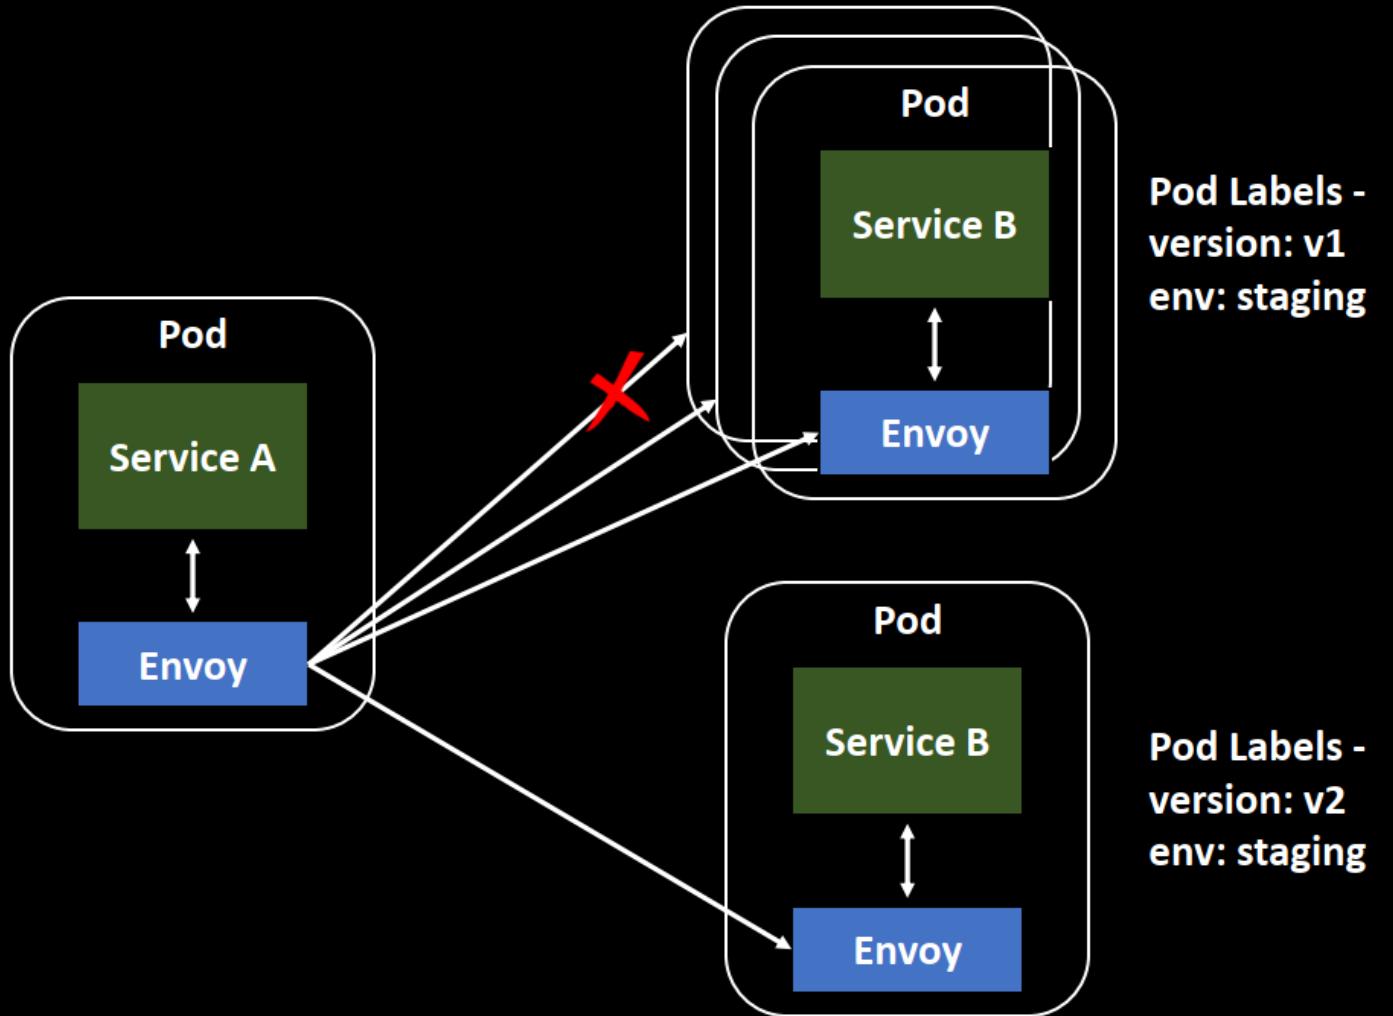
```
# Limits the number of concurrent  
connections and requests  
  
kind: DestinationRule  
metadata:  
  name: serviceC  
spec:  
  hosts:  
  - serviceC  
  trafficPolicy:  
    connectionPool:  
      http:  
        http1MaxPendingRequests: 10  
        maxRequestsPerConnection: 1  
      tcp:  
        maxConnections: 1
```



Outlier Detection

```
# Detect faulty instances in the
pool & remove from traffic routing

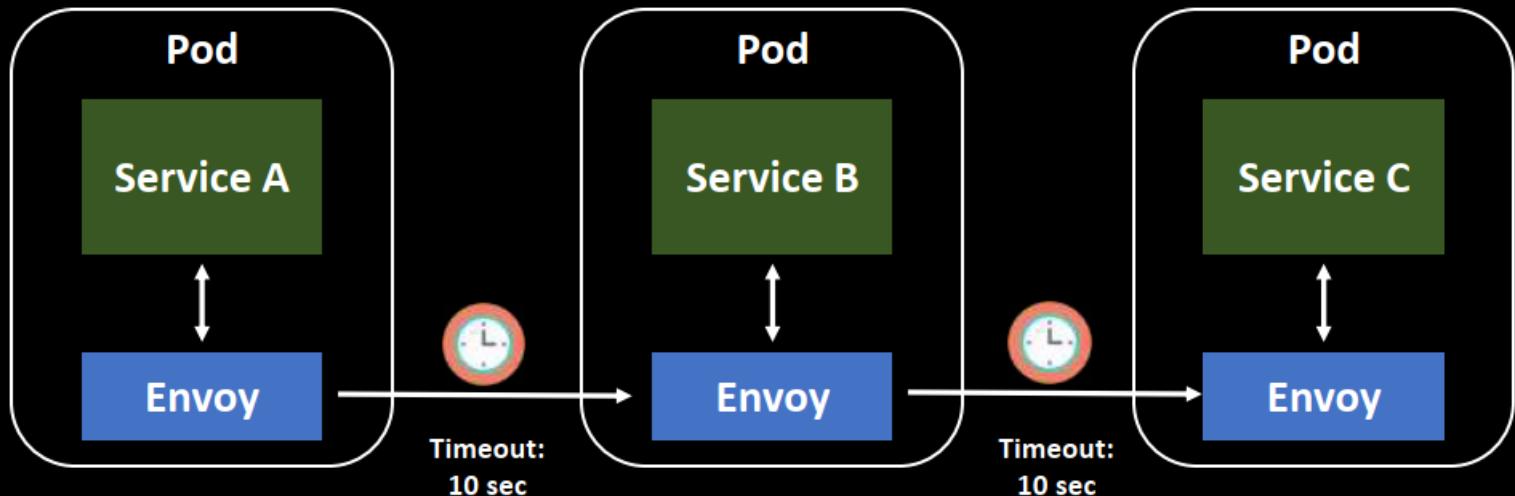
kind: DestinationRule
metadata:
  name: serviceB
spec:
  hosts:
  - serviceB
  trafficPolicy:
    outlierDetection:
      baseEjectionTime: 20s
      consecutiveErrors: 3
      interval: 10s
      maxEjectionPercent: 100
```



Timeout

Timeout strategy for service communication over network

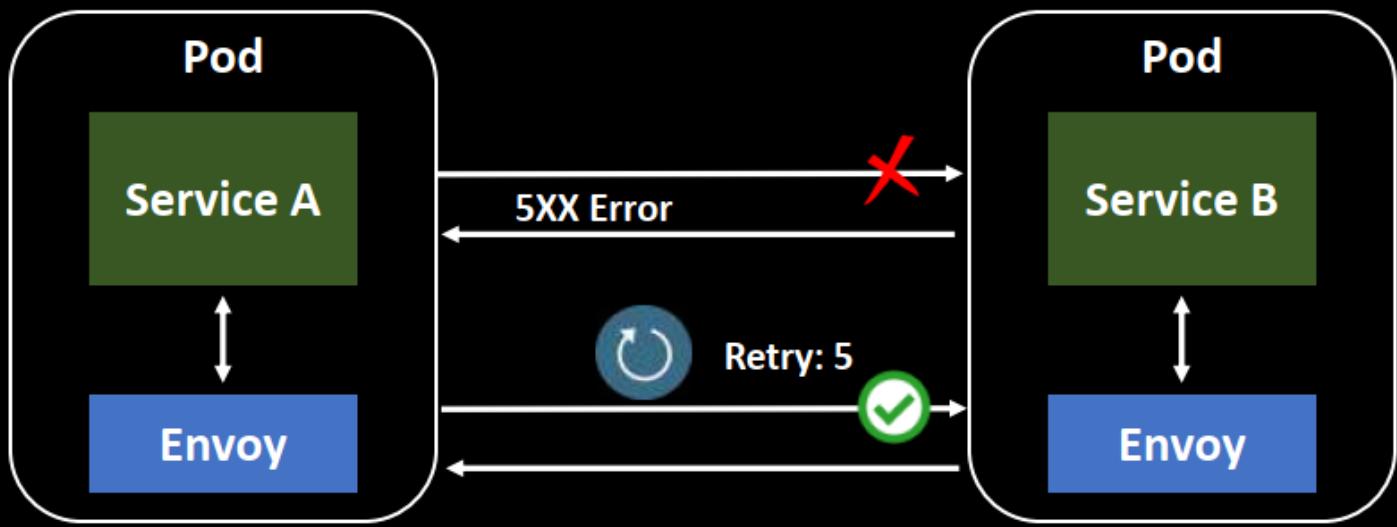
```
kind: virtualService
metadata:
  name: serviceB
spec:
  hosts:
  - serviceB
  http:
  - route:
    - destination:
        host: serviceB
    timeout: 10s
```



Istio Retry Policy

```
# Retry strategy for service communication over network
```

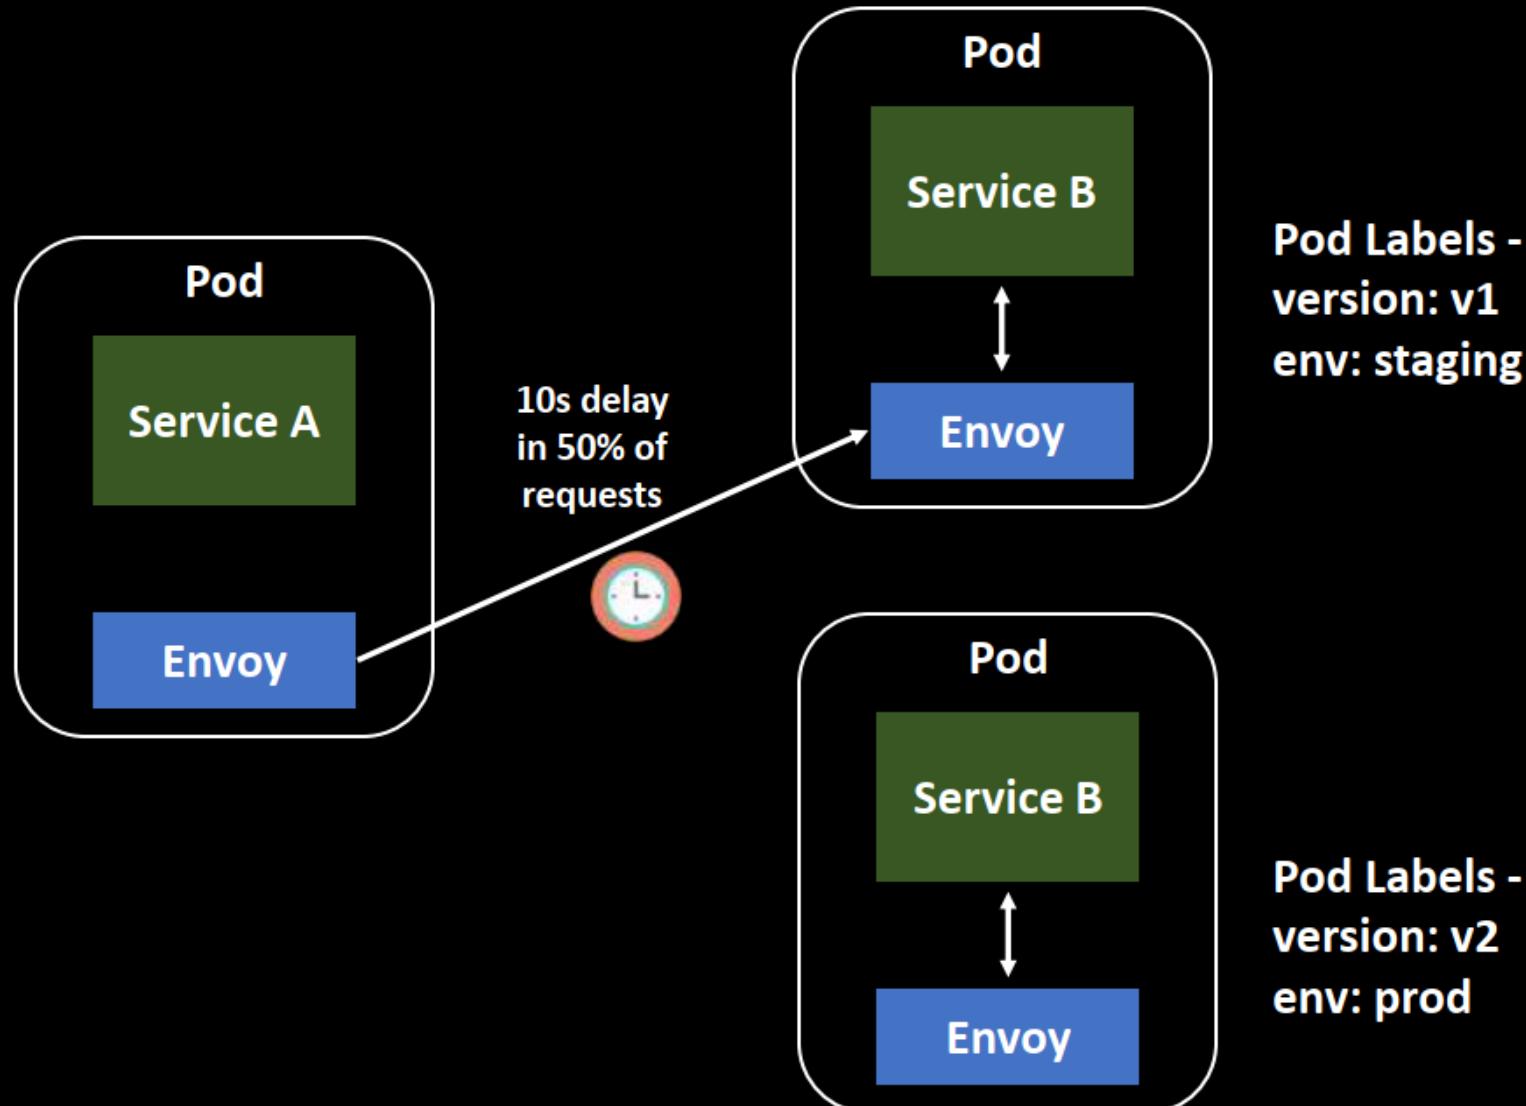
```
kind: VirtualService
metadata:
  name: serviceB
spec:
  hosts:
    - serviceB
  http:
    - route:
        - destination:
            host: serviceB
  retries:
    attempts: 3
    perTryTimeout: 2s
```



Chaos Testing – Inject Delays

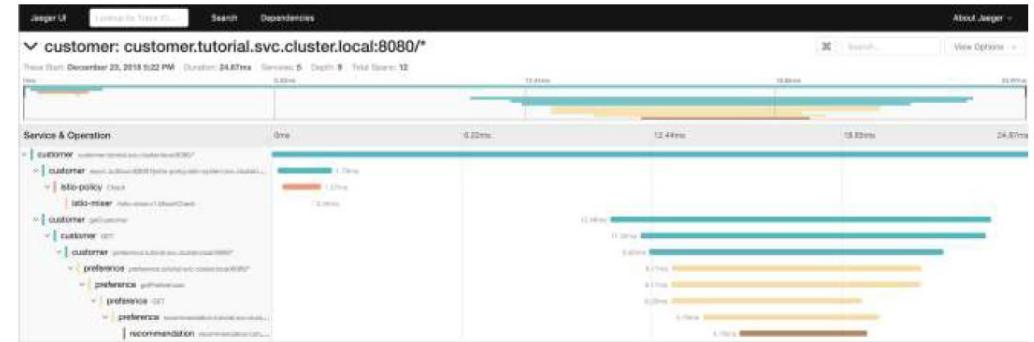
```
# Create rule to delay traffic to ServiceB v1
```

```
kind: virtualService
metadata:
  name: serviceB
spec:
  hosts:
    - serviceB
  http:
    - fault:
        delay:
          fixedDelay: 10s
          percent: 50
      route:
        - destination:
            host: serviceB
            subset: v1
```



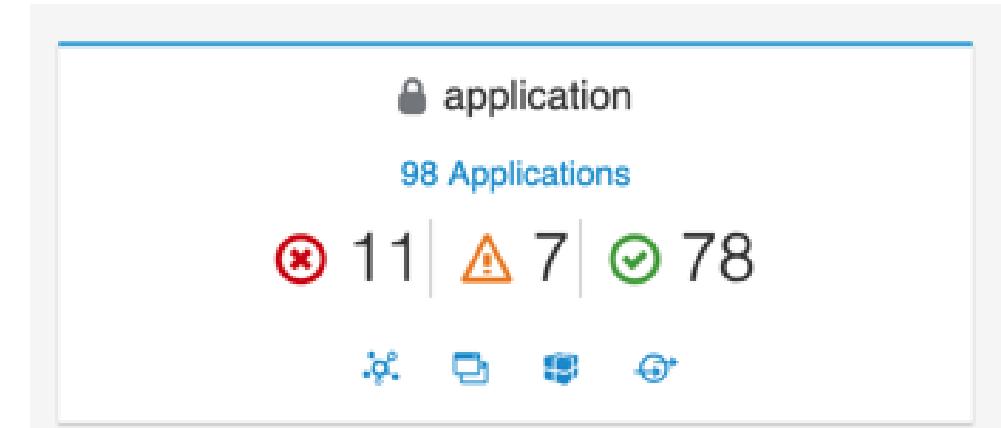
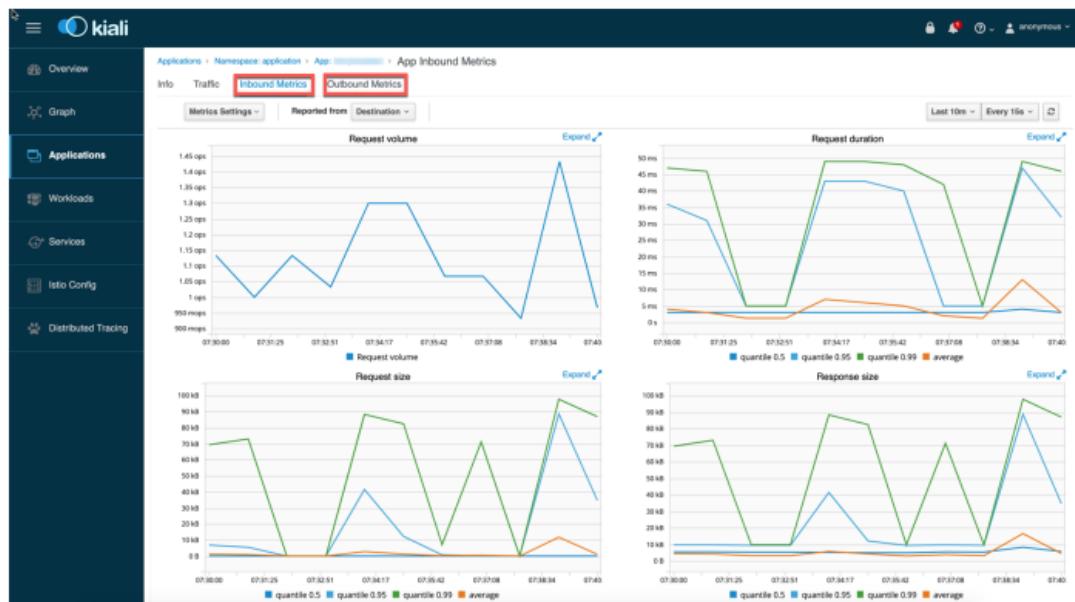
Observability - Tracing and Metrics

- Jaeger (Outsourced by Uber)

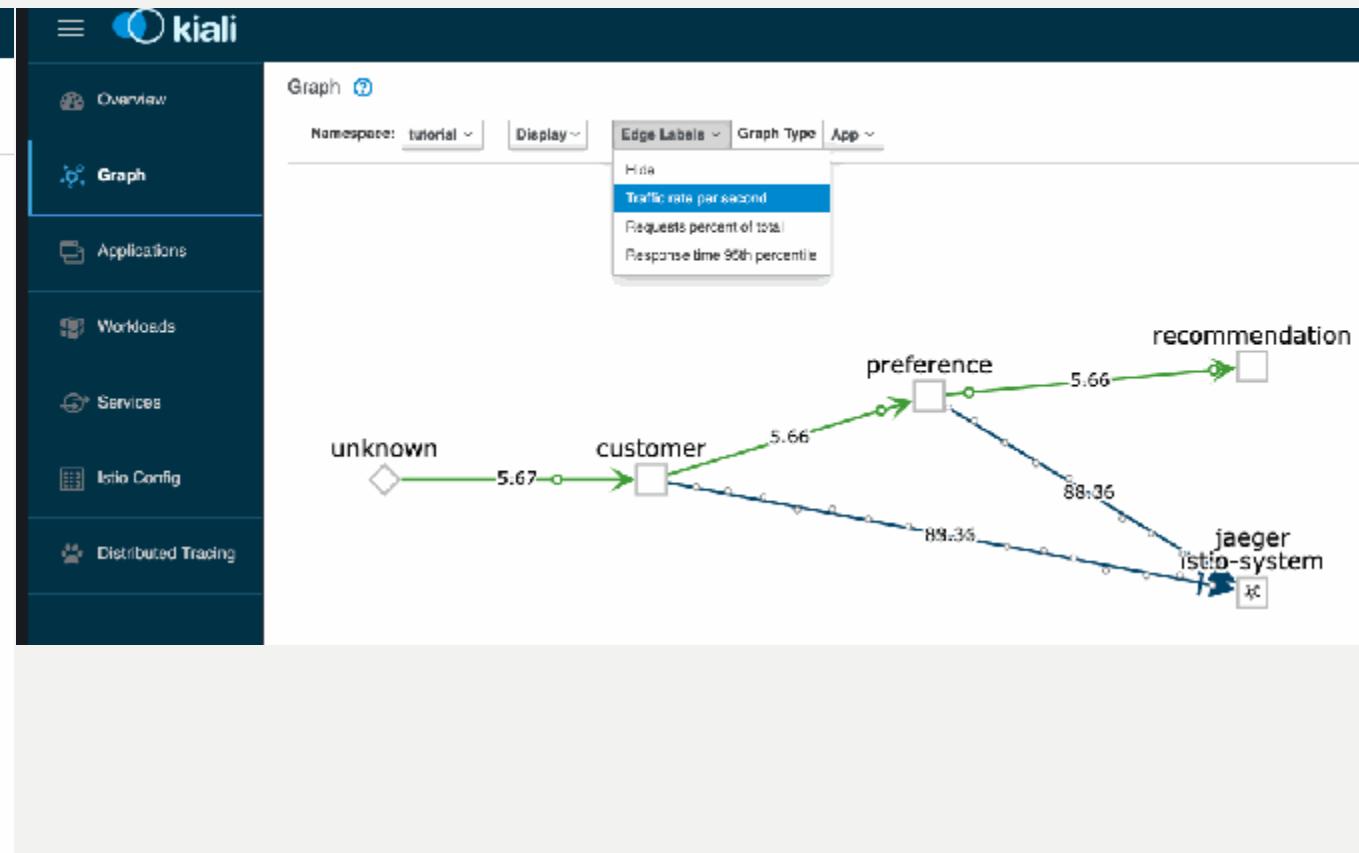
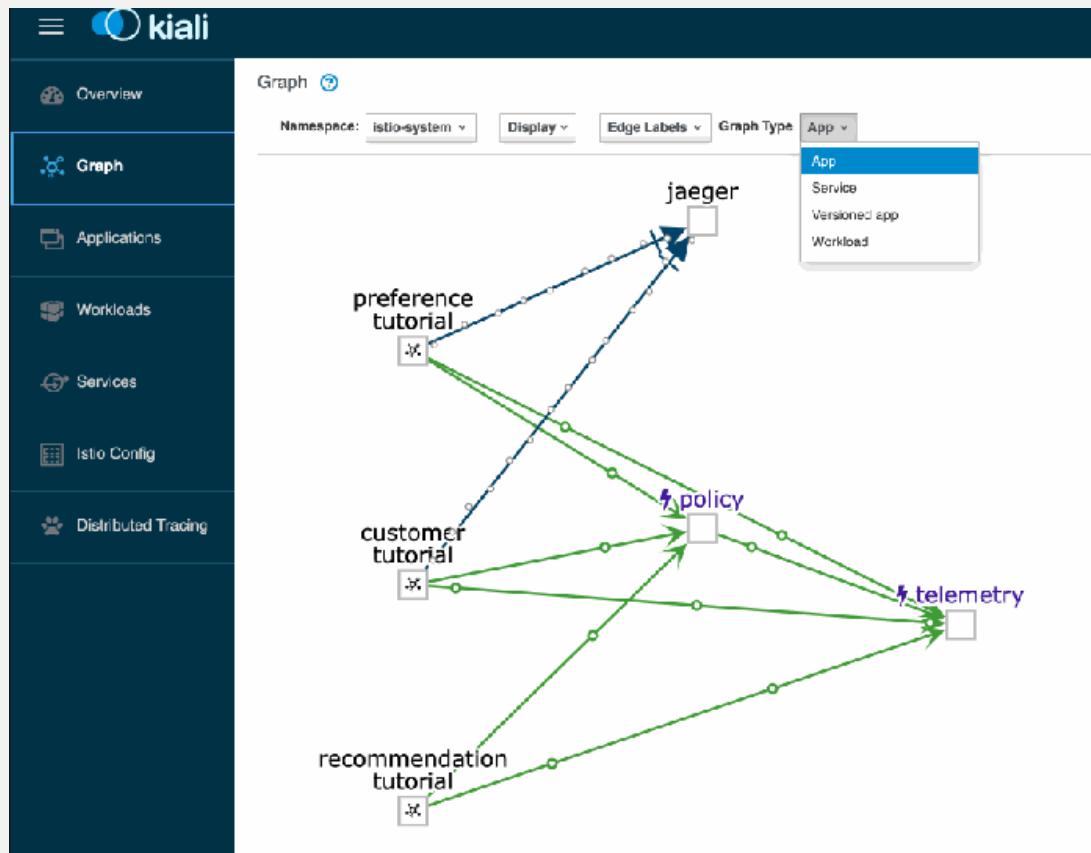


Observability: Service Graphs

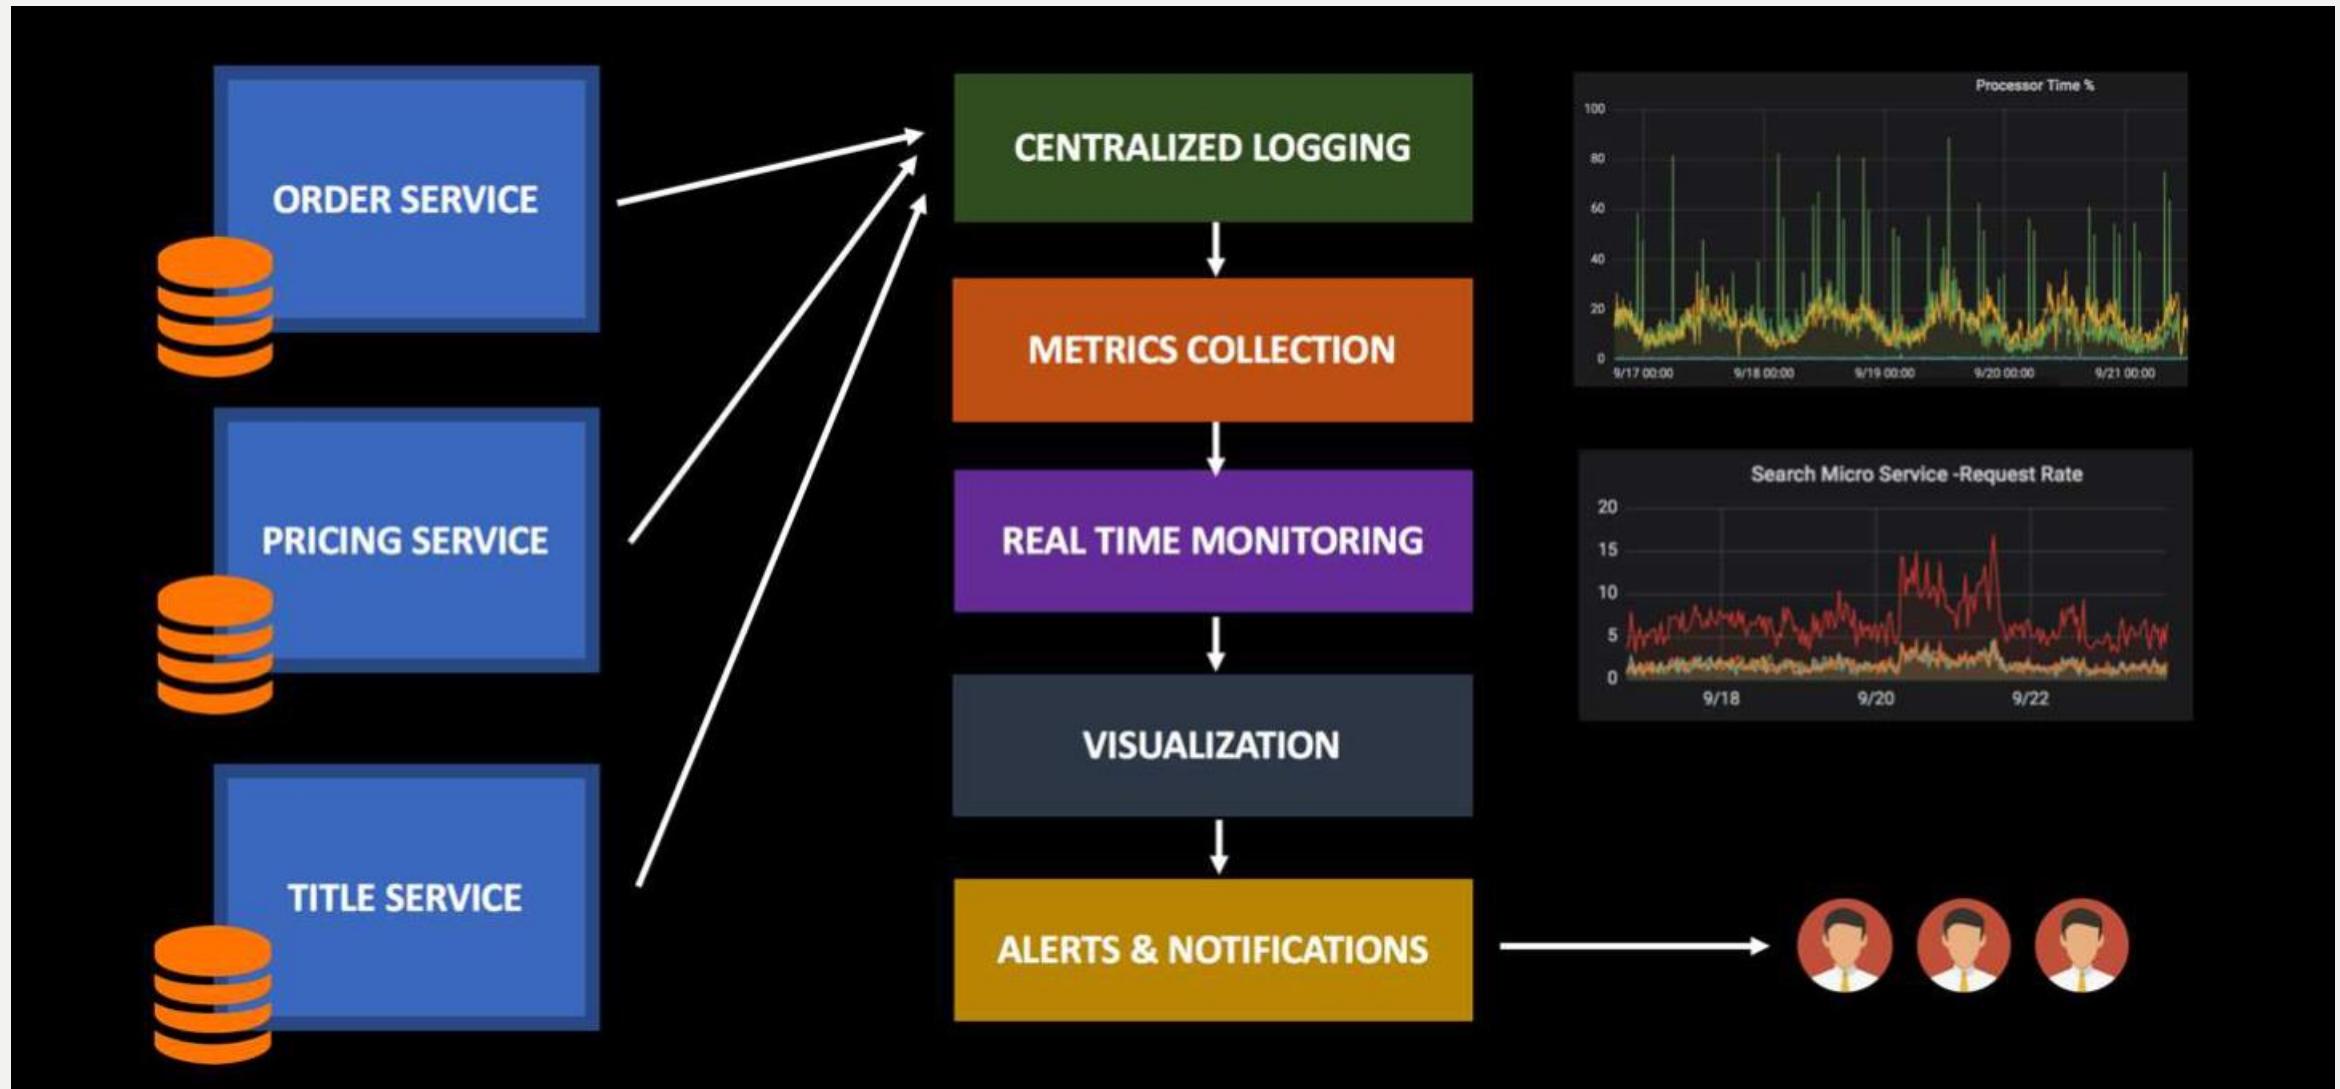
Visualization



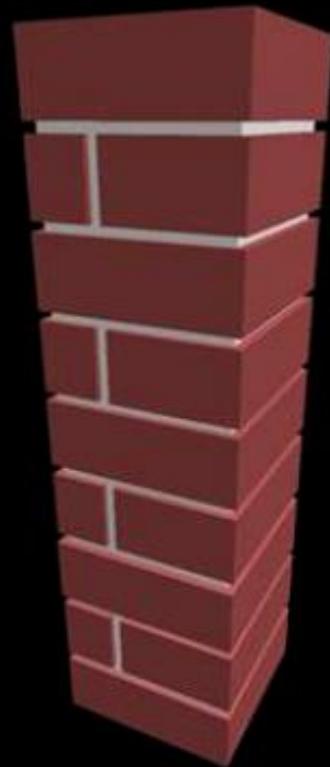
Visualization - Service Graphs



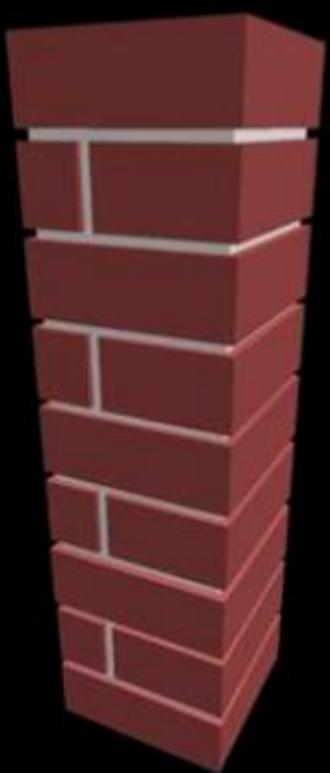
<https://dzone.com/articles/visualizing-the-istio-service-mesh-using-kiali>



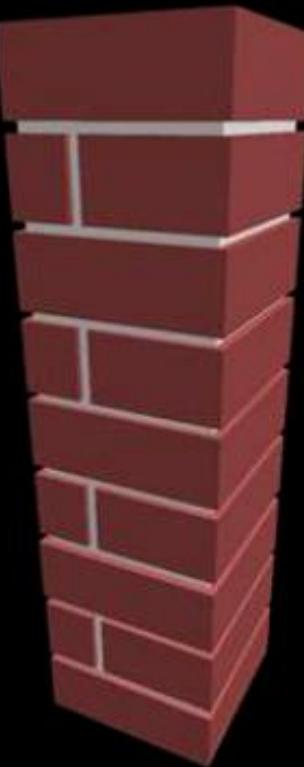
LOGGING

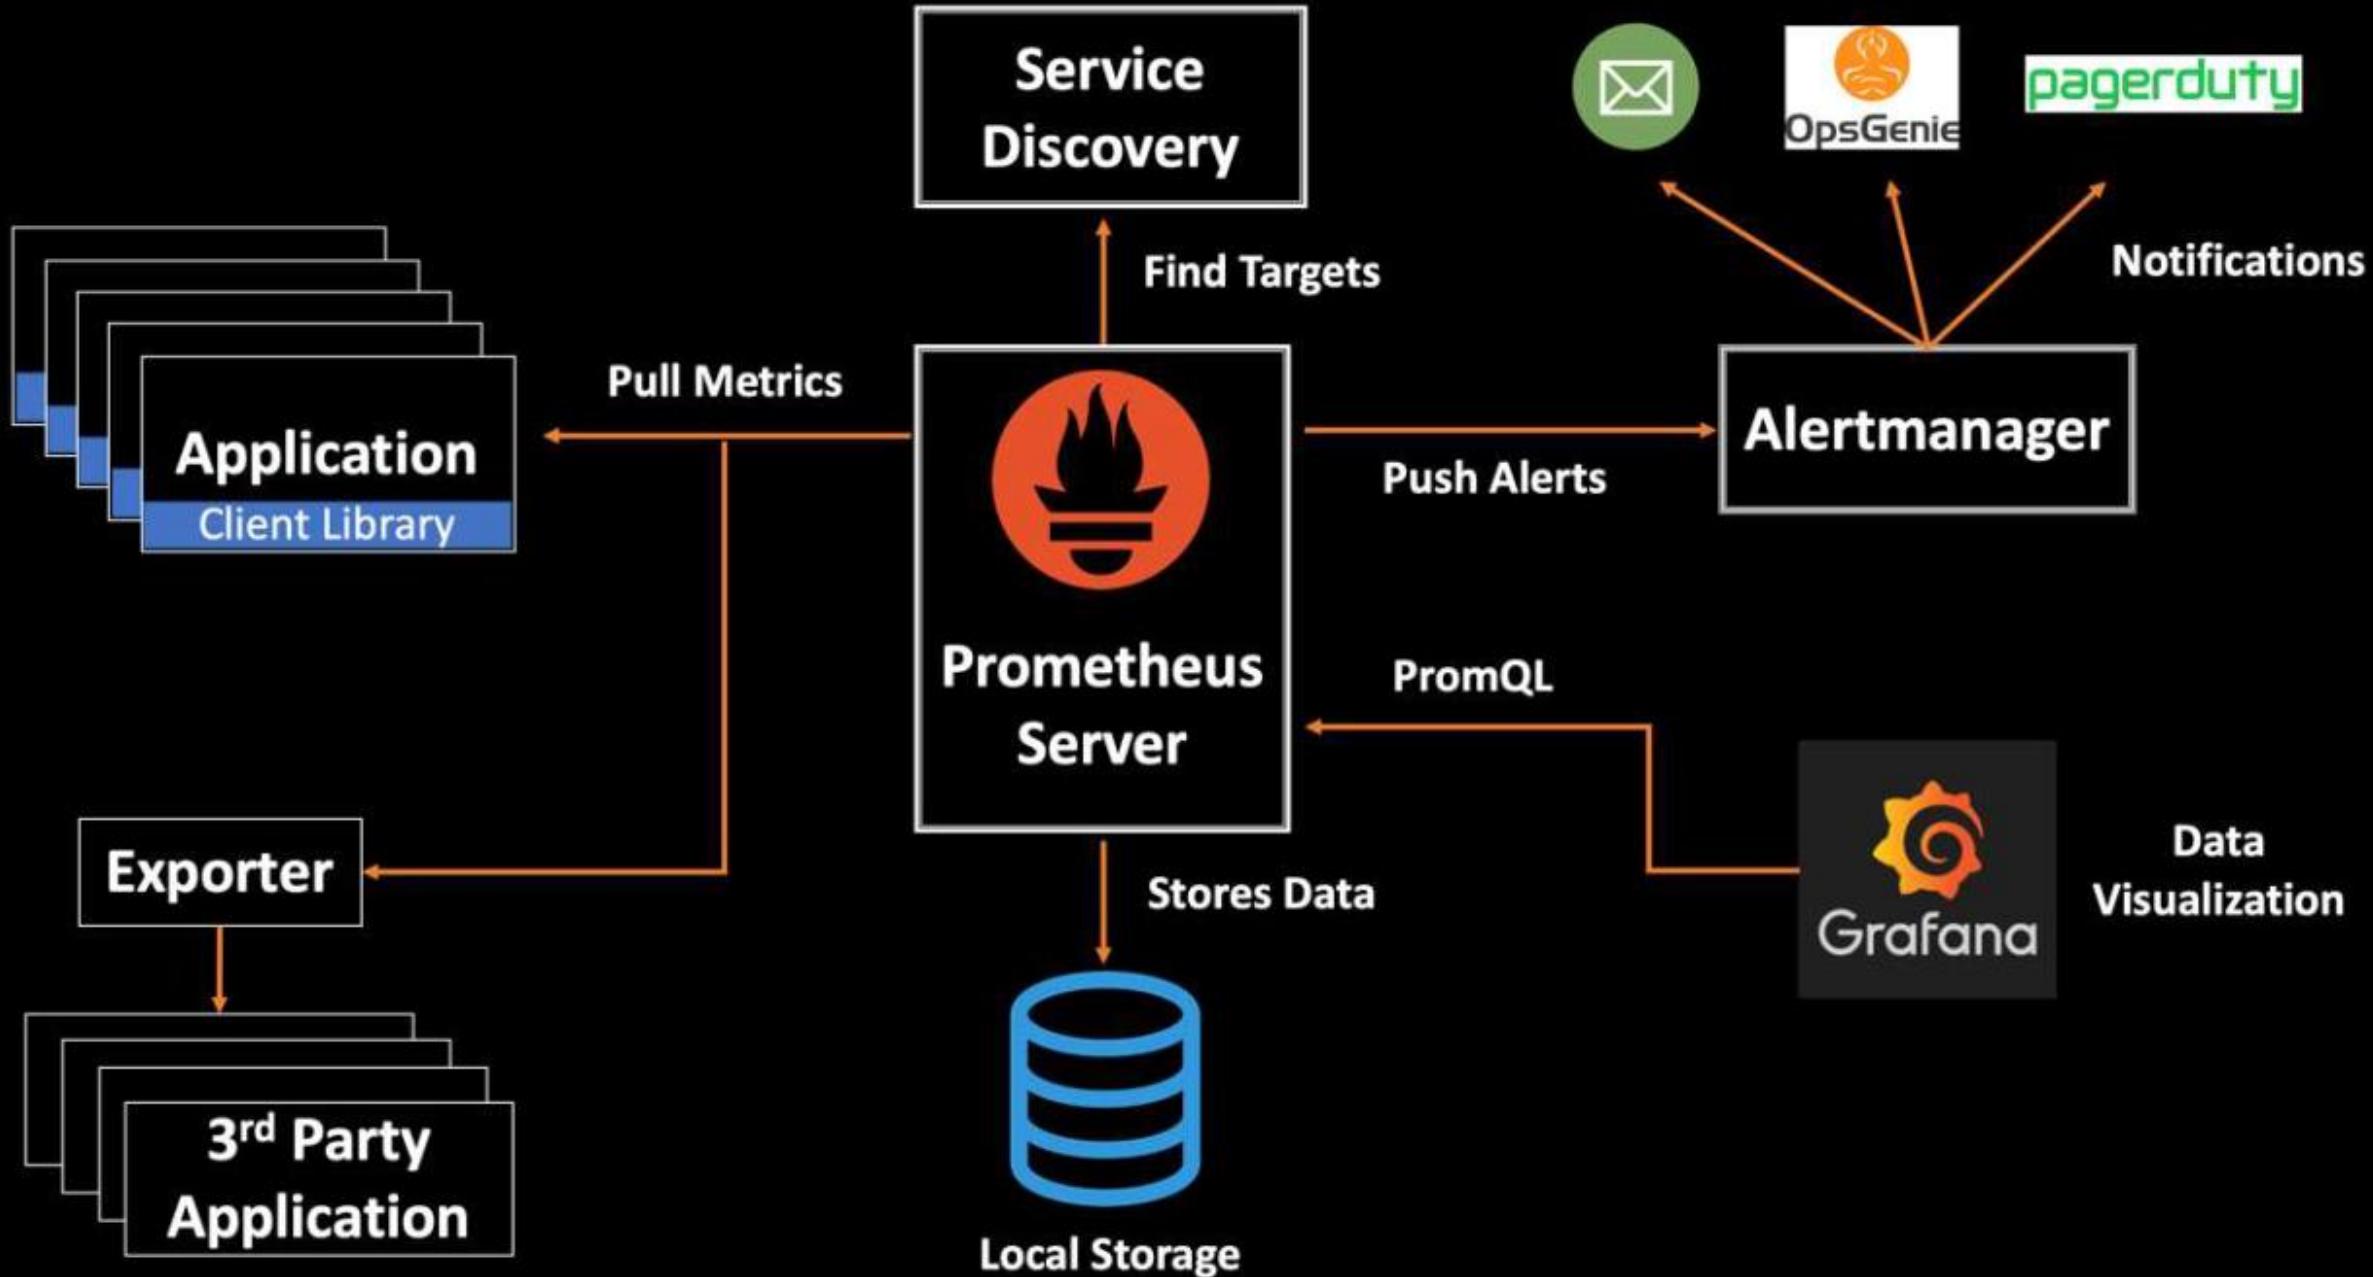


METRICS

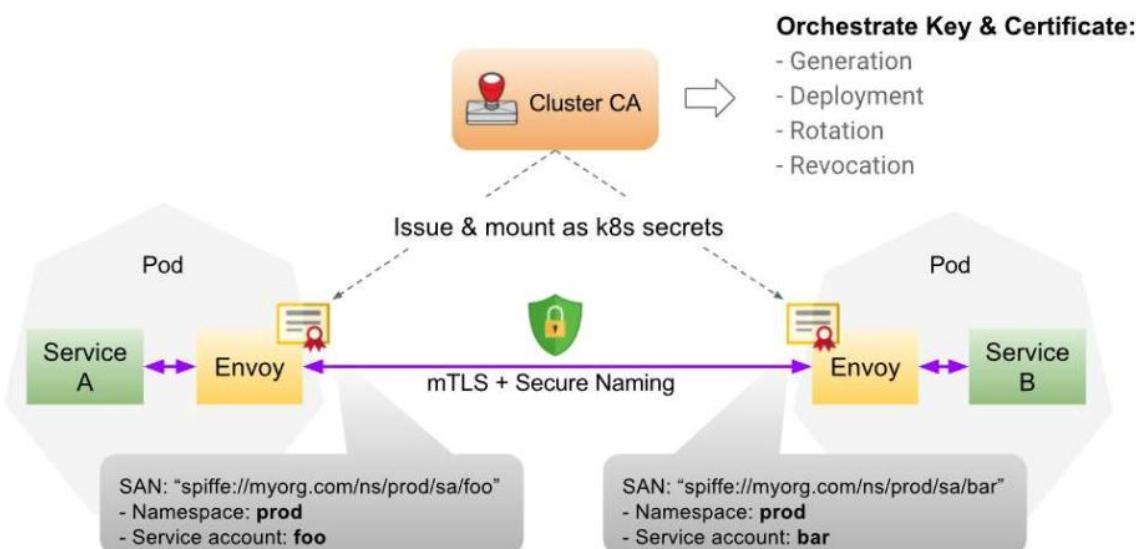


TRACING





Security – Enforce mTLS



Security - Enforce mTLS

```
apiVersion: "authentication.istio.io/v1alpha1"
kind: "Policy"
metadata:
  name: "default"
  namespace: "tutorial"
spec:
  peers:
  - mtls: {}
```

```
apiVersion: "networking.istio.io/v1alpha3"
kind: "DestinationRule"
metadata:
  name: "default"
  namespace: "tutorial"
spec:
  host: "*.{tutorial}.svc.cluster.local"
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL
```

Security – Access Control with Mixer Policies

```
metadata:  
  name: direct-cnn-through-egress-gateway  
  
spec:  
  hosts:  
    - edition.cnn.com  
  gateways:  
    - istio-egressgateway  
    - mesh  
  http:  
    - match:  
        - gateways:  
            - mesh  
        port: 80  
  route:  
    - destination:  
        host: istio-egressgateway.istio-system.svc.cluster.local  
        subset: cnn
```

<https://istio.io/latest/blog/2018/egress-monitoring-access-control/>

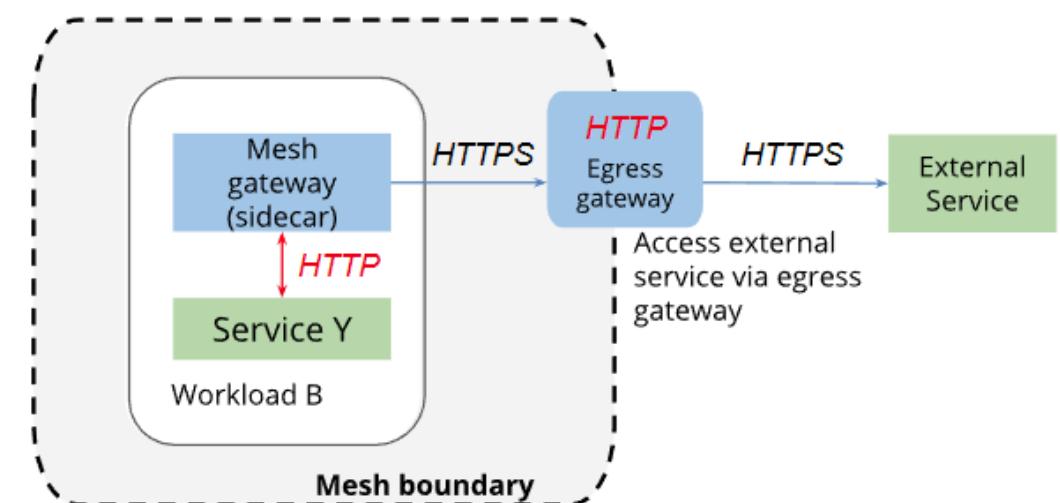
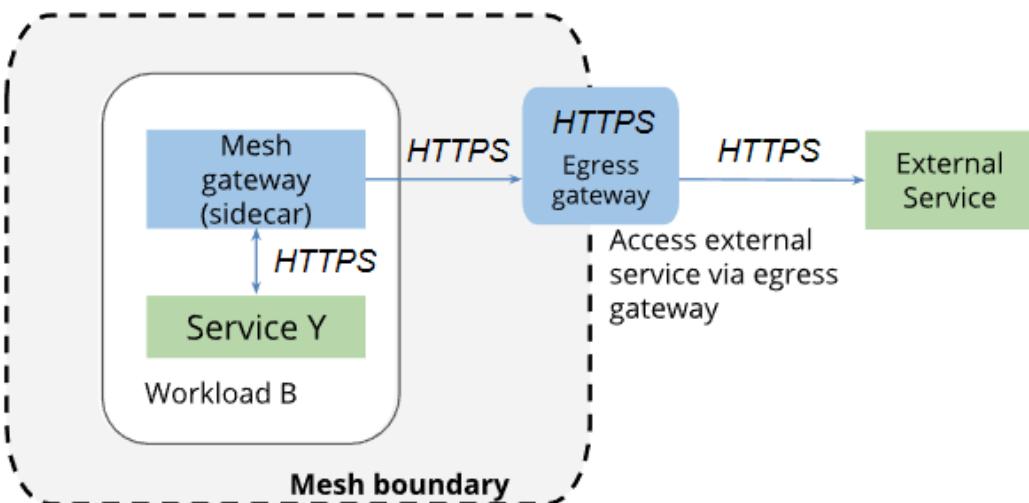


```
- match:  
  - gateways:  
    - istio-egressgateway  
  port: 443  
  uri:  
    regex: "/health|/sport"  
  route:  
    - destination:  
        host: edition.cnn.com  
        port:  
          number: 443  
        weight: 100
```

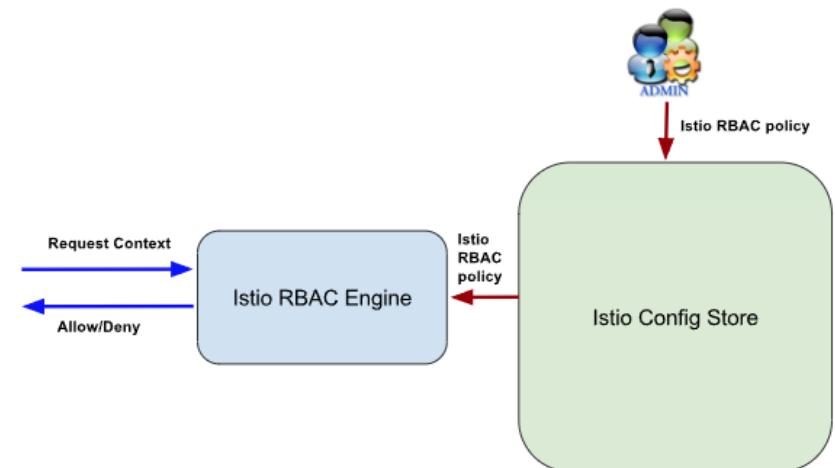
EOF

Different ways to implement Egress security with Istio

<https://istio.io/latest/blog/2018/egress-monitoring-access-control/>



Security – Role-Based Access Control (RBAC) with ISTIO



Security – RBAC



Istio RBAC policy

Istio RBAC introduces `ServiceRole` and `ServiceRoleBinding`, both of which are defined as Kubernetes CustomResourceDefinition (CRD) objects.

- `ServiceRole` defines a role for access to services in the mesh.
- `ServiceRoleBinding` grants a role to subjects (e.g., a user, a group, a service).

`ServiceRole`

A `ServiceRole` specification includes a list of rules. Each rule has the following standard fields:

- **services**: A list of service names, which are matched against the `action.service` field of the “requestcontext”.
- **methods**: A list of method names which are matched against the `action.method` field of the “requestcontext”. In the above “requestcontext”, this is the HTTP or gRPC method. Note that gRPC methods are formatted in the form of “packageName.serviceName/methodName” (case sensitive).
- **paths**: A list of HTTP paths which are matched against the `action.path` field of the “requestcontext”. It is ignored in gRPC case.

A `ServiceRole` specification only applies to the `namespace` specified in `metadata` section. The “services” and “methods” are required fields in a rule. “paths” is optional. If not specified or set to “*”, it applies to “any” instance.



Concepts

- ▶ What is Istio?
- ▶ Traffic Management
- ▼ Security
 - Authentication Policy
 - Mutual TLS
 - Authentication
 - Istio Role-Based Access Control (RBAC)
- ▶ Policies and Telemetry

Setup

Tasks

Guides

Performance and Scalability

Reference

ServiceRoleBinding <https://istio.io/v0.8/docs/concepts/security/rbac/>

A **ServiceRoleBinding** specification includes two parts:

- **roleRef** refers to a **ServiceRole** resource **in the same namespace**.
- A list of **subjects** that are assigned the role.

A subject can either be a “user”, or a “group”, or is represented with a set of “properties”. Each entry (“user” or “group” or an entry in “properties”) must match one of fields (“user” or “groups” or an entry in “properties”) in the “subject” part of the “requestcontext” instance.

Here is an example of **ServiceRoleBinding** resource “test-binding-products”, which binds two subjects to ServiceRole “product-viewer”:

- user “alice@yahoo.com”.
- “reviews.abc.svc.cluster.local” service in “abc” namespace.

```
apiVersion: "config.istio.io/v1alpha2"
kind: ServiceRoleBinding
metadata:
  name: test-binding-products
  namespace: default
spec:
  subjects:
    - user: "alice@yahoo.com"
    - properties:
        service: "reviews.abc.svc.cluster.local"
        namespace: "abc"
  roleRef:
    kind: ServiceRole
    name: "products-viewer"
```

In the case that you want to make a service(s) publicly accessible, you can use set the subject to **user: “*”**. This will assign a **ServiceRole** to all users/services.

- Summary
 - uServices need Service Mesh
 - Istio Provides a lot of Great Features