

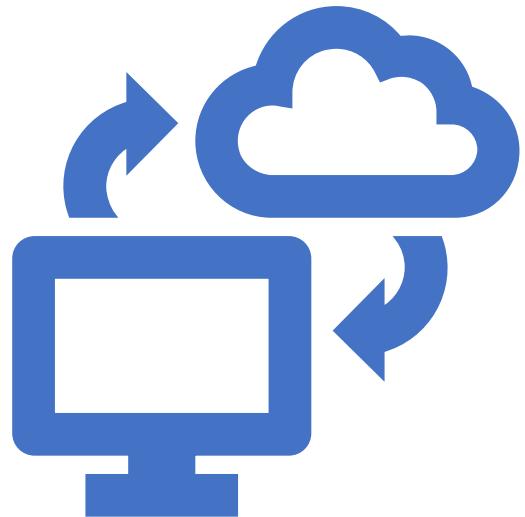
MapReduce Cloud Computing Models & Hadoop

- **ECE 4150, CS 4893**
- **Vijay Madisetti**
- **Fall 2024**
- **Georgia Tech**



MR – A Model of Computation Developed for Cloud

- We did not specifically look at new models of computation in earlier models discussed in class.
- Today, we will look at a new model of computing of that is specially adapted to distributed cloud computing ->>
MR



Types of Logs Collected by Enterprises

- **Hardware & Networking Devices** – servers, routers, switches, bridges, hubs, firewalls, etc. – often distributed across multiple datacenters
- **Operating Systems** – Windows & Unix based operating systems event logs in various levels of information detail, containing registry, event, file system, Syslog, SNMP, NetFlow and other types of logs.
- **Database Systems** – transaction, query, recovery, configuration and audit log
- **Applications** – HTTP weblogs, FTP logs, Log4J, JMS, JMX, Net events, Portal logs, but mainly the core business applications that often have their own homegrown log formats
- **Customer Facing Data** – logs reporting on product usage, click stream data, shopping card logs, online transaction logs
- **Virtualization & Cloud** – hypervisor logs, VM logs, cloud app logs
- **Outside Datacentre** – Manufacturing, logistics logs, CDRs & IPDRs, power consumption, RFID or GPS data logs

https://www.youtube.com/watch?v=BPC_mCINSXk&feature=emb_logo

Big Data Sets are Very Large

Example: 20+ billion web pages \times 20KB = 400+ terabytes

- One computer can read 30-35 MB/sec from disk
 - ~four months to read the web
- ~1,000 hard drives just to store the web
- Even more to *do something with the data*

Parallel Computing

Good news: same problem with 1000 machines, < 3 hours

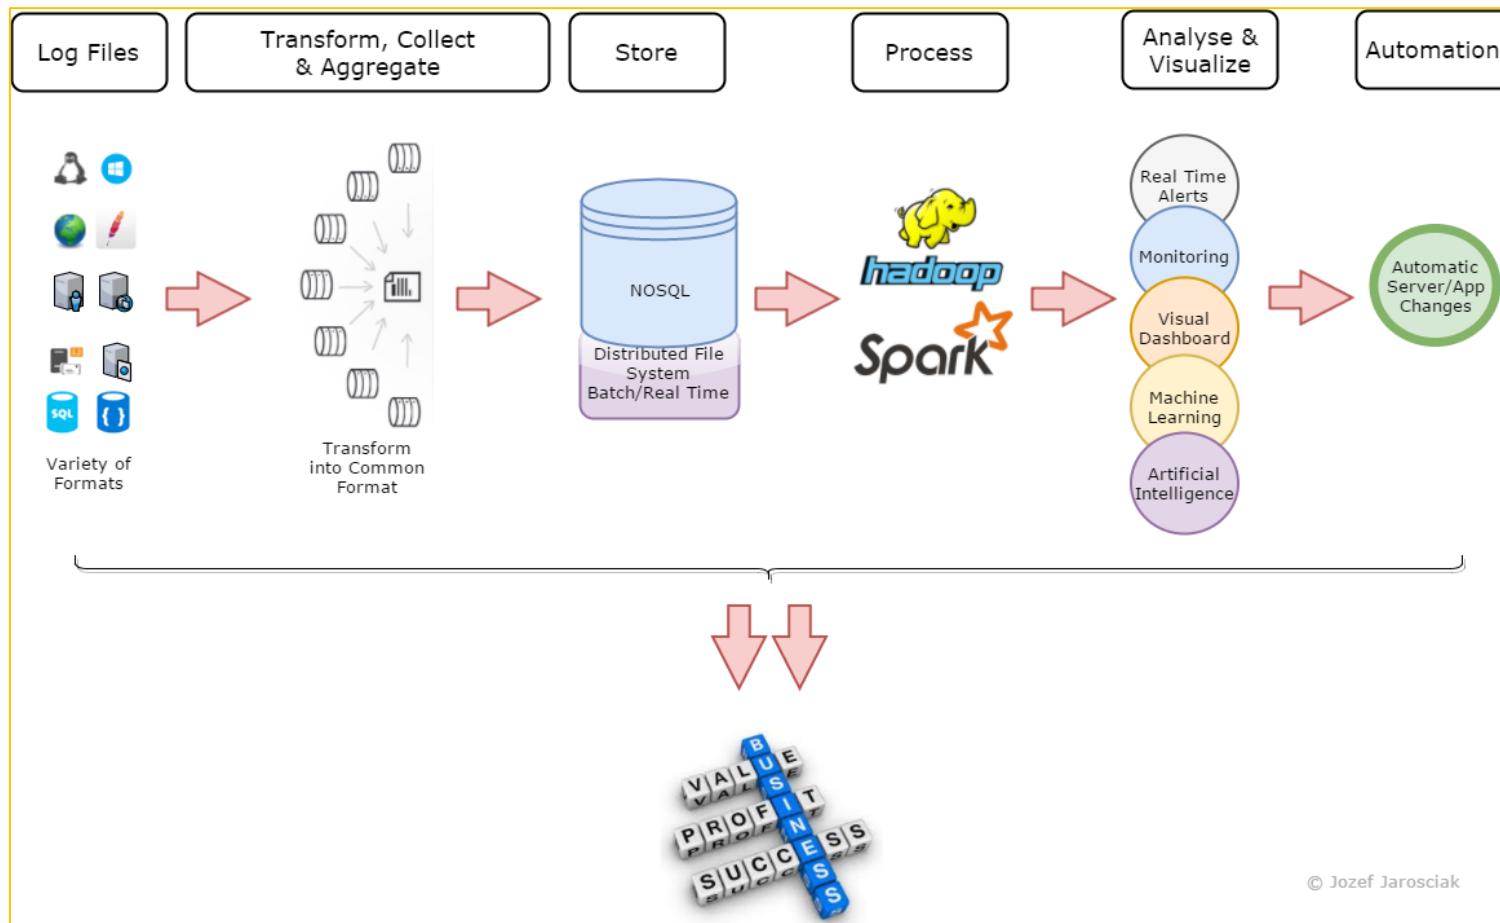
Bad news: programming work

- communication and coordination
- recovering from machine failure (all the time!)
- status reporting
- debugging
- optimization
- locality

Bad news II: repeat for every problem you want to solve

How can we make this easier?

Enterprise Data Analytics Flow



What is MapReduce MR ?

A simple programming model that applies to many large-scale computing problems



Hide messy details in MapReduce runtime library:

- automatic parallelization
- load balancing
- network and disk transfer optimization
- handling of machine failures
- robustness
- **improvements to core library benefit all users of library!**

MR approach

Read a lot of data

Map: extract something you care about from each record

Shuffle and Sort

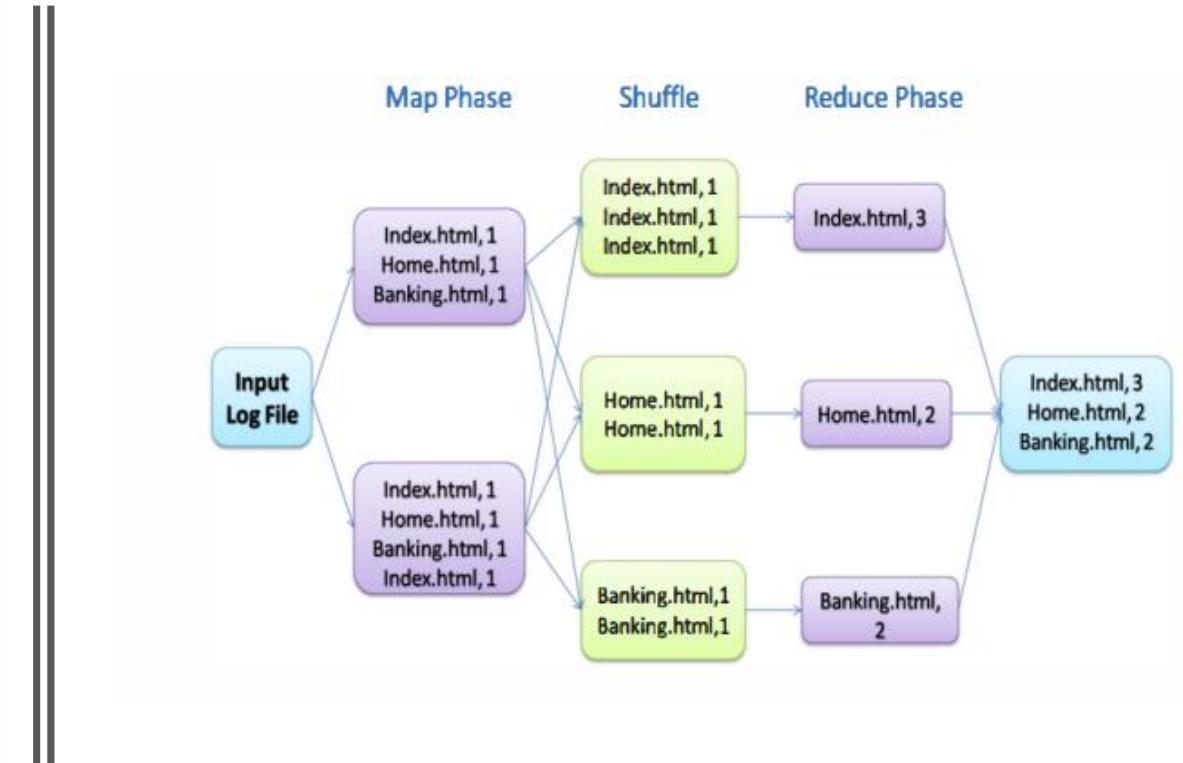
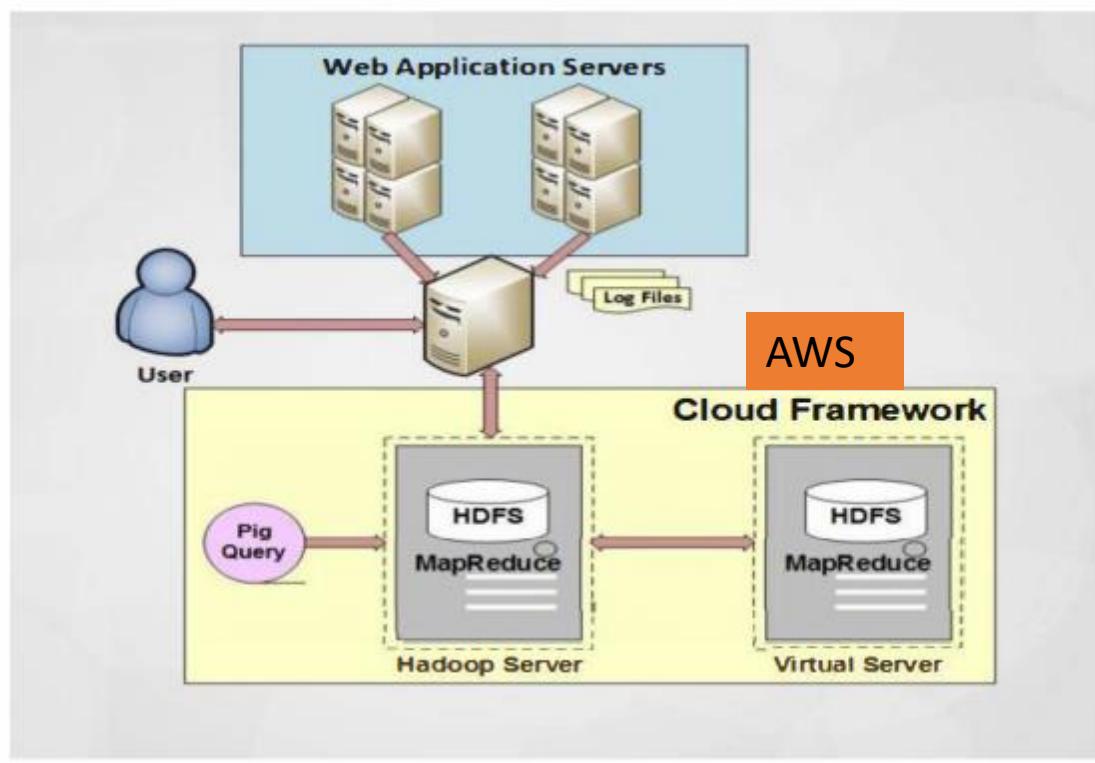
Reduce: aggregate, summarize, filter, or transform

Write the results

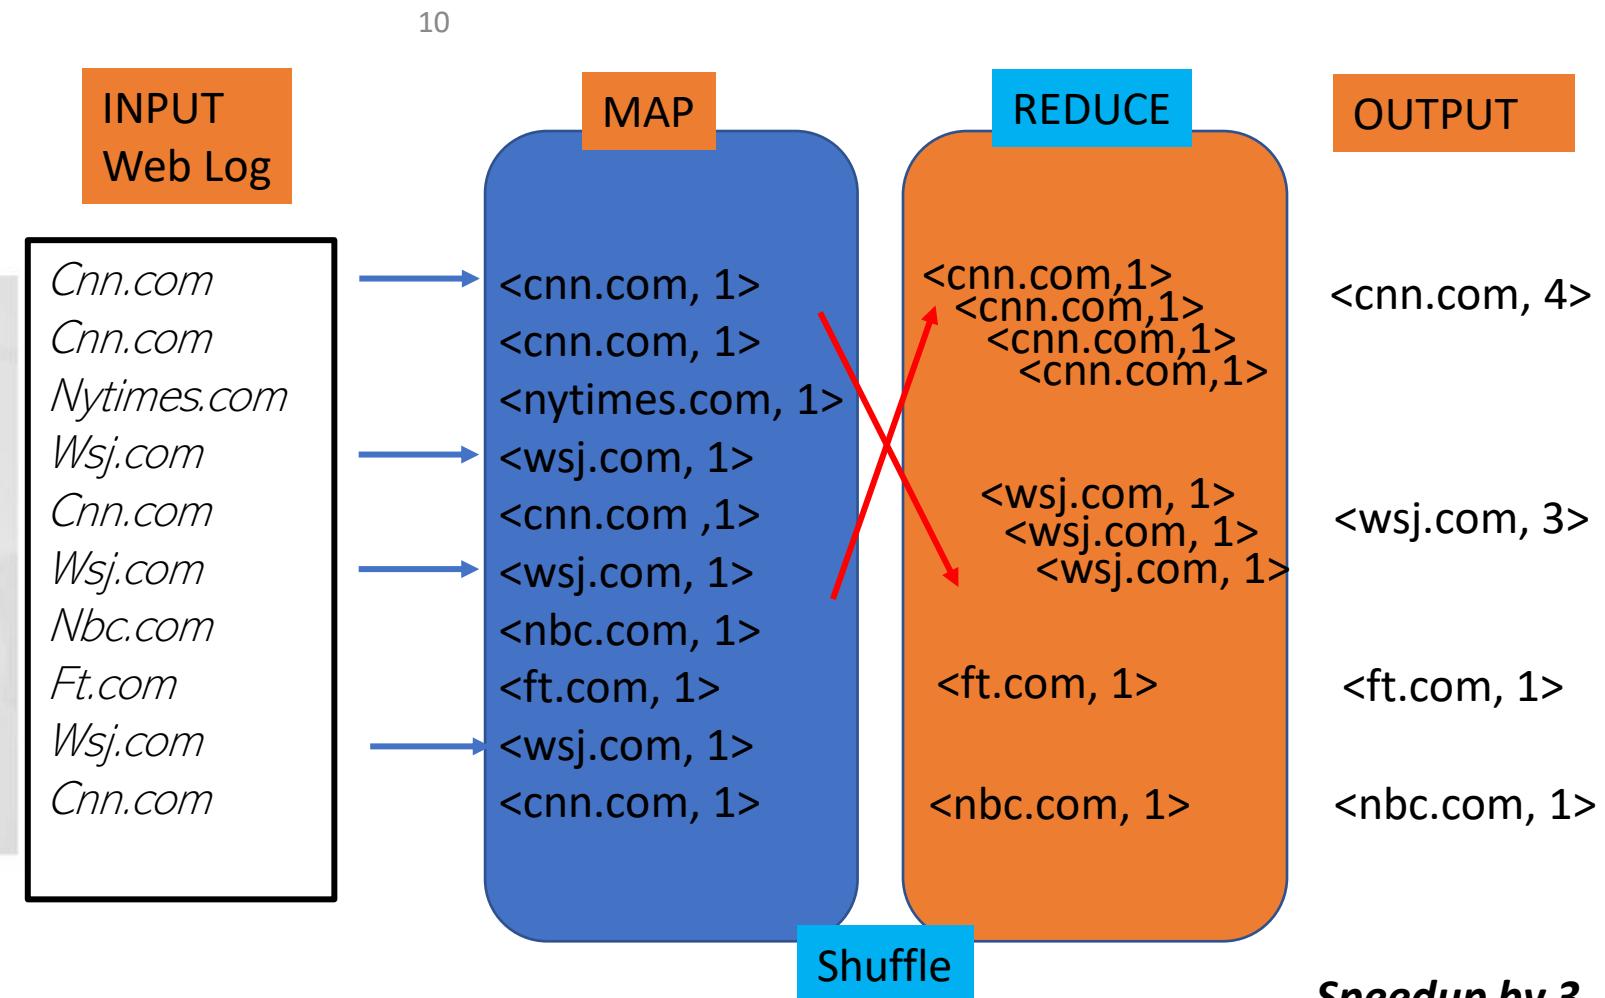
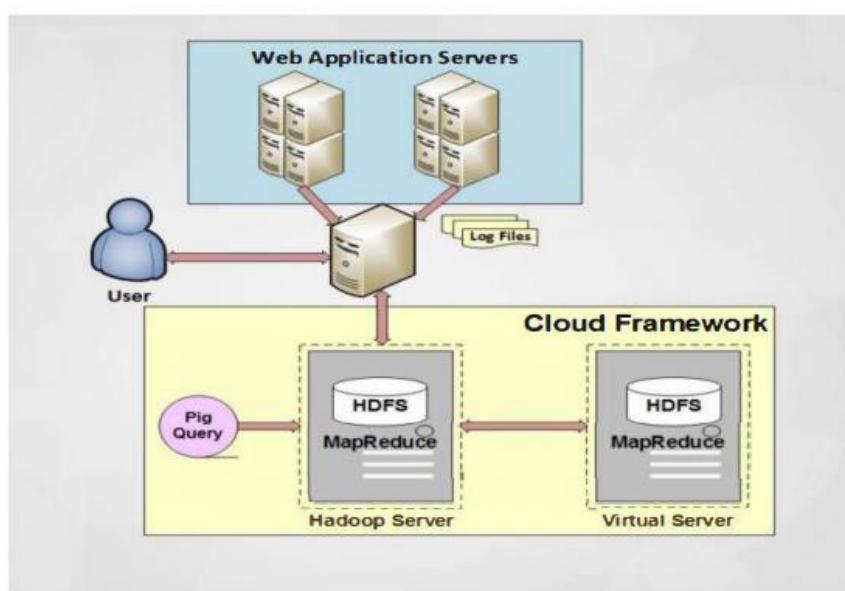
Outline stays the same,

Map and Reduce change to fit the problem

How Does MR work?

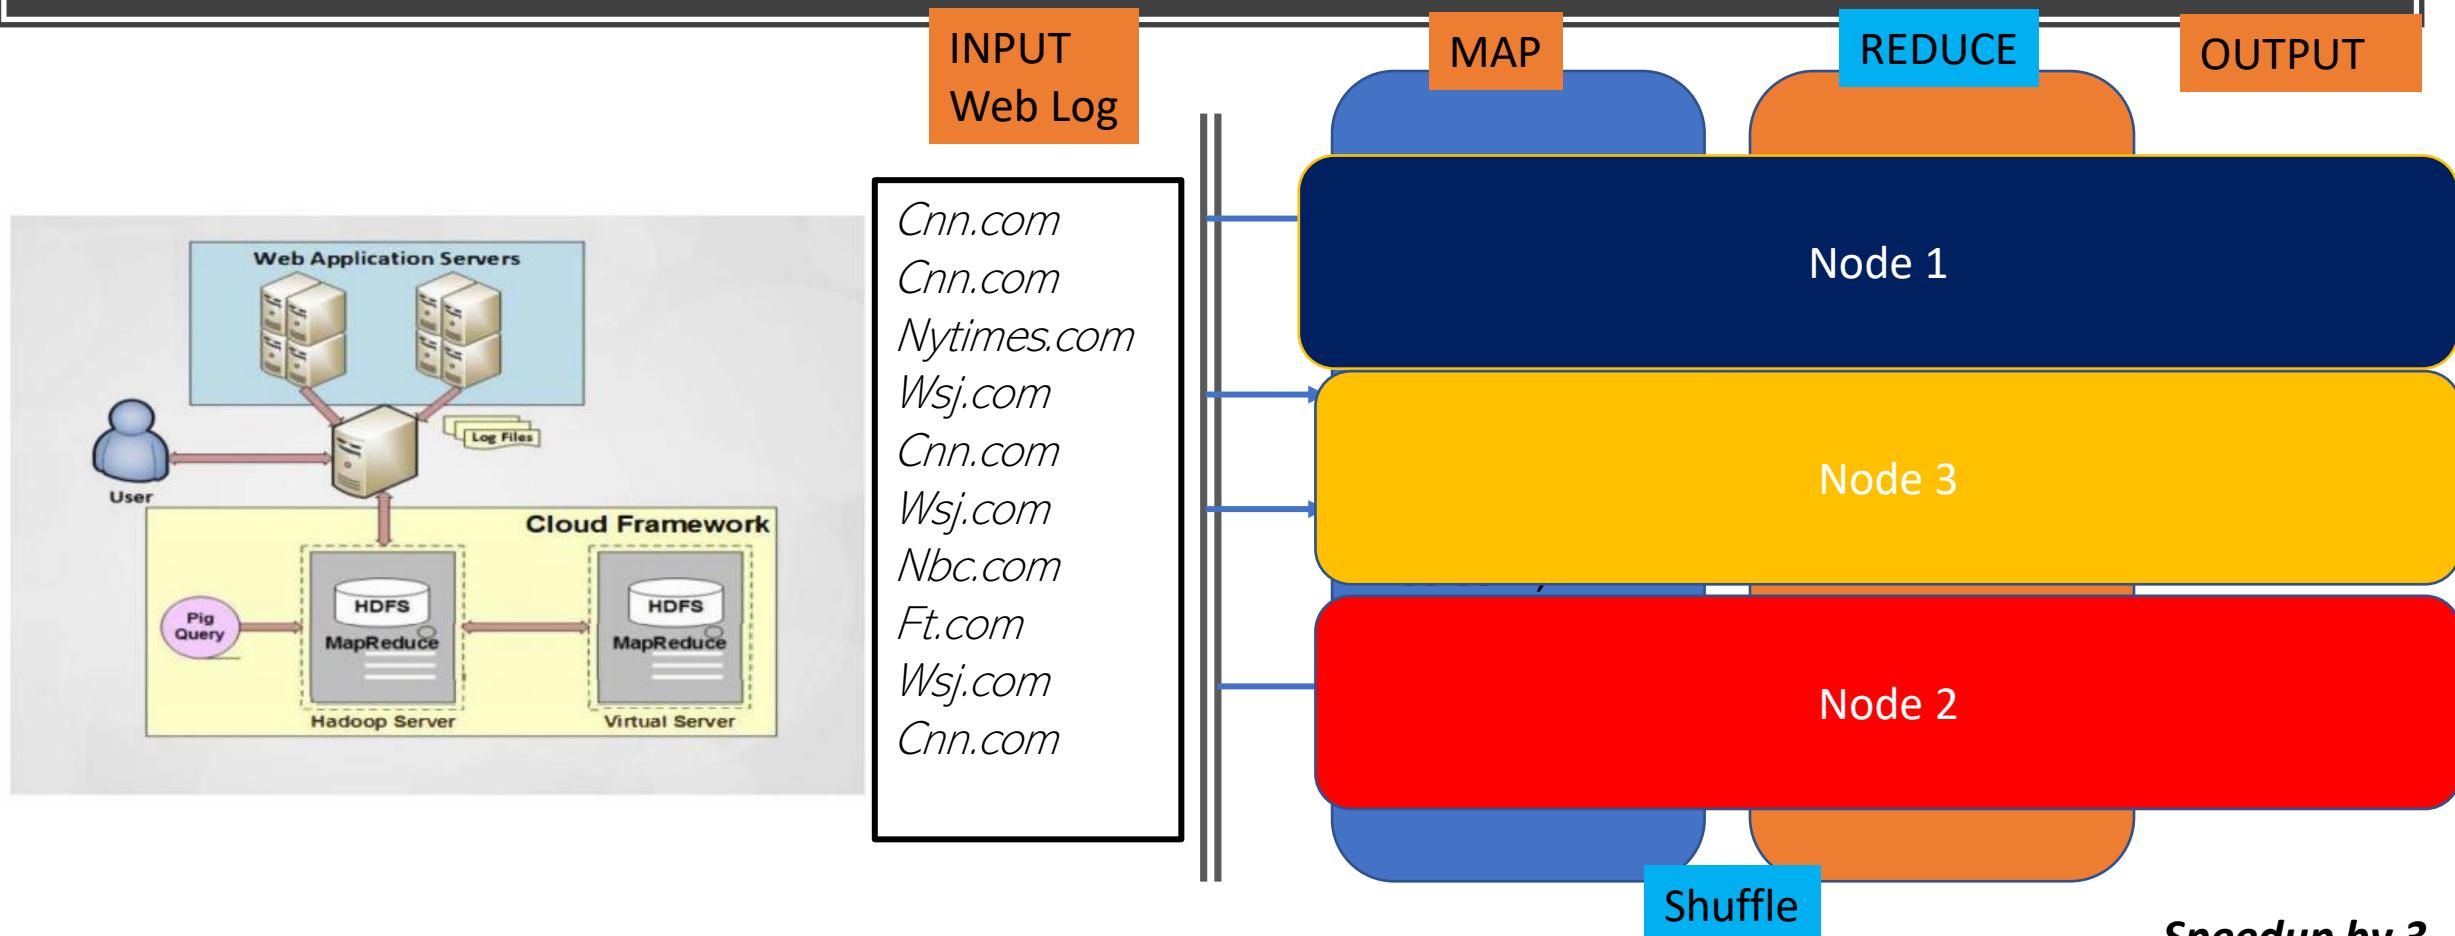


Hadoop/MR in Website Logs Analytics



Hadoop/MR in Website Logs Analytics

11



MR Template

Basic data type: the key-value pair (k, v).

For example, key = URL, value = HTML of the web page.

Programmer specifies two primary methods:

- **Map:** $(k, v) \mapsto \langle (k_1, v_1), (k_2, v_2), (k_3, v_3), \dots, (k_n, v_n) \rangle$
- **Reduce:** $(k', \langle v'_1, v'_2, \dots, v'_{n'} \rangle) \mapsto \langle (k', v''_1), (k', v''_2), \dots, (k', v''_{n''}) \rangle$

All v' with same k' are reduced together.

(Remember the invisible “Shuffle and Sort” step.)

Typical Training Exercise at



A typical exercise for a new Google engineer in his or her first week

Input: files with one document per record

Specify a *map* function that takes a key/value pair

key = document URL

value = document contents

Output of map function is (potentially many) key/value pairs.

In our case, output (word, “1”) once per word in the document

“document1”, “to be or not to be”



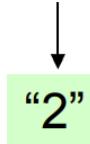
“to”, “1”
“be”, “1”
“or”, “1”

Word Frequencies in a Web Page

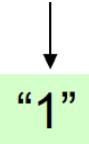
MapReduce library gathers together all pairs with the same key
(shuffle/sort)

Specify a *reduce* function that combines the values for a key
In our case, compute the sum

key = "be"
values = "1", "1"



key = "not"
values = "1"



key = "or"
values = "1"



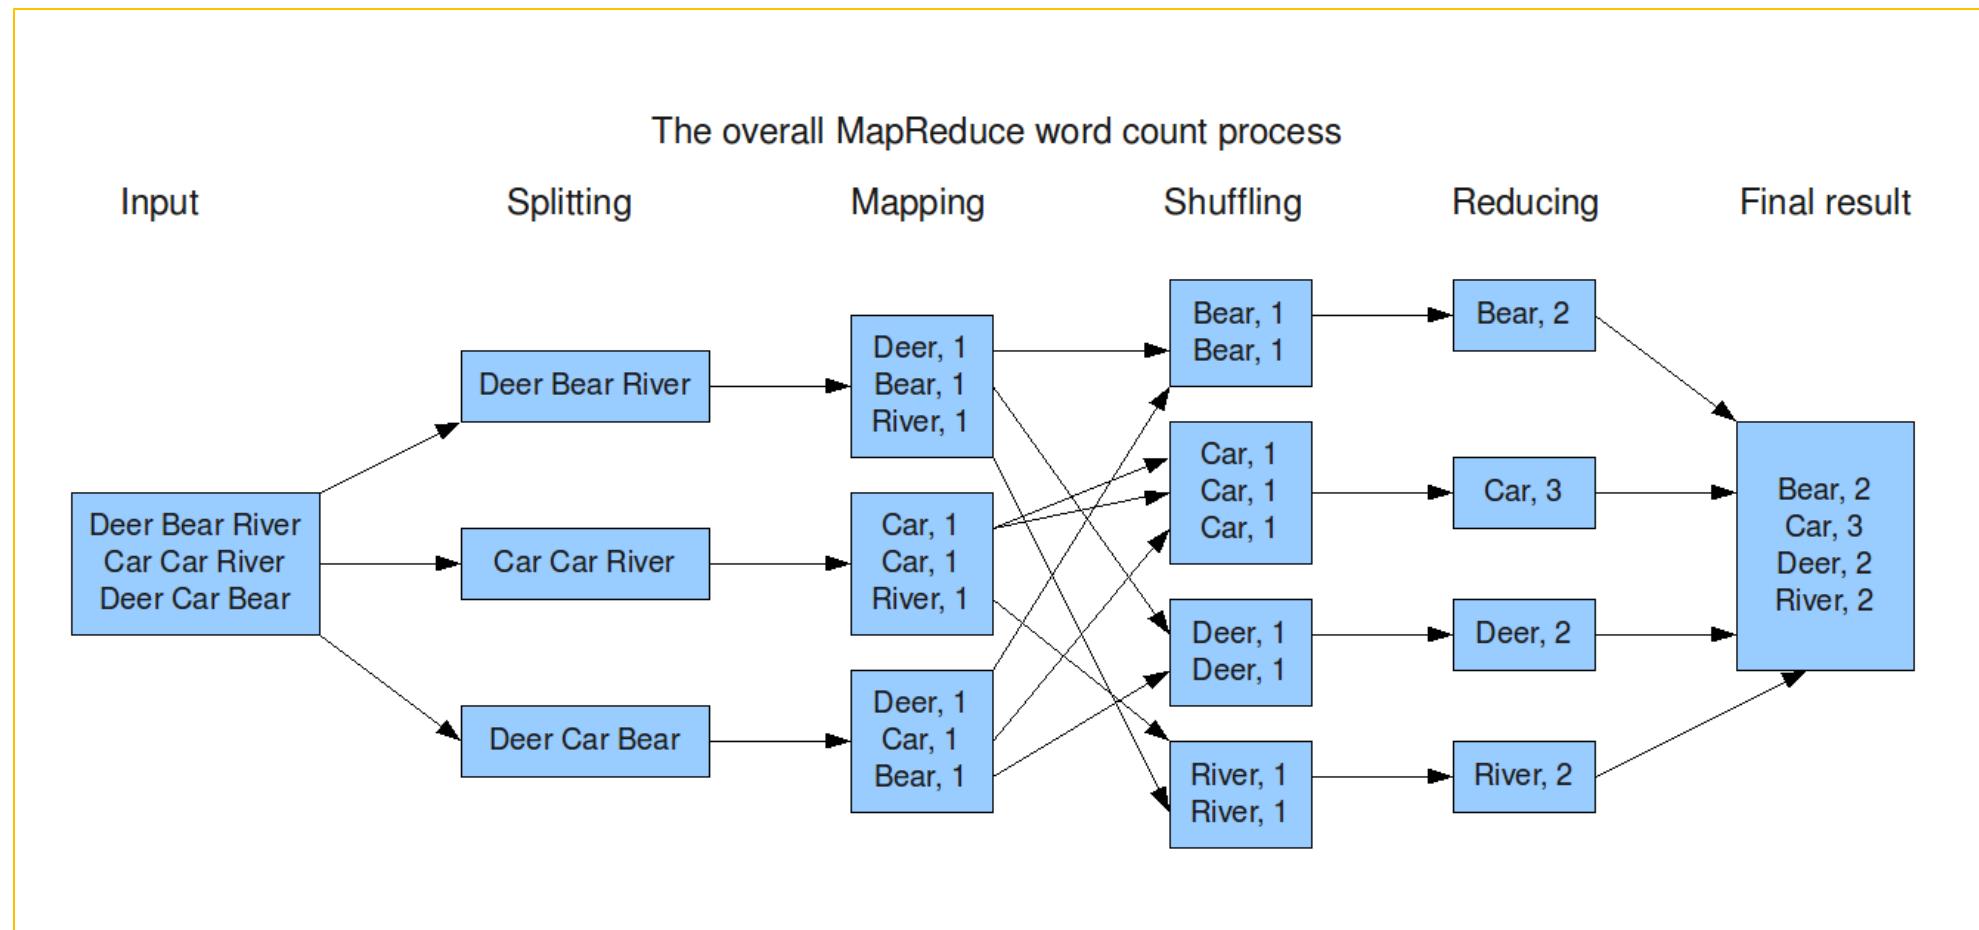
key = "to"
values = "1", "1"



Output of reduce (usually 0 or 1 value) paired with key and saved

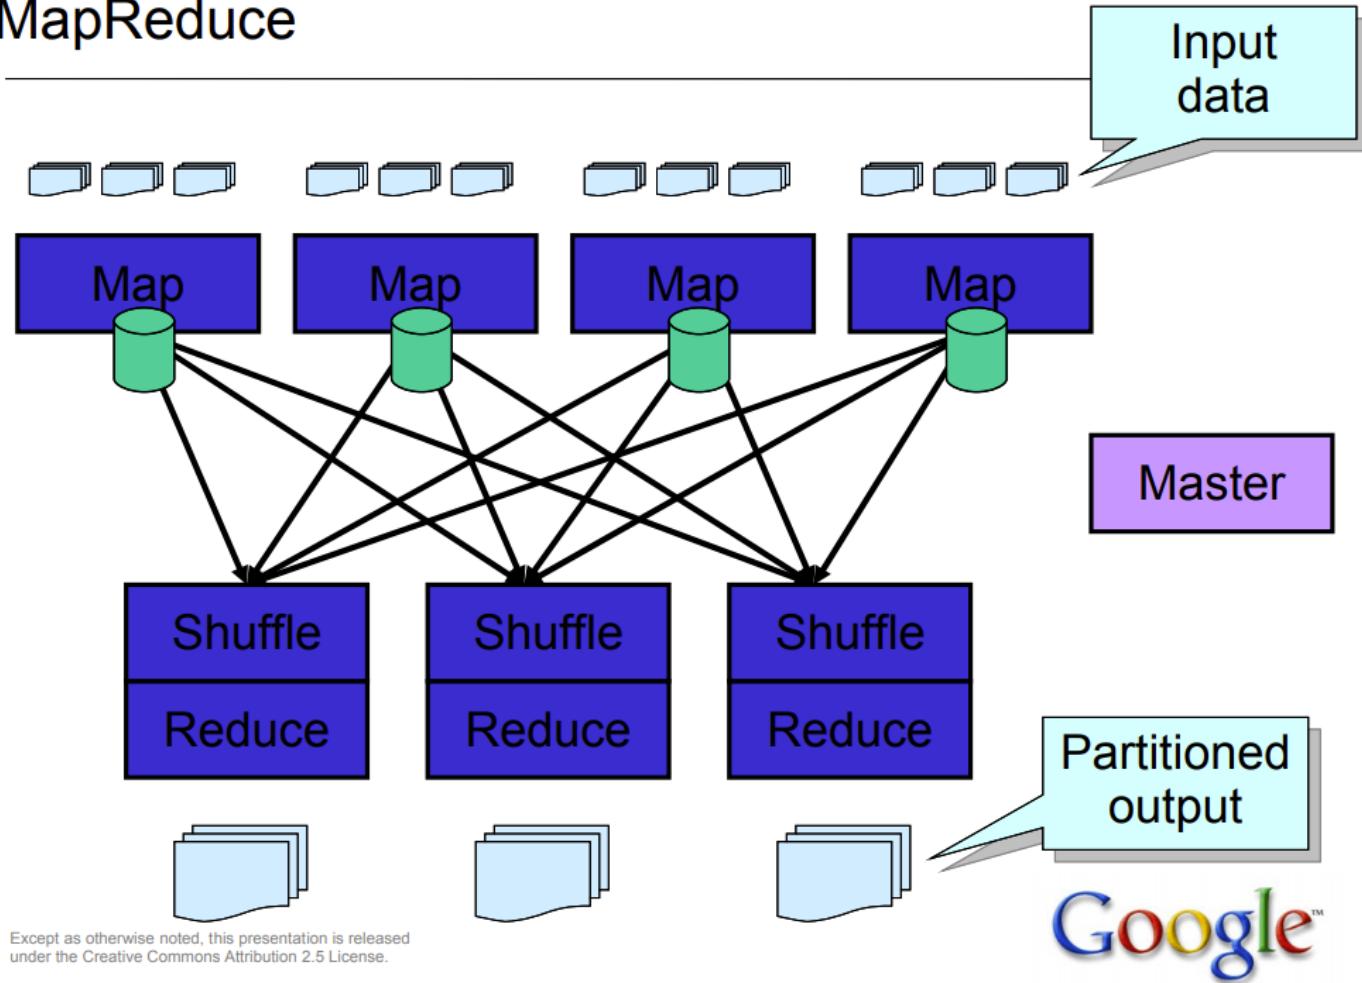
"be", "2"
"not", "1"
"or", "1"
"to", "2"

MapReduce: Word Count



MapReduce model at Google

MapReduce

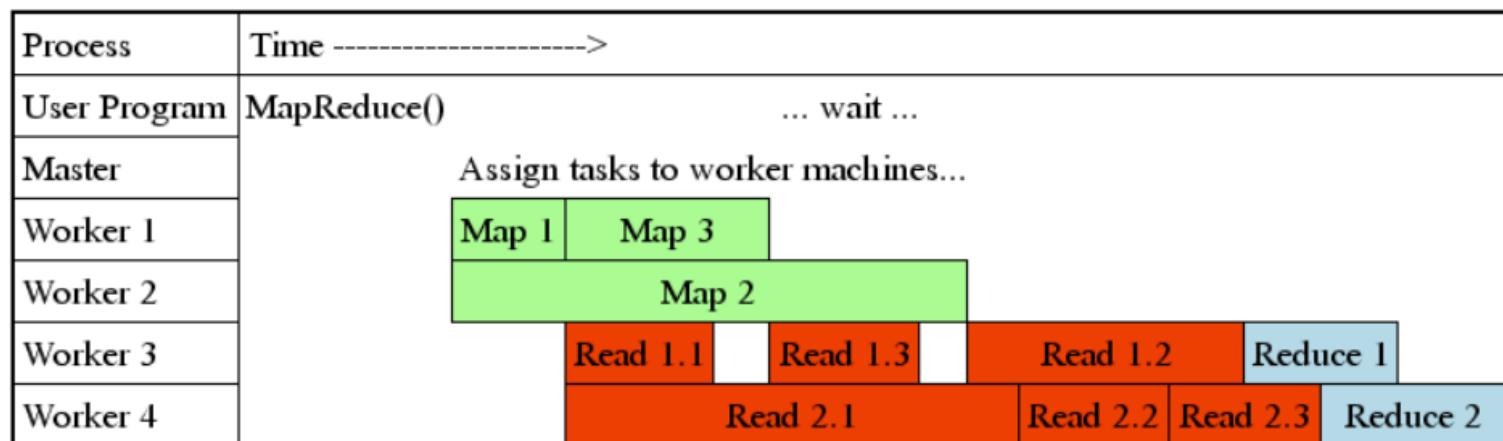


Except as otherwise noted, this presentation is released under the Creative Commons Attribution 2.5 License.

MapReduce Model at Google

Fine granularity tasks: many more map tasks than machines

- Minimizes time for fault recovery
- Can pipeline shuffling with map execution
- Better dynamic load balancing

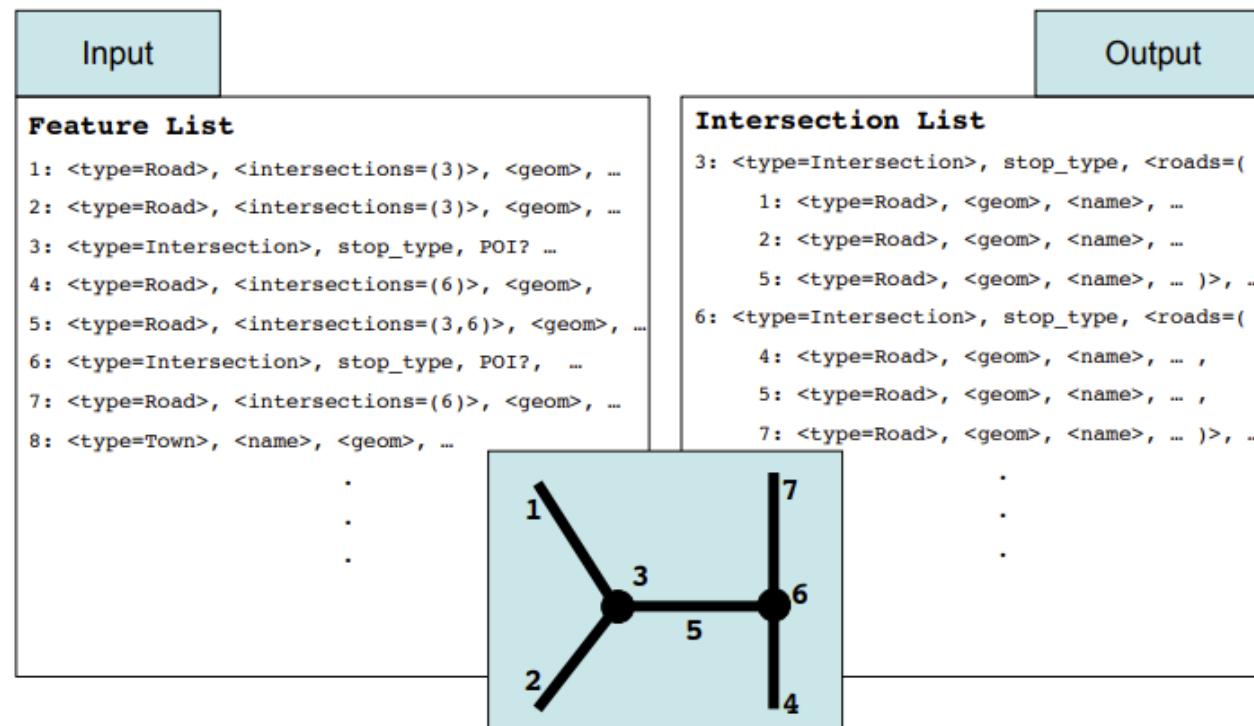


Except as otherwise noted, this presentation is released
under the Creative Commons Attribution 2.5 License.

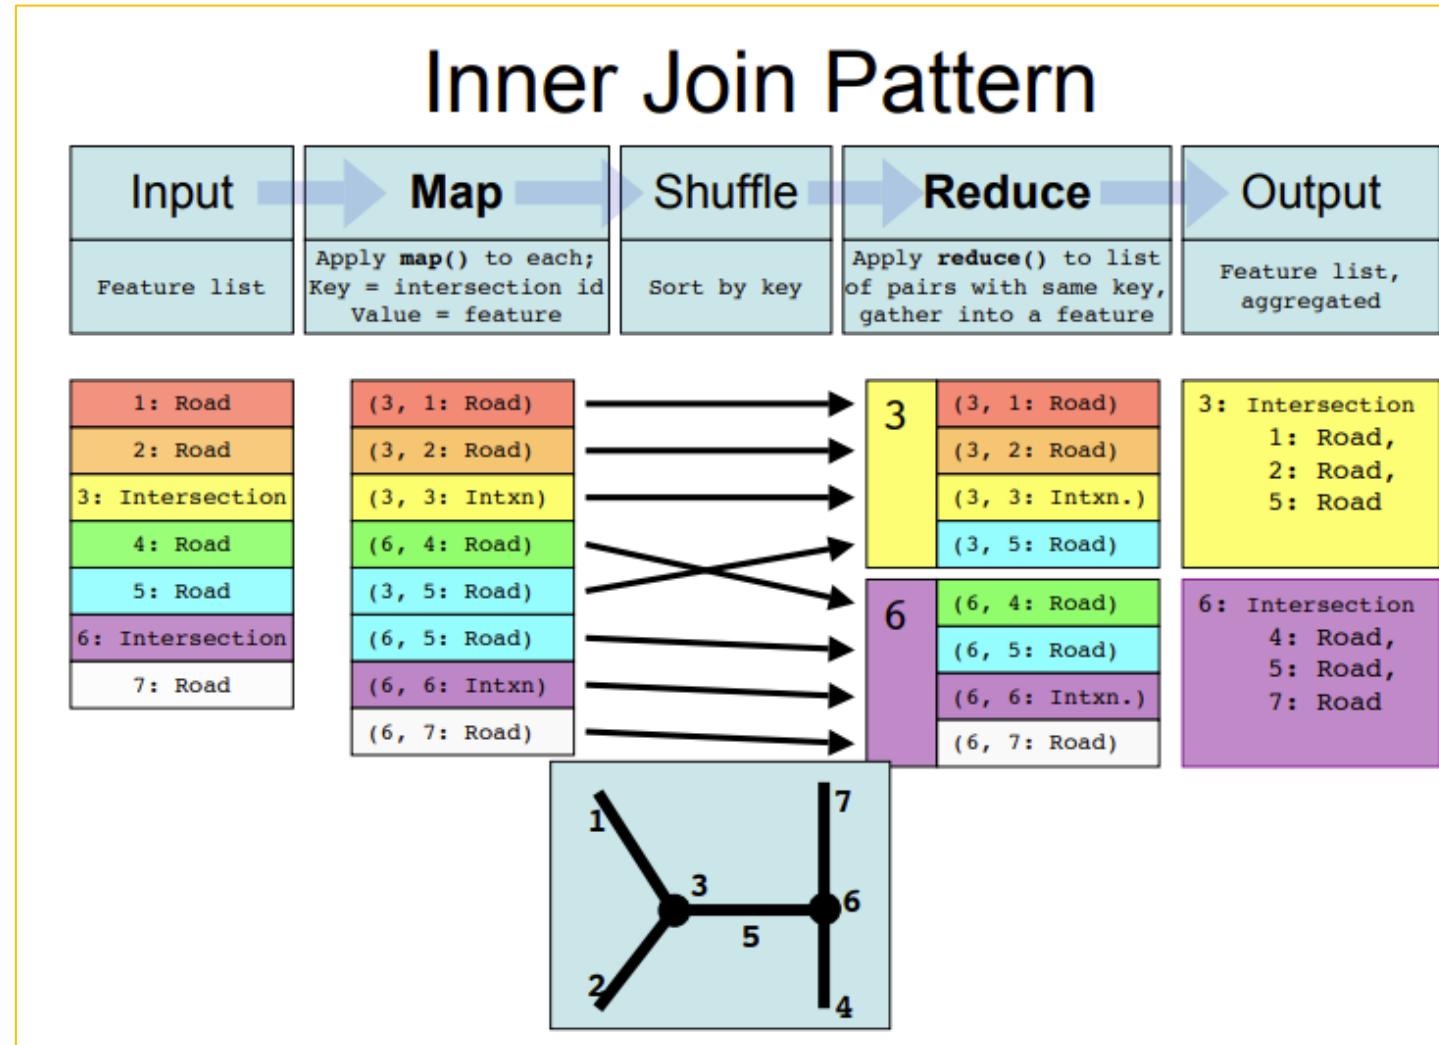


Google Maps using MR

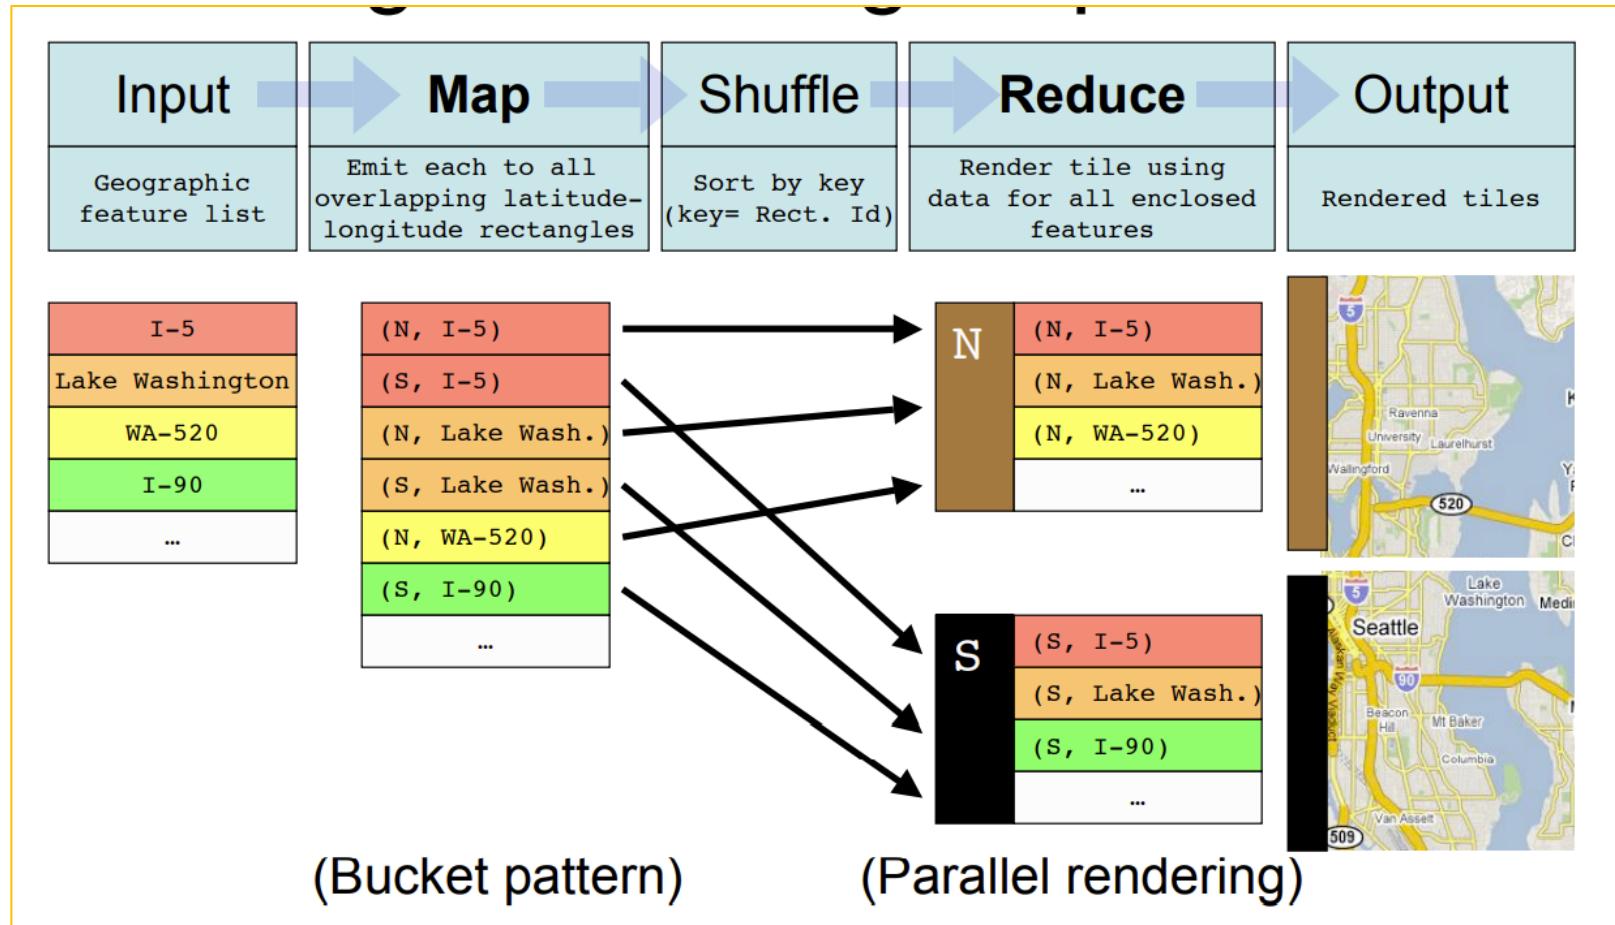
Pointer Following (or) Joining



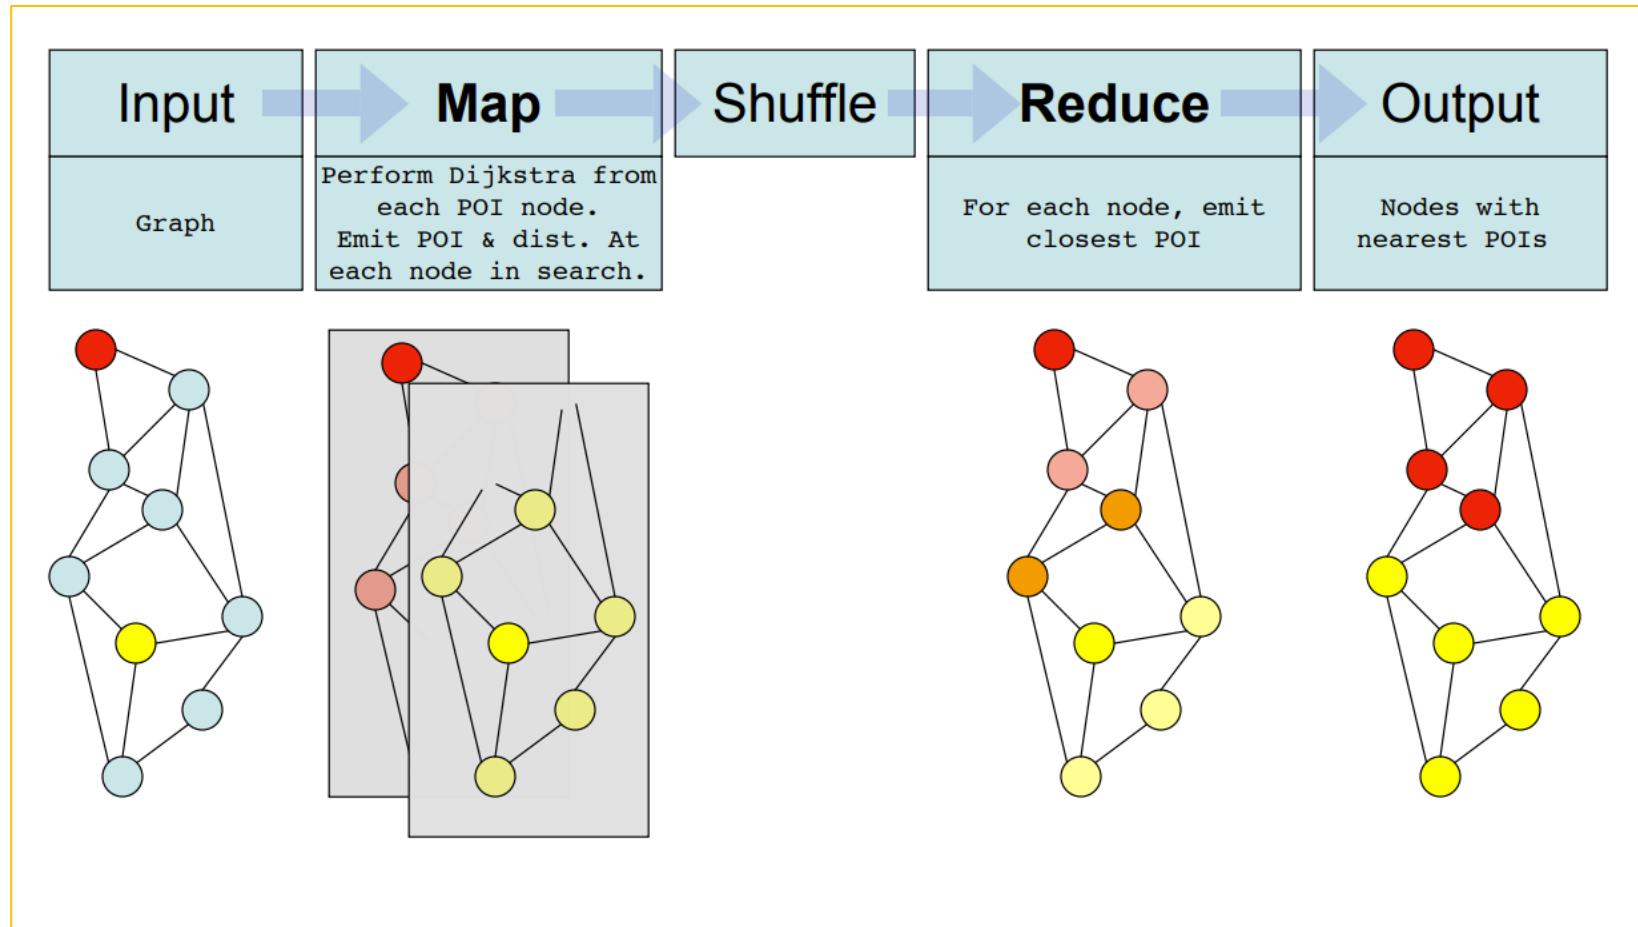
MR in Google Maps

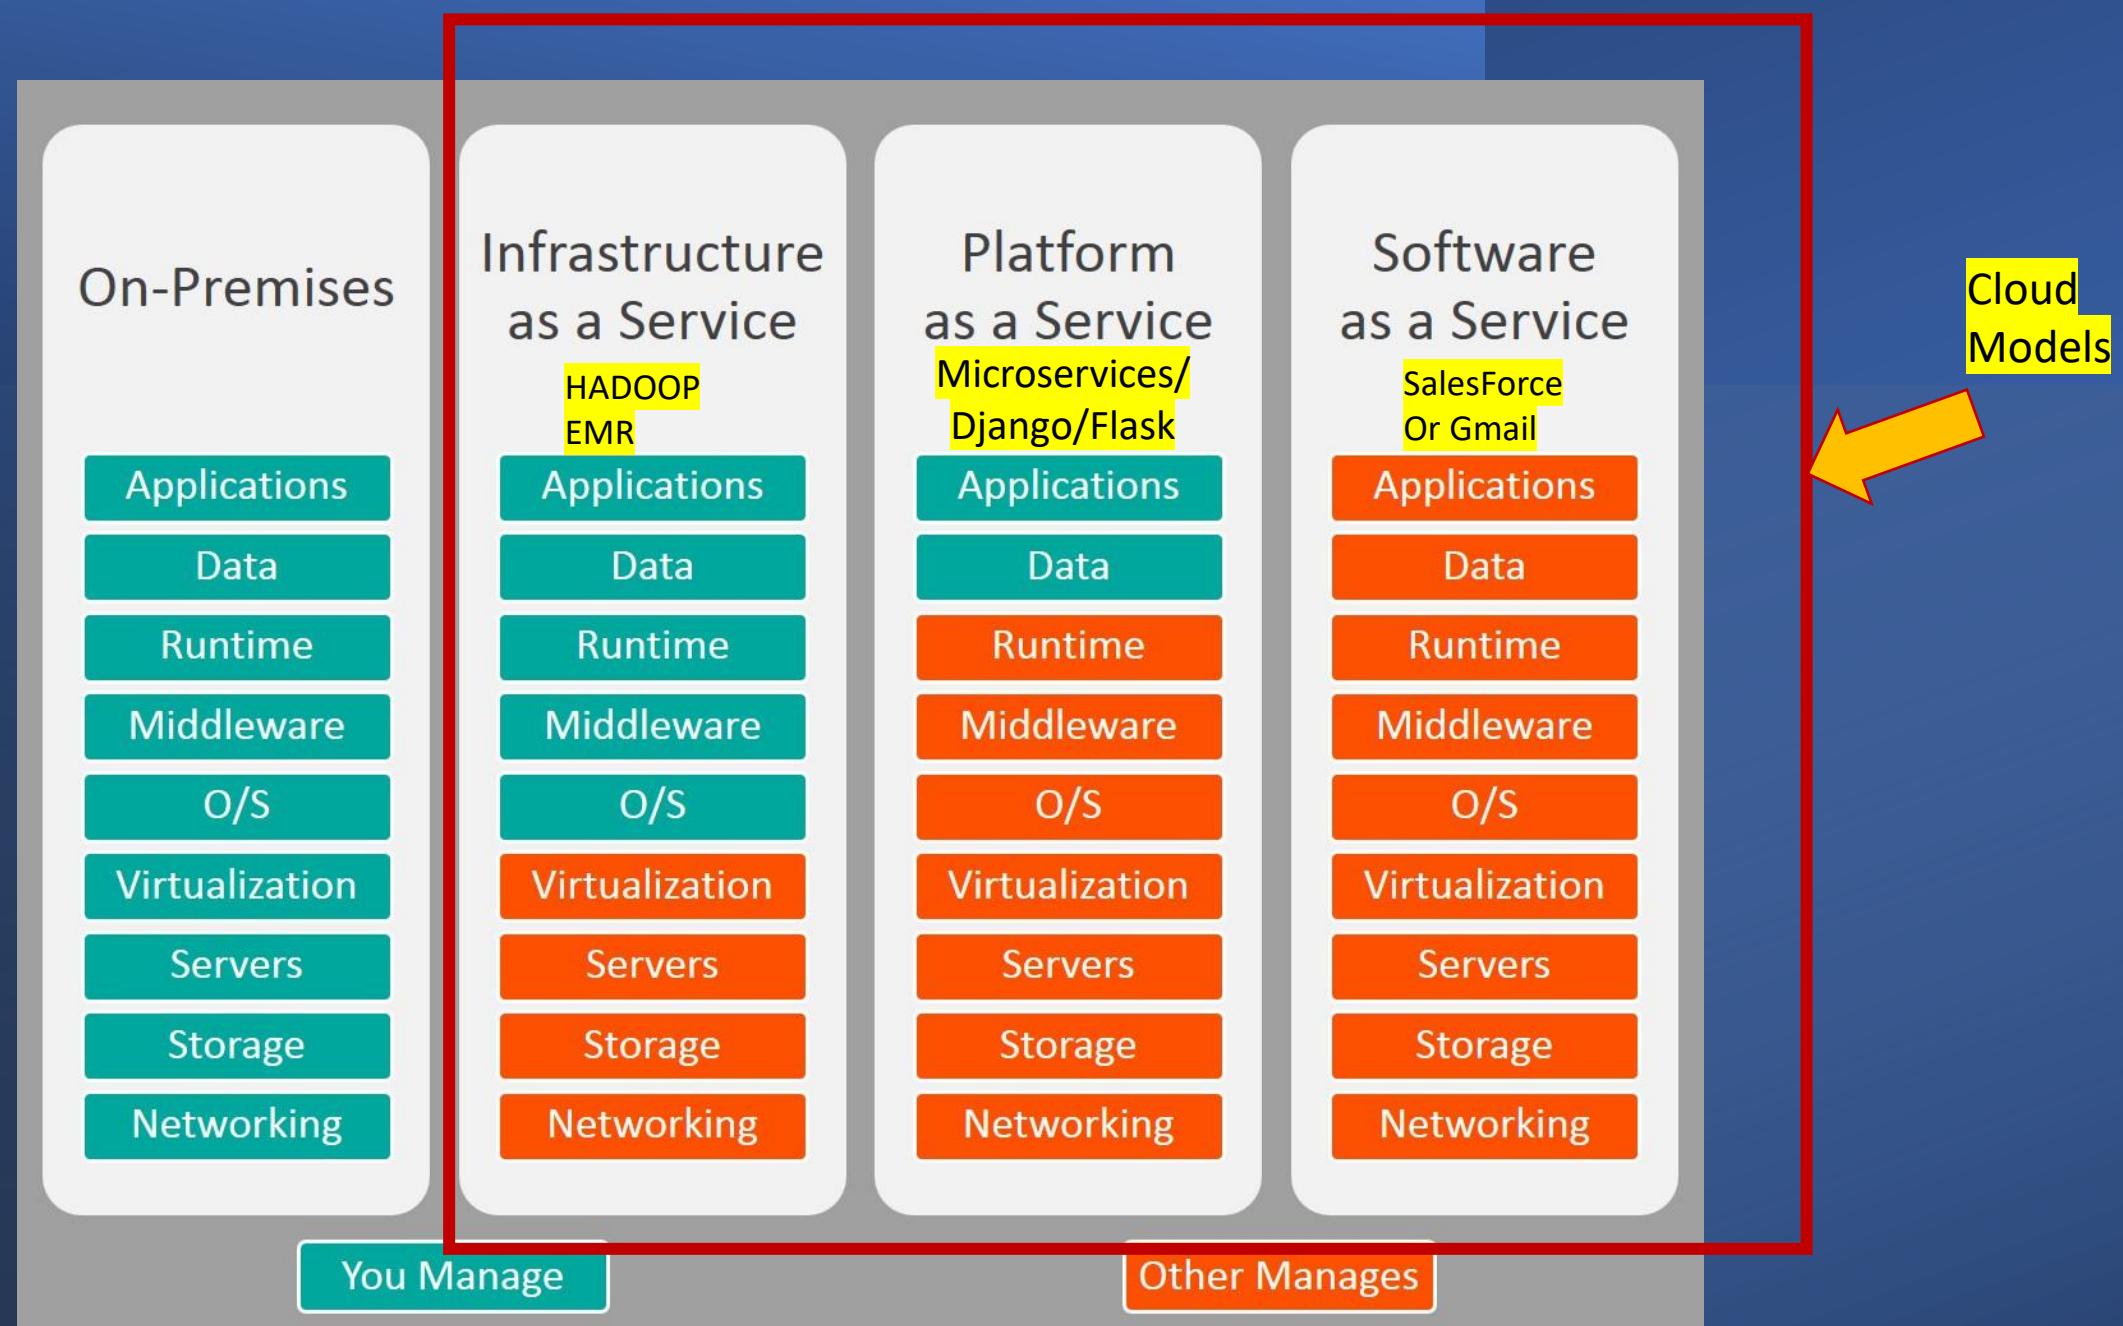


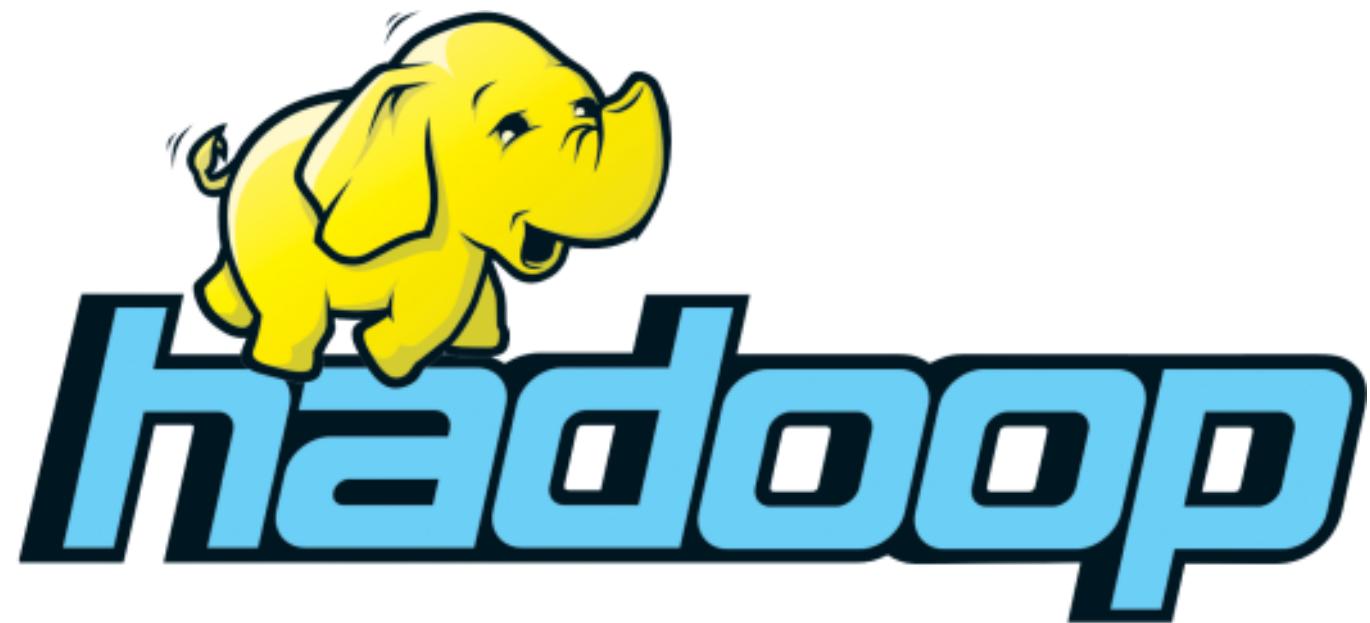
Distributing MR across map tiles



Finding Point of Interest (POI) in Google Maps







Hadoop

- Based on work done by Google in the early 2000s
 - “The Google File System” in 2003
 - “MapReduce: Simplified Data Processing on Large Clusters” in 2004
- The core idea was to distribute the data as it is initially stored
 - Each node can then perform computation on the data it stores without moving the data for the initial processing

Core Hadoop Concepts

- Applications are written in a high-level programming language
 - No network programming or temporal dependency
- Nodes should communicate as little as possible
 - A “shared nothing” architecture
- Data is spread among the machines in advance
 - Perform computation where the data is already stored as often as possible

Hadoop High-Level Overview

- When data is loaded onto the system it is divided into blocks
 - Typically , 64MB or 128MB
- Tasks are divided into two phases
 - Map tasks which are done on small portions of data where the data is stored
 - Reduce tasks which combine data to produce the final output
- A master program allocates work to individual nodes

Hadoop Overview



Responsible for storing data on the cluster



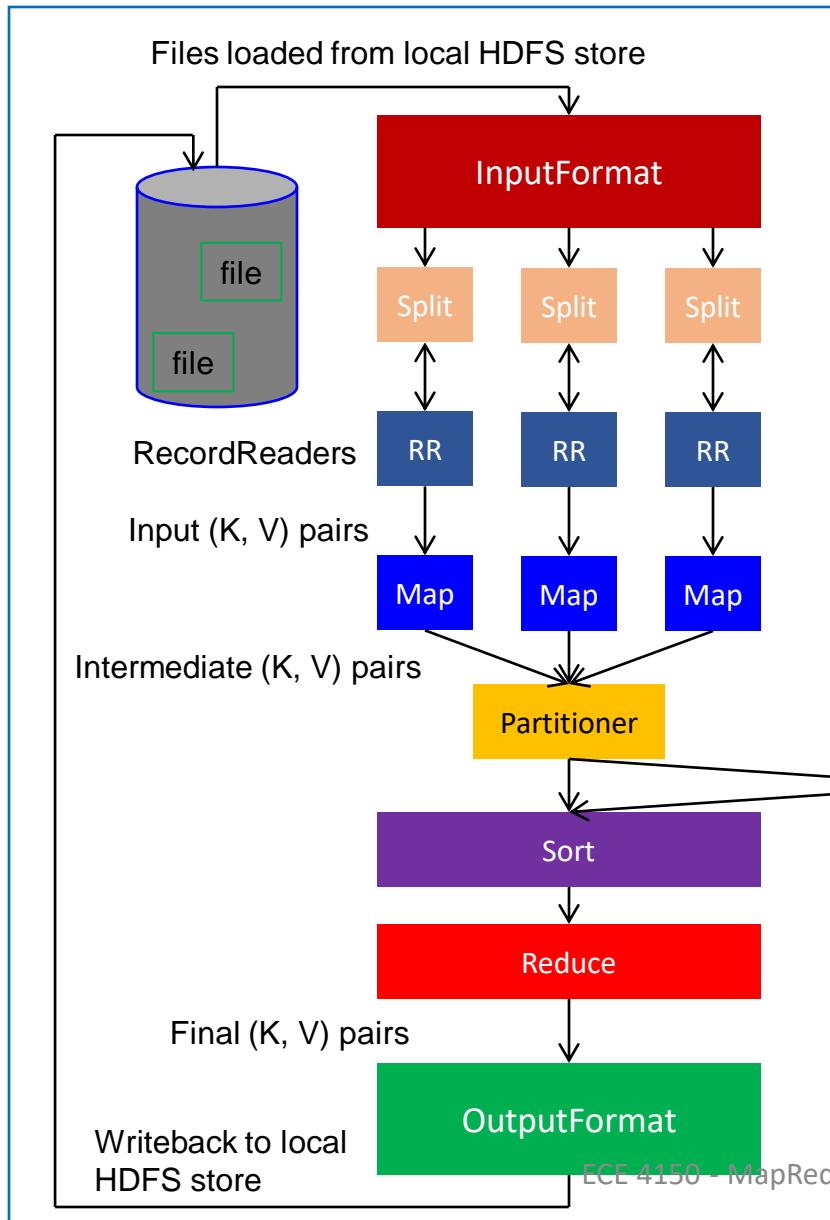
Data files are split into blocks and distributed across the nodes in the cluster



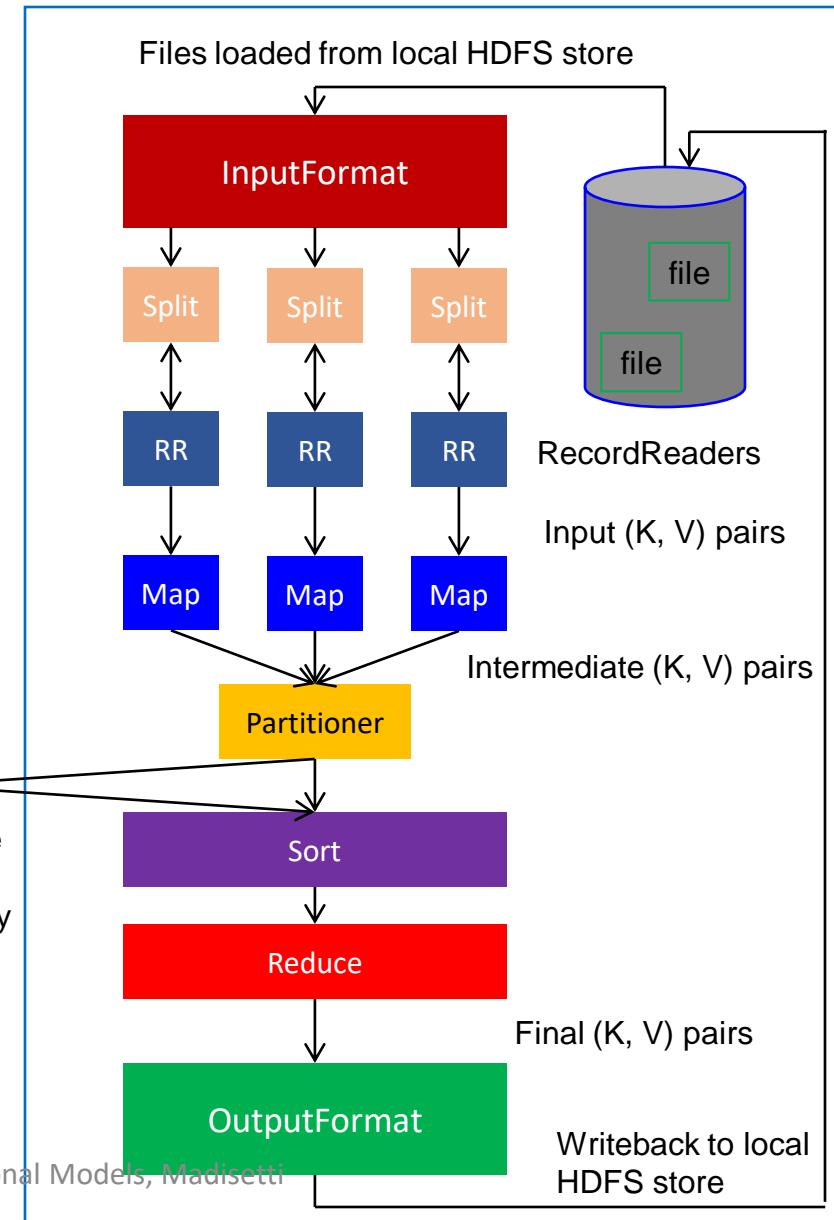
Each block is replicated multiple times

Hadoop MapReduce: A Closer Look

Node 1



Node 2



HDFS Basic Concepts

- HDFS is a file system written in Java based on the Google's GFS
- Provides redundant storage for massive amounts of data

MRJob Example for Log Analysis

Input

```
root@Acer-VKM: ~/vkm
0,Will,33,385
1,Jean-Luc,26,2
2,Hugh,55,221
3,Deanna,40,465
4,Quark,68,21
5,Weyoun,59,318
6,Gowron,37,220
7,Will,54,307
8,Jadzia,38,380
9,Hugh,27,181
10,Odo,53,191
11,Ben,57,372
12,Keiko,54,253
13,Jean-Luc,56,444
14,Hugh,43,49
15,Rom,36,49
16,Weyoun,22,323
17,Odo,35,13
18,Jean-Luc,45,455
19,Geordi,60,246
20,Odo,67,220
--More--(3%)
```

MR Job

```
root@Acer-VKM:~/vkm# more Friends-By-Age-2.py
from mrjob.job import MRJob

class MRFriendsByAge(MRJob):

    def mapper(self, _, line):
        (ID, name, age, numFriends) = line.split(',')
        yield age, float(numFriends)

    def reducer(self, age, numFriends):
        total = 0
        numElements = 0
        for x in numFriends:
            total += x
            numElements += 1

        yield age, total / numElements

if __name__ == '__main__':
    MRFriendsByAge.run()
```

<Index, Name, Age, Number of Friends>

Output

```
root@Acer-VKM: ~/vkm
Streaming final output from /tmp/Friends-By-Age-2
...
"18"      343.375
"19"      213.27272727272728
"20"      165.0
"21"      350.875
"22"      206.42857142857142
"23"      246.3
"24"      233.8
"25"      197.45454545454547
"26"      242.05882352941177
"27"      228.125
"28"      209.1
"29"      215.91666666666666
"30"      235.8181818181818
"31"      267.25
"32"      207.9090909090909
"33"      325.33333333333333
"34"      245.5
"35"      211.625
"36"      246.6
"37"      249.33333333333334
```

<Age, Number of Friends>

AWS MapReduce Support from 2009

Screenshot of the AWS News Blog post titled "Announcing Amazon Elastic MapReduce" from March 2009.

The page header includes the AWS logo, navigation links (Products, Solutions, Pricing, Documentation, Learn, Partner Network, AWS Marketplace, Customer Enablement, Events, Explore More), and a search bar. A prominent orange button says "Create an AWS Account".

The main content starts with a brief introduction: "Over the past 3 or 4 years, scientists, researchers, and commercial developers have recognized and embraced the [MapReduce](#) programming model. Originally described in a landmark paper, the MapReduce model is ideal for processing large data sets on a cluster of processors. It is easy to scale up a MapReduce application to jobs of arbitrary size by simply adding more compute power. Here's a very simple overview of the data flow in a typical MapReduce job:

Diagram illustrating the data flow in a typical MapReduce job:

```
graph LR; ID((Input Data)) --> IDP1((Input Data Part 1)); ID --> IDPN((Input Data Part N)); IDP1 --> MI1[Map Instance #1]; IDPN --> MIn[Map Instance #N]; MI1 --> RI[Reduce Instance]; MIn --> RI; RI --> OD((Output Data));
```

The right sidebar features a "Follow" section with links to social media platforms (Twitter, Facebook, LinkedIn, Twitch, RSS Feed, Email Updates) and a promotional box for "AWS Pi Week 2021".

[Announcing Amazon Elastic MapReduce | AWS News Blog](#) in 2009



Hadoop Architecture for Cloud

ECE 4150

Vijay Madisetti

How does the cloud computer look like ?

Central Ideas of Hadoop Ecosystem

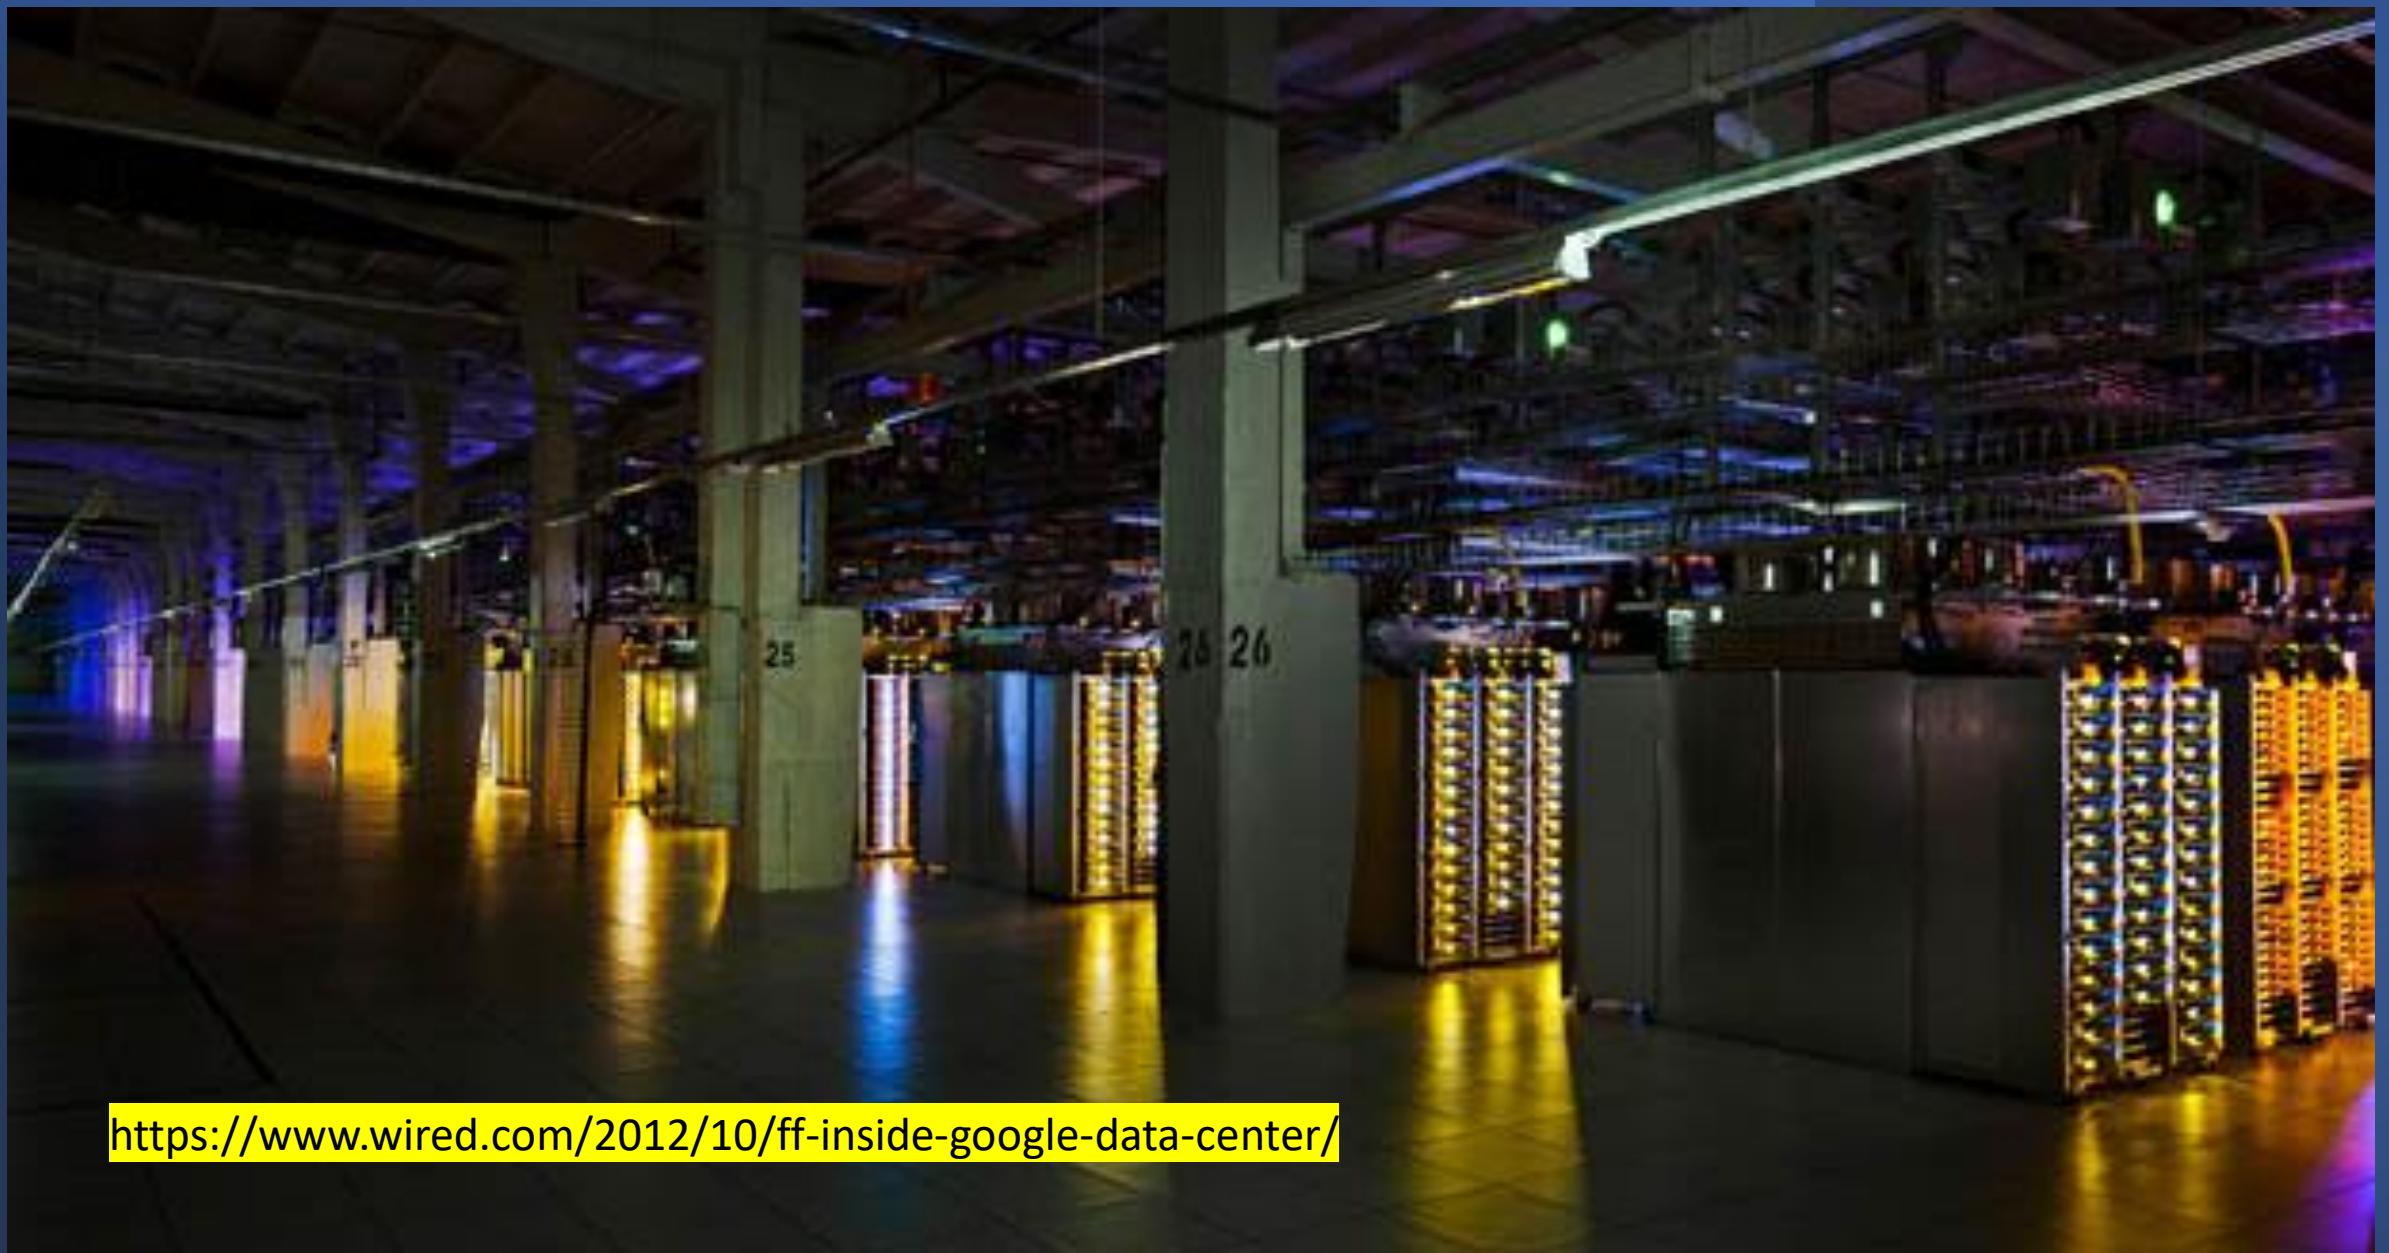


**DON'T TELL ME DETAILS!!
I DON'T CARE!!**





<https://www.wired.com/2012/10/ff-inside-google-data-center/>



<https://www.wired.com/2012/10/ff-inside-google-data-center/>











Water

Interesting videos on data centers

- Microsoft GFS Datacenter Tour (Youtube)
 - <http://www.youtube.com/watch?v=hOxA111pQIw>
- Timelapse of a Datacenter Construction on the Inside (Fortune 500 company)
 - <http://www.youtube.com/watch?v=ujO-xNvXj3g>



Hadoop Ecosystem

oozie
(Work flow)

HCatalog
Table & schema Management

 Pig
(Scripting)
 Hive
(Sql Query)

 Mahout
(Machine Learning)
 Drill
(Interactive Analysis)

 AVRO
(JSON)
 Thrift
(Cross Language Service)

 APACHE HBASE
HBASE
(Columnar Store)

 Sqoop
(Data Collection)

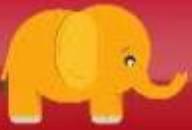
 Zookeeper
(Coordination)

 Ambari
Apache Ambari
(Management & Monitoring)

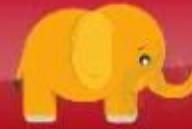
Mapreduce
(Data Processing)



Yarn
(Cluster Resource Management)



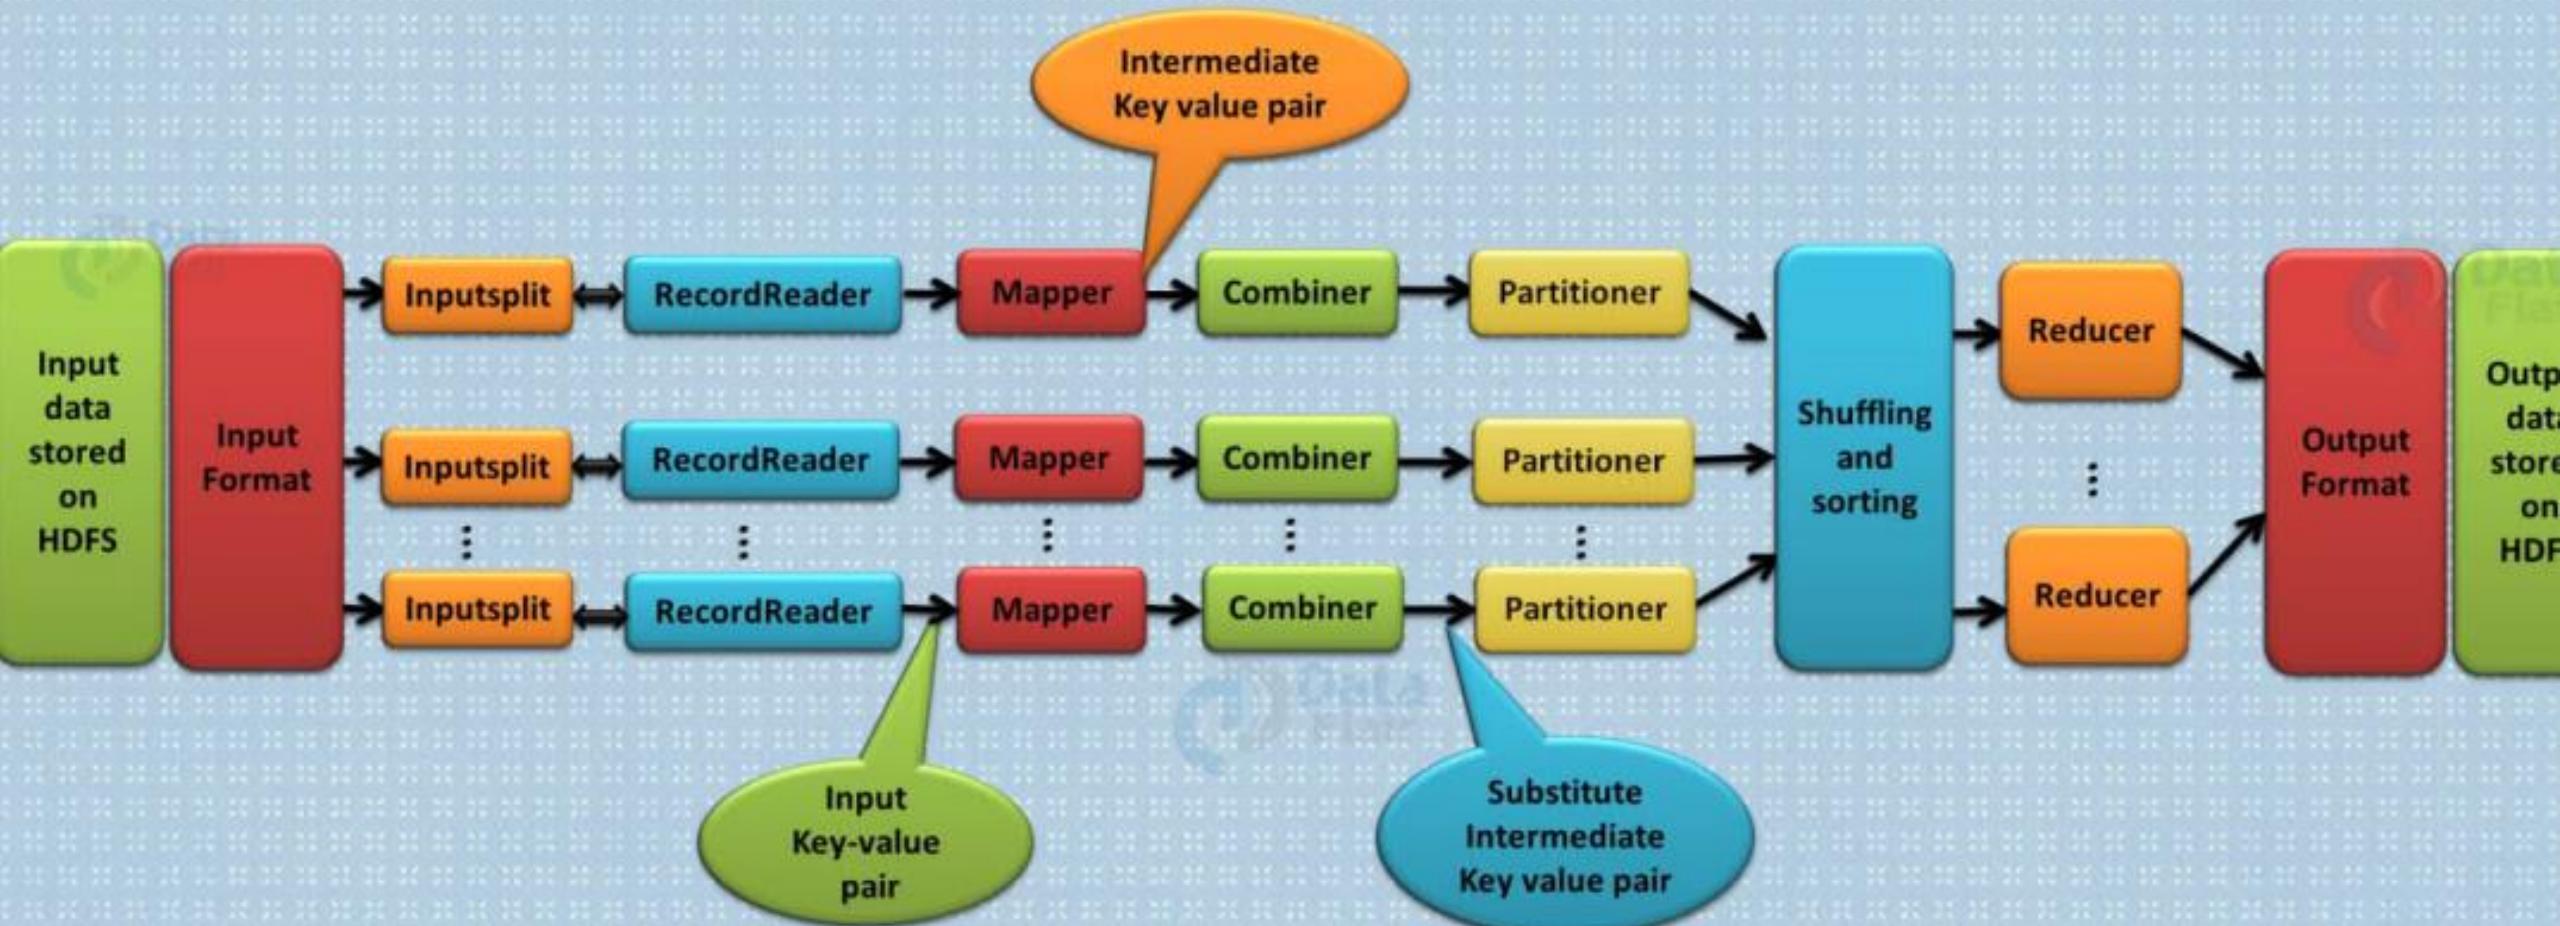
HDFS
(Hadoop Distributed File system)



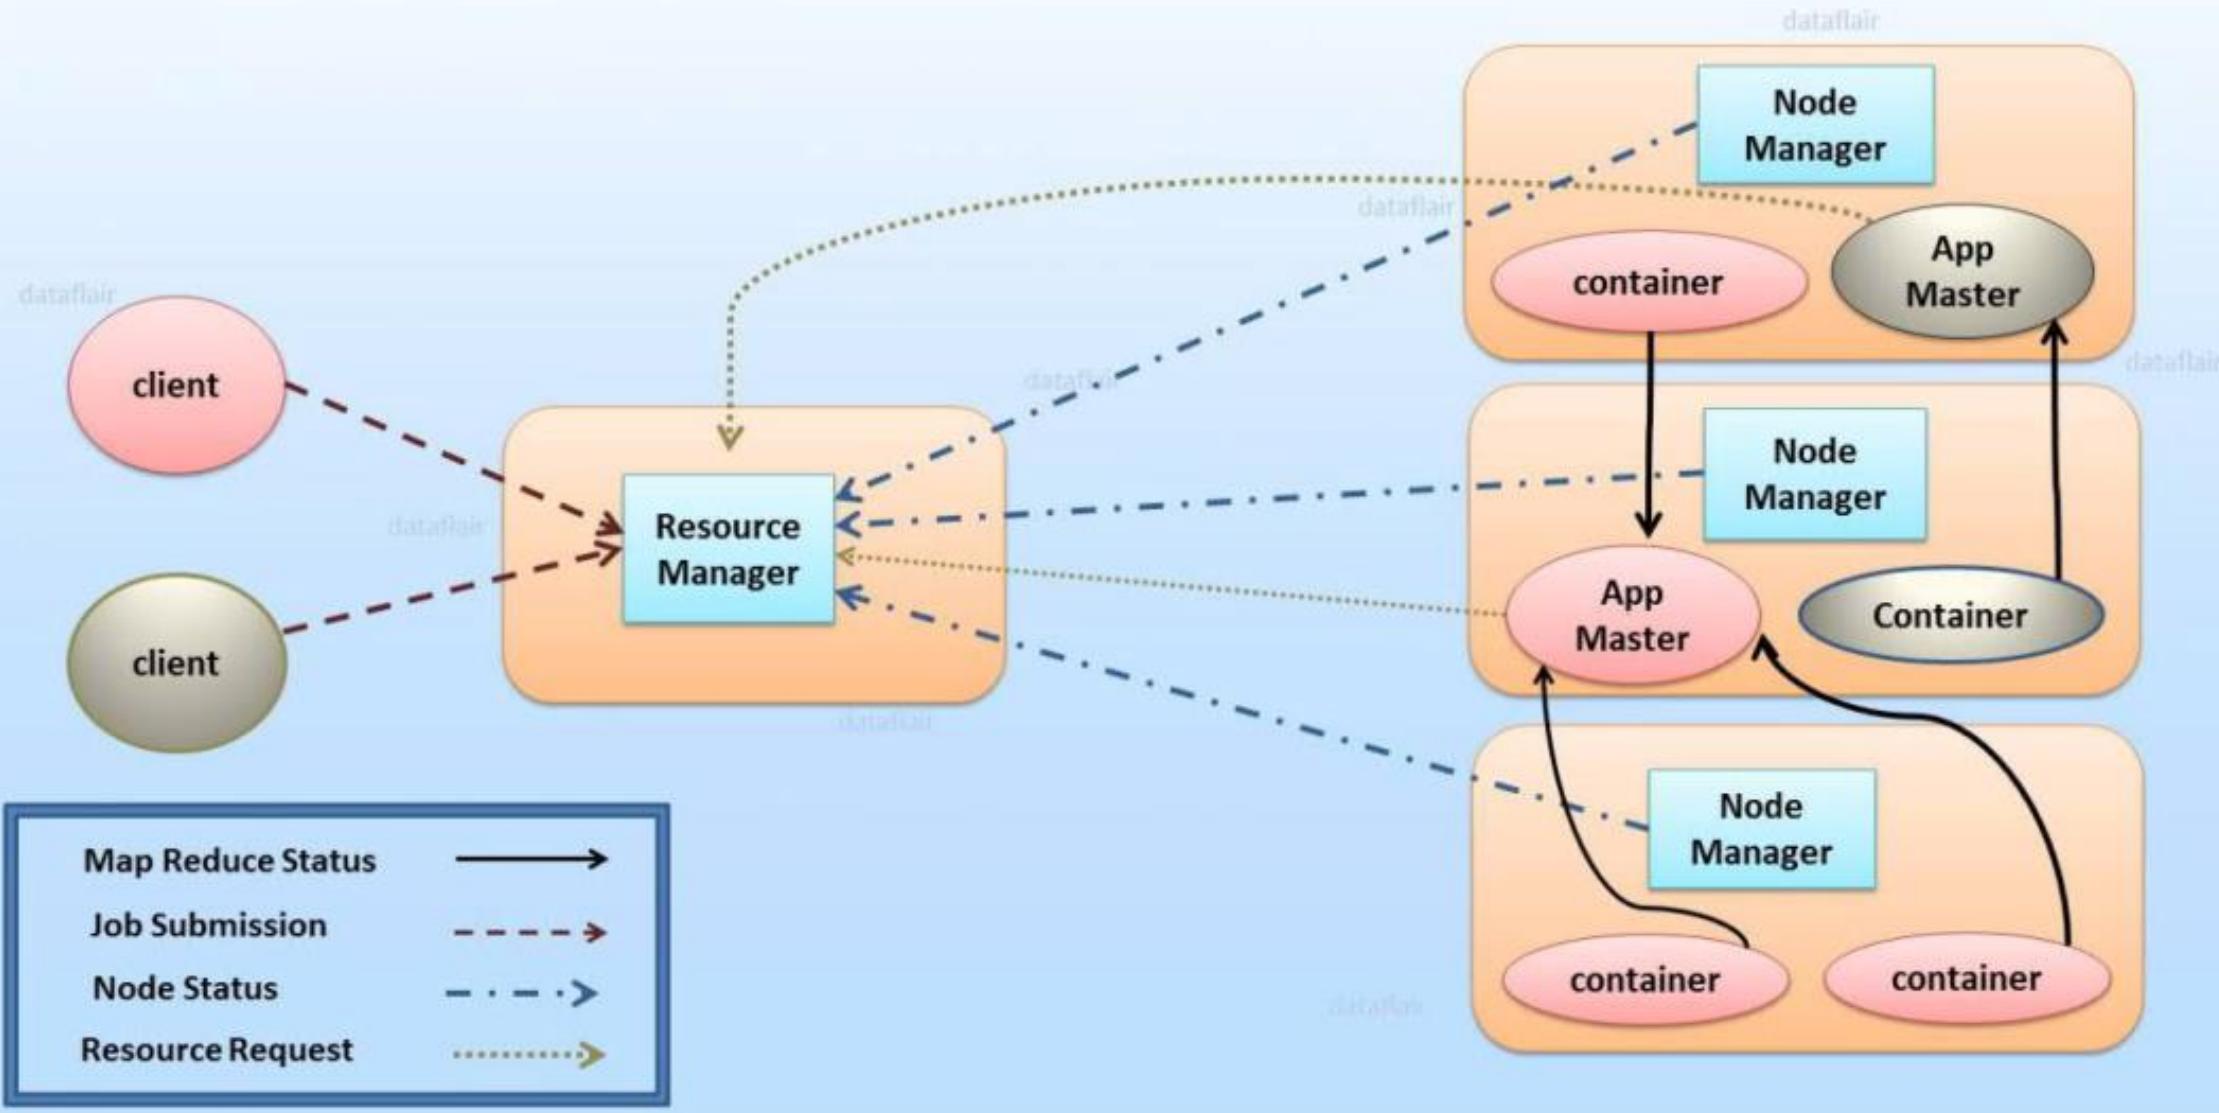
Hadoop Ecosystem and Their Components

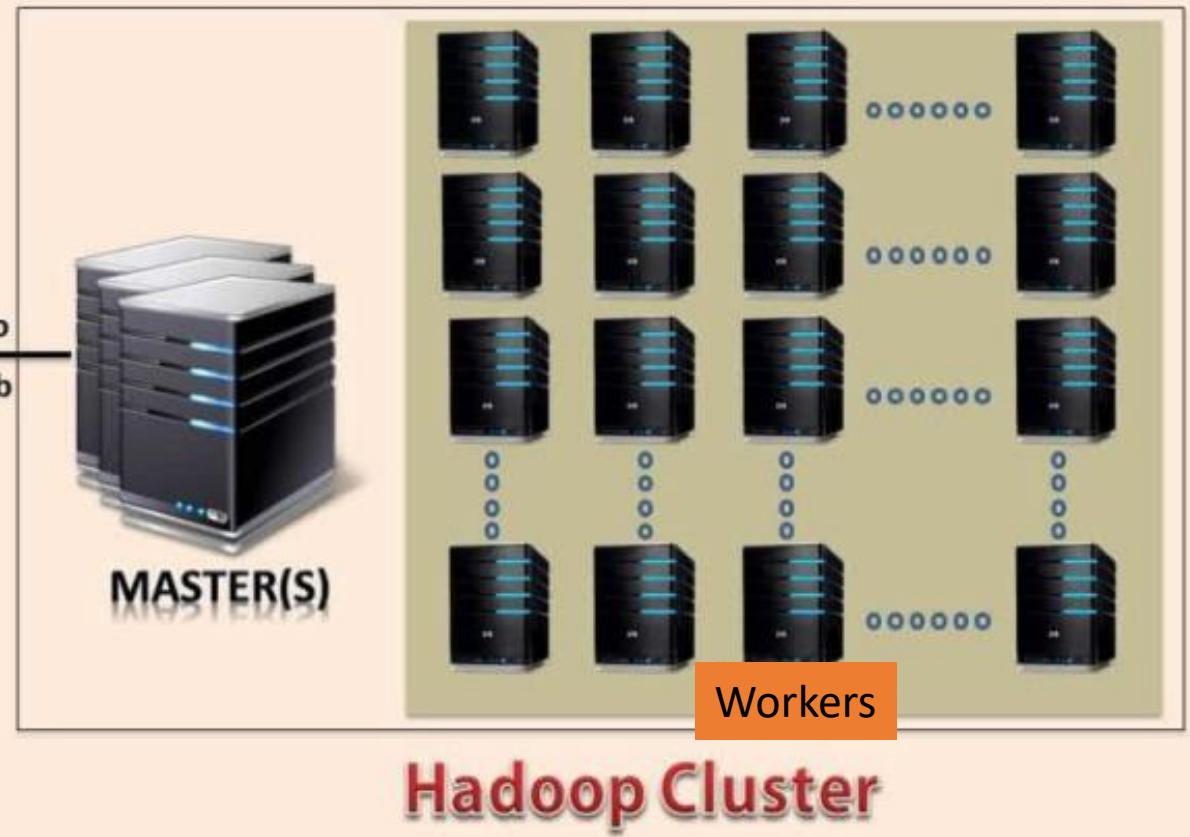
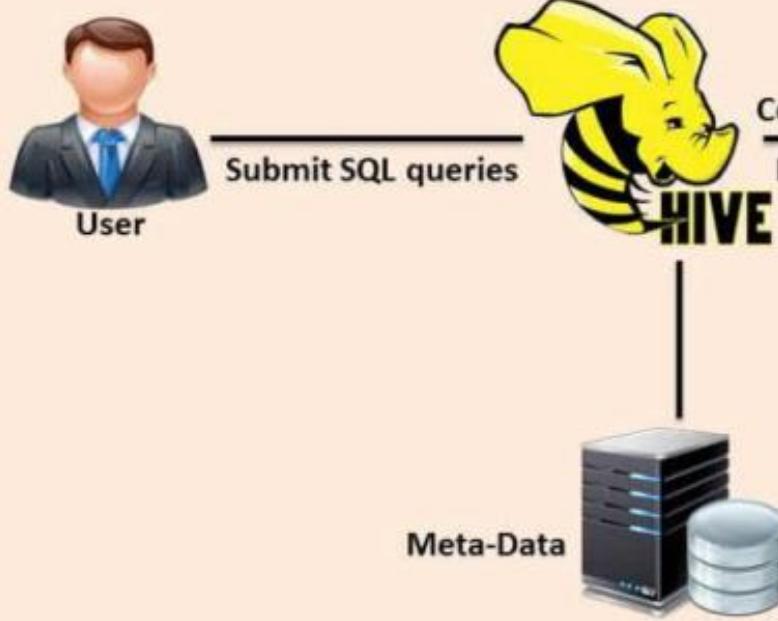


MapReduce Job Execution Flow

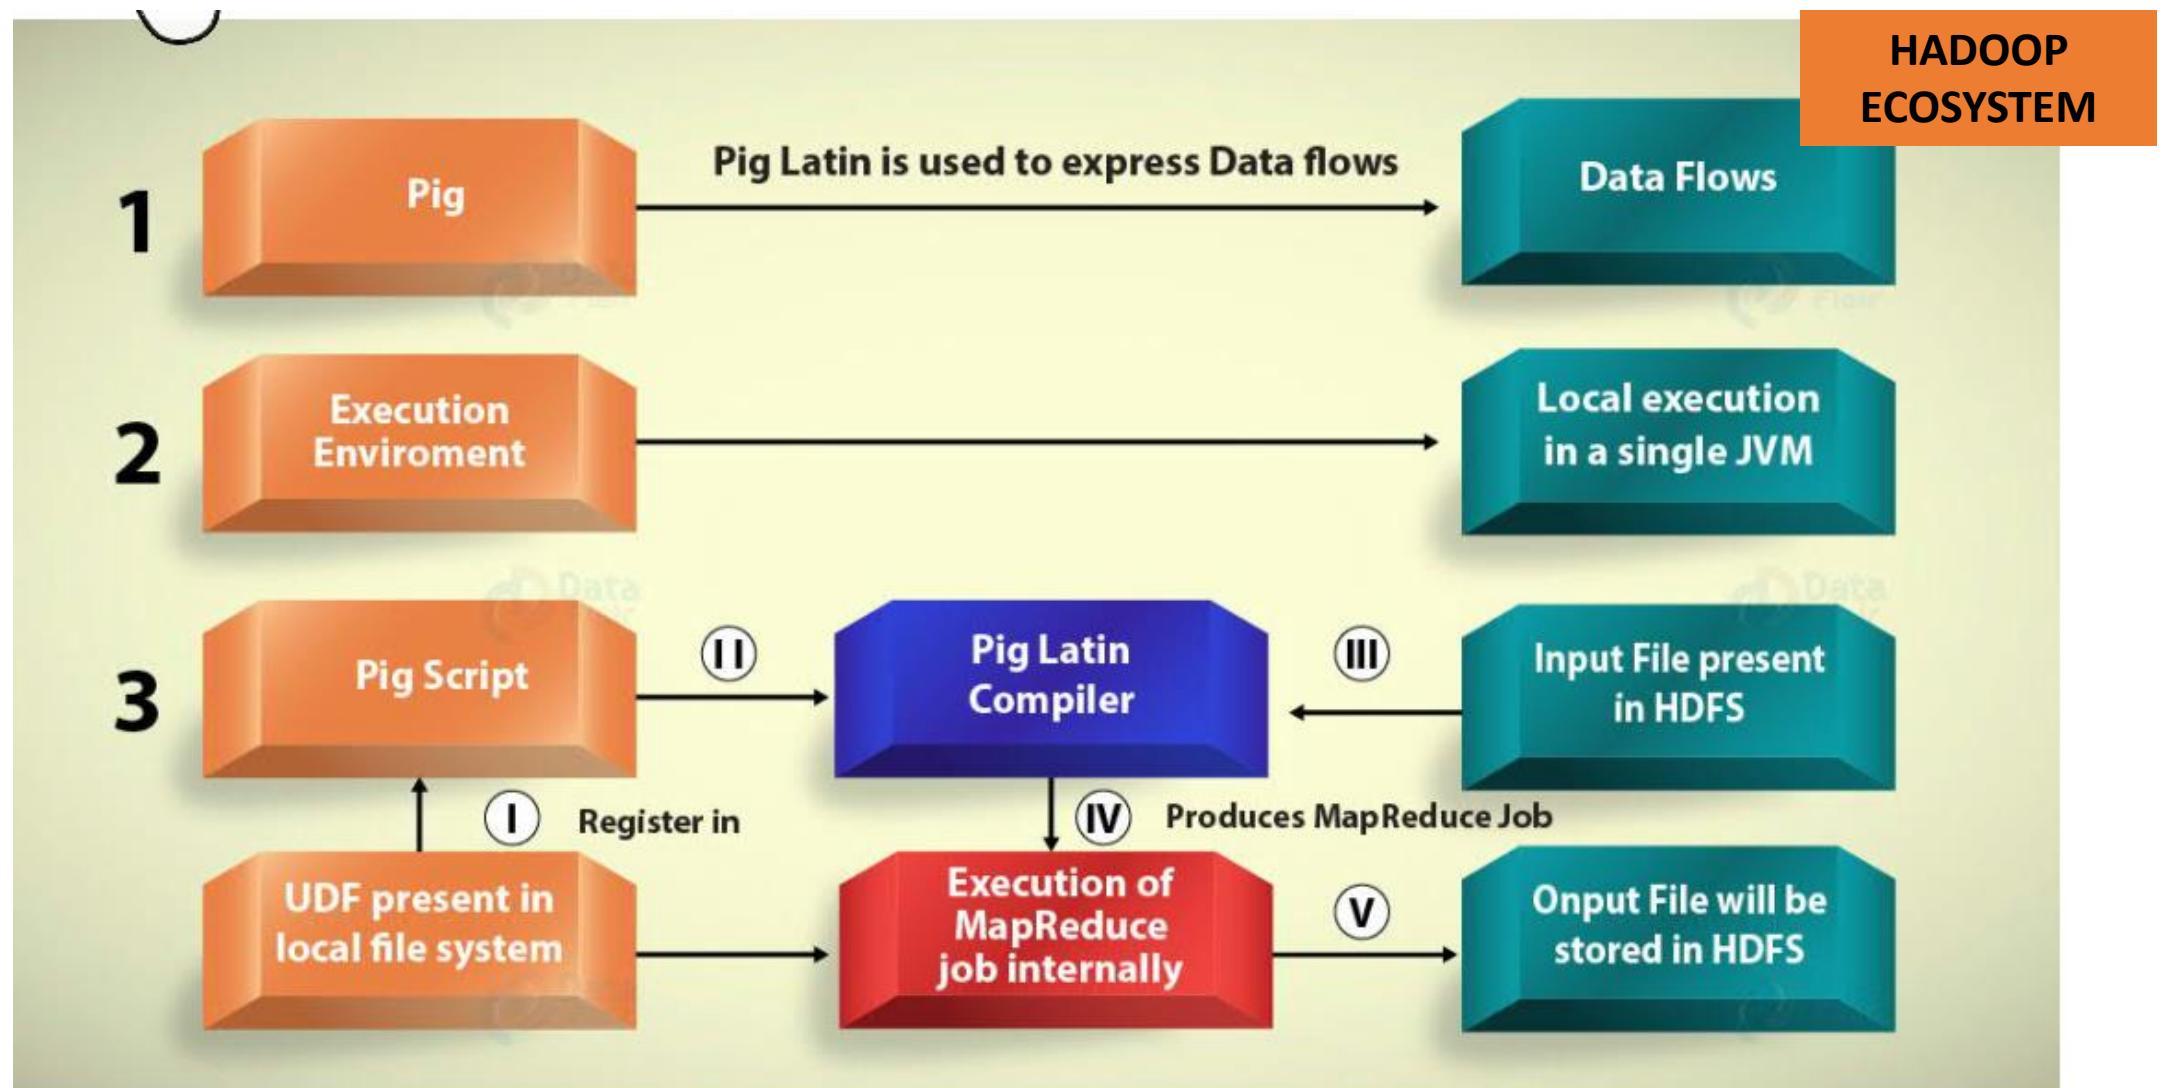


Apache Hadoop YARN-Architecture

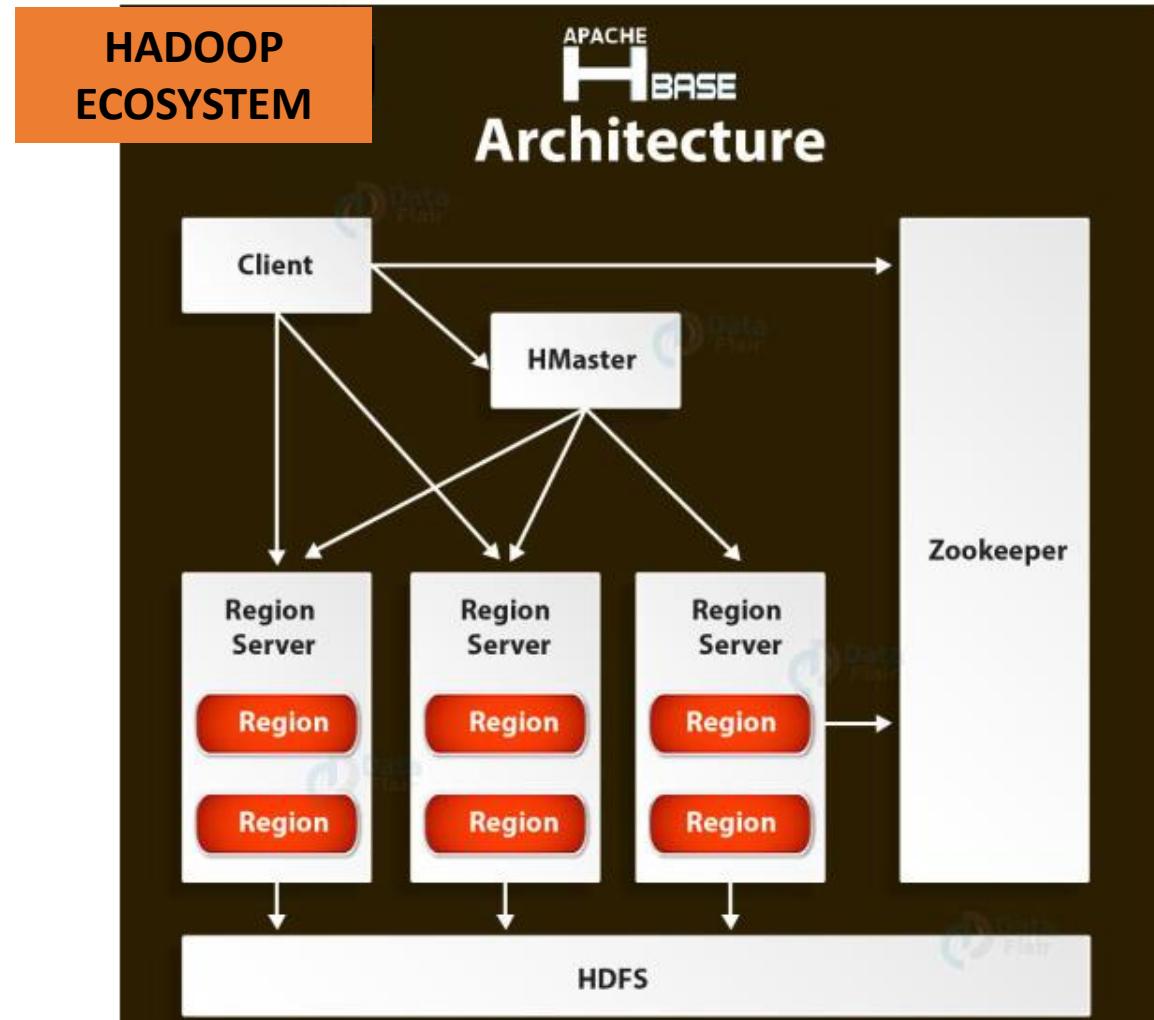




HADOOP
ECOSYSTEM



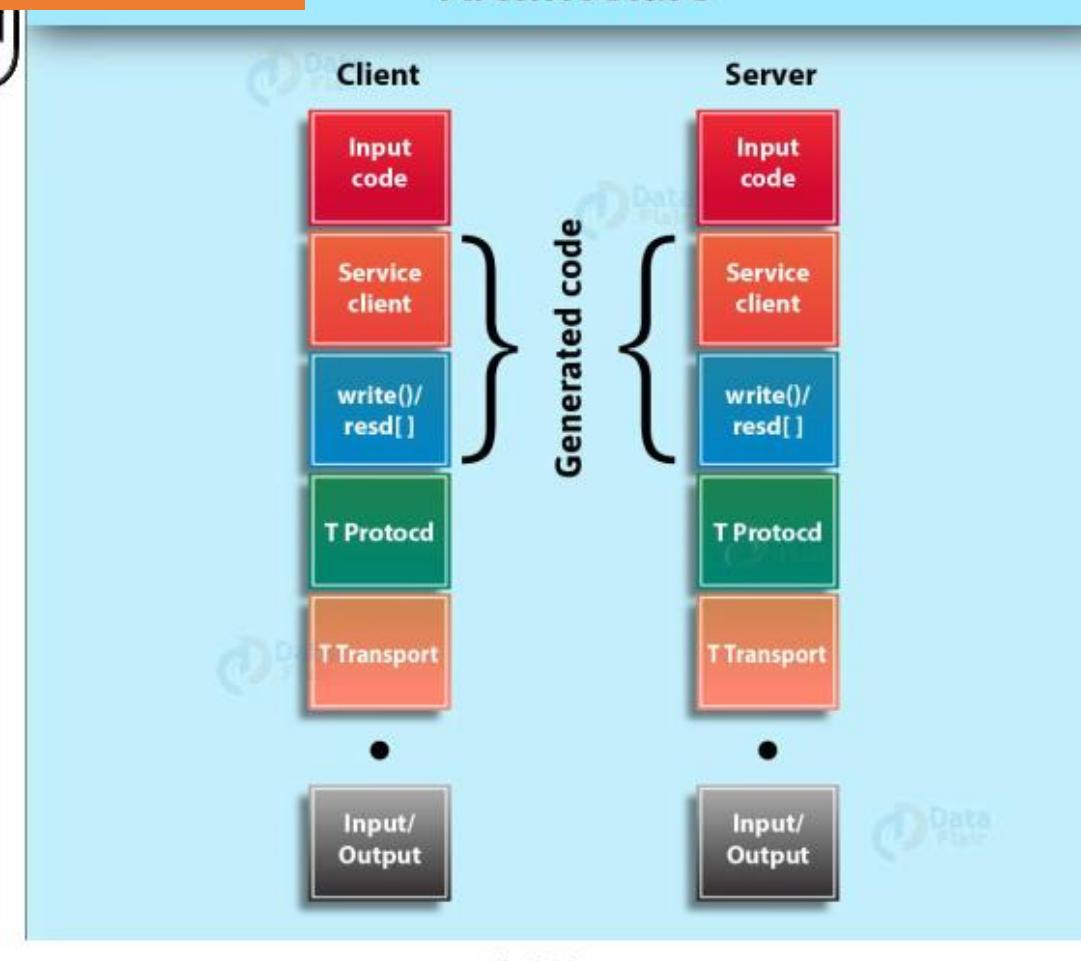
Pig Diagram



HBase Diagram

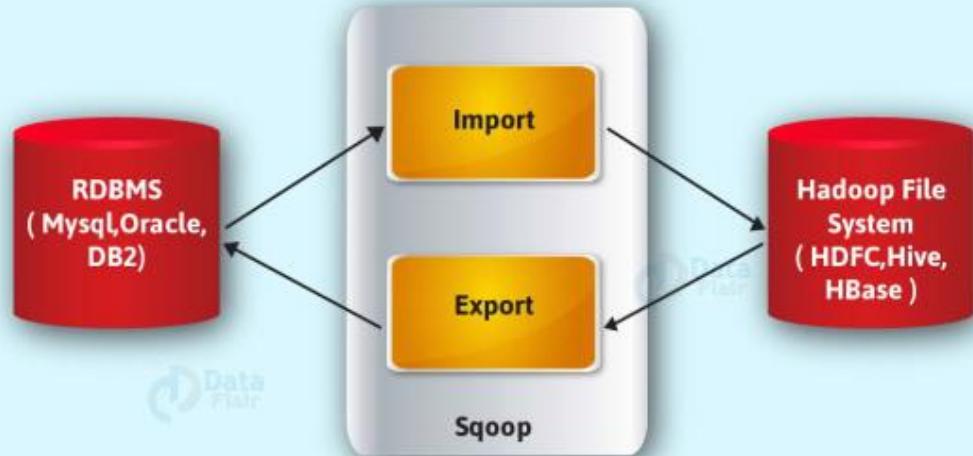
HADOOP ECOSYSTEM

Apache Thrift Architecture

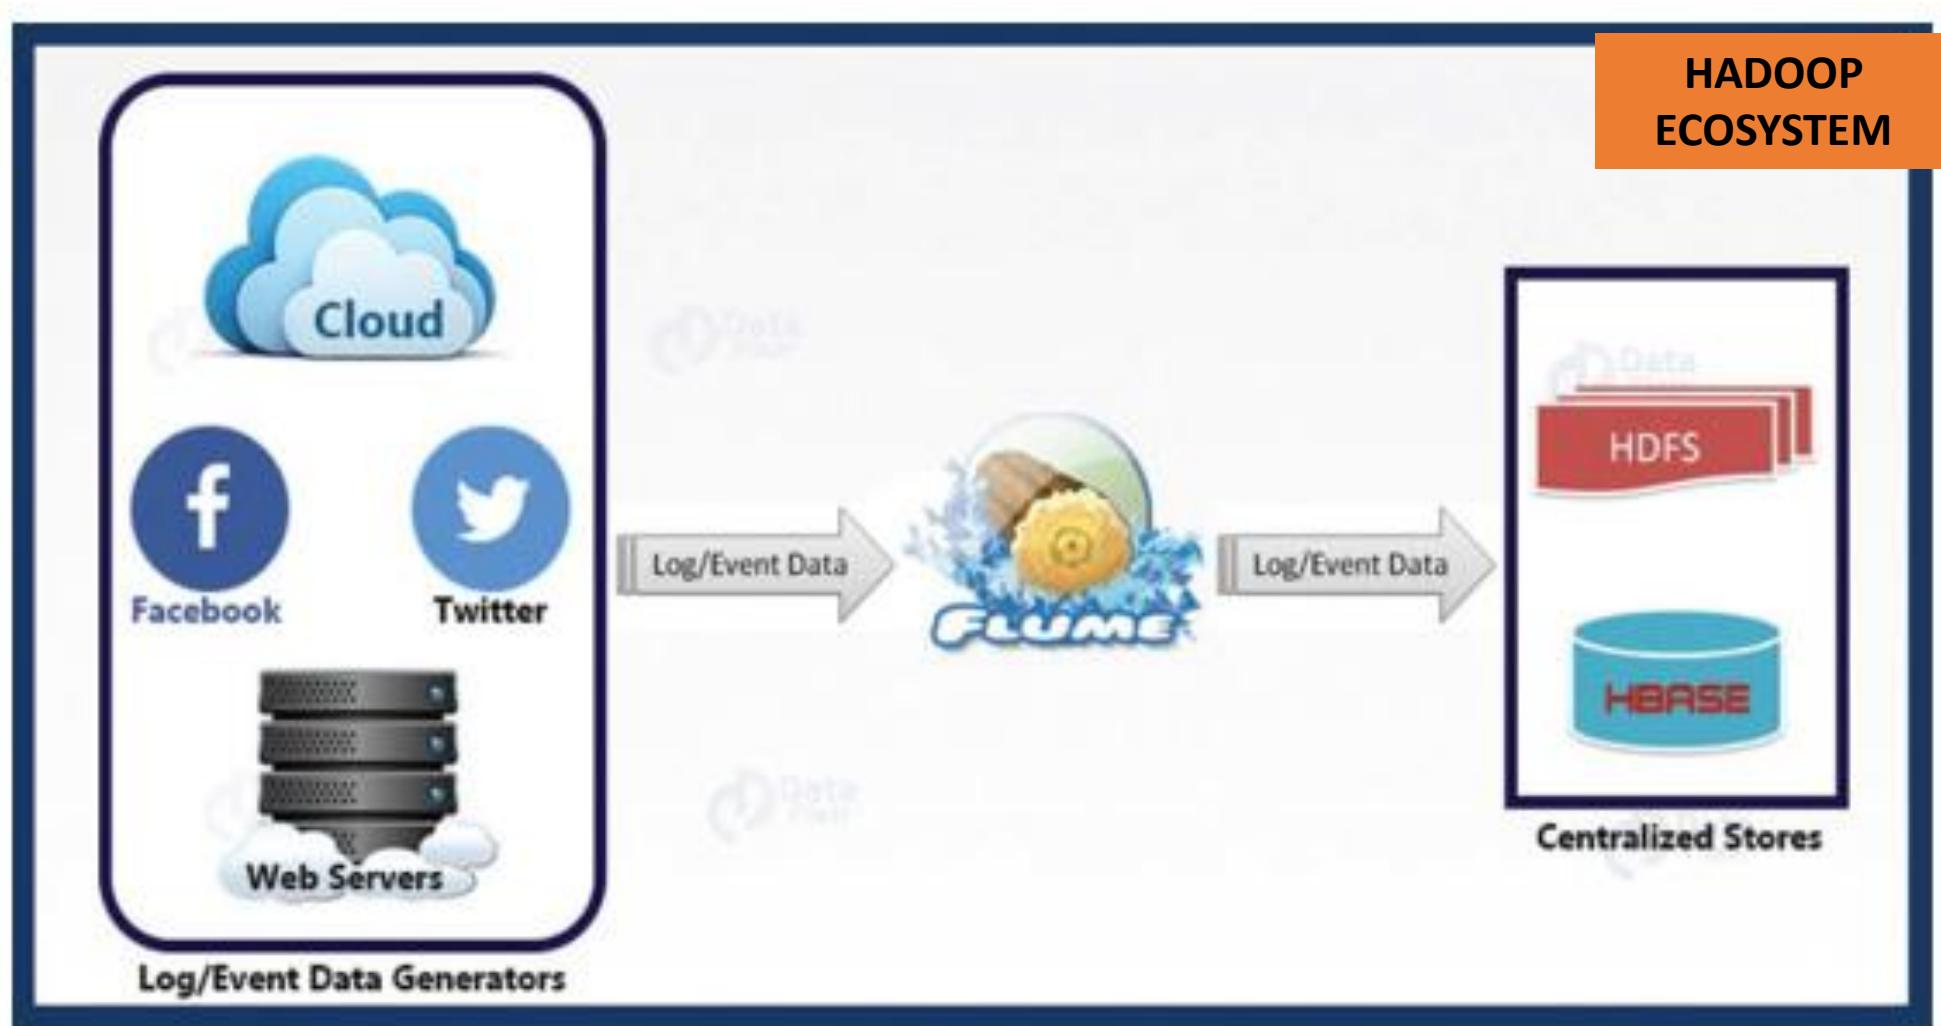


HADOOP ECOSYSTEM

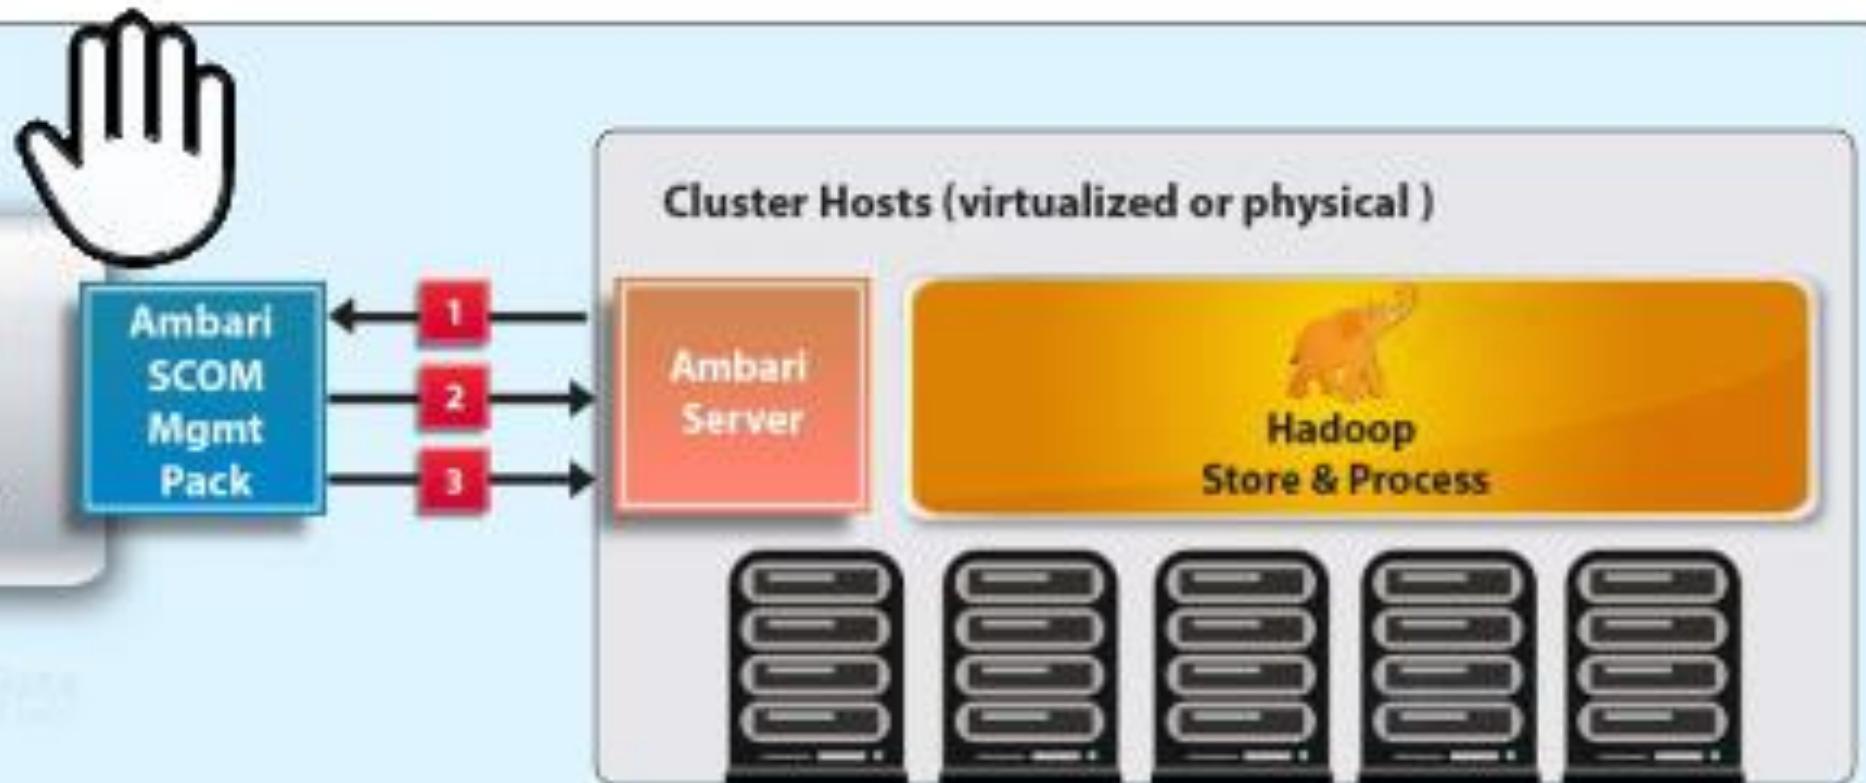
How Sqoop Works



Apache Sqoop Diagram

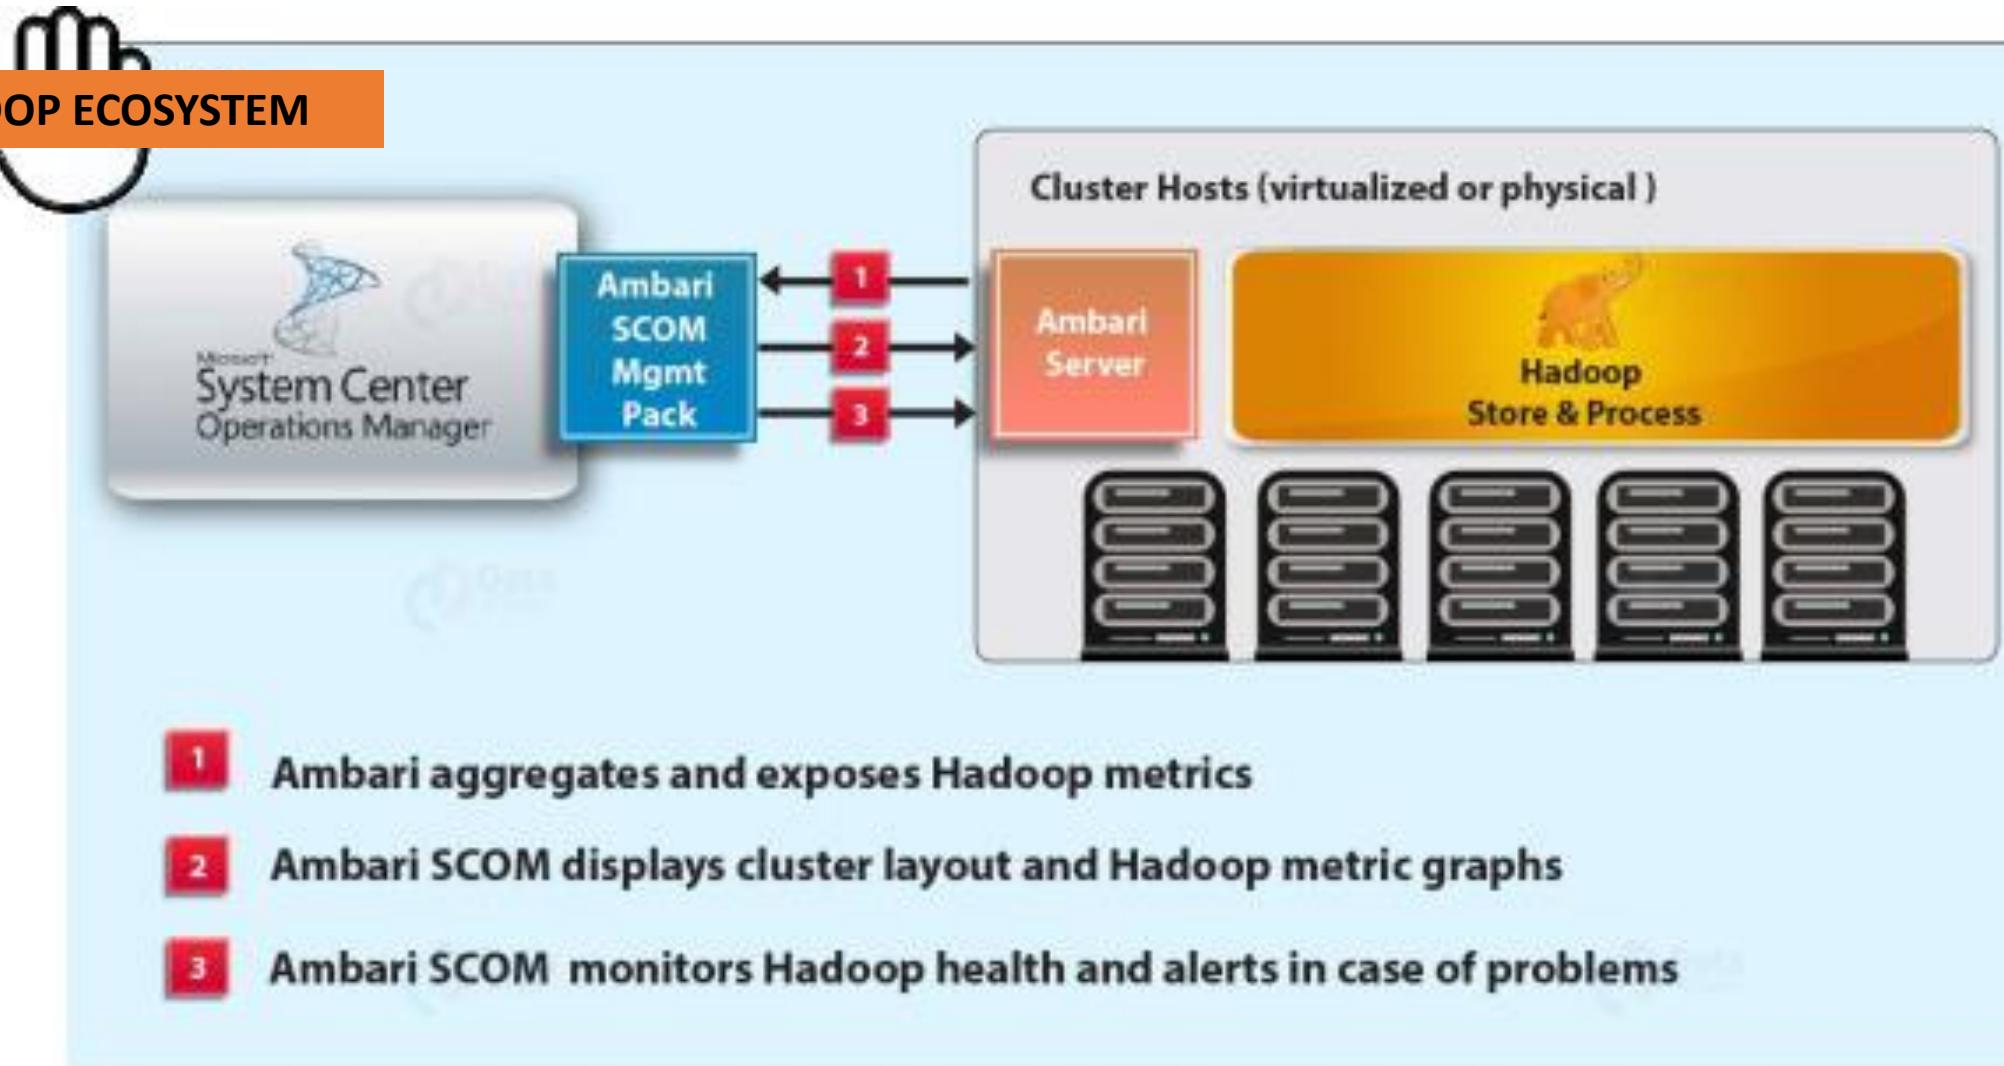


Apache Flume

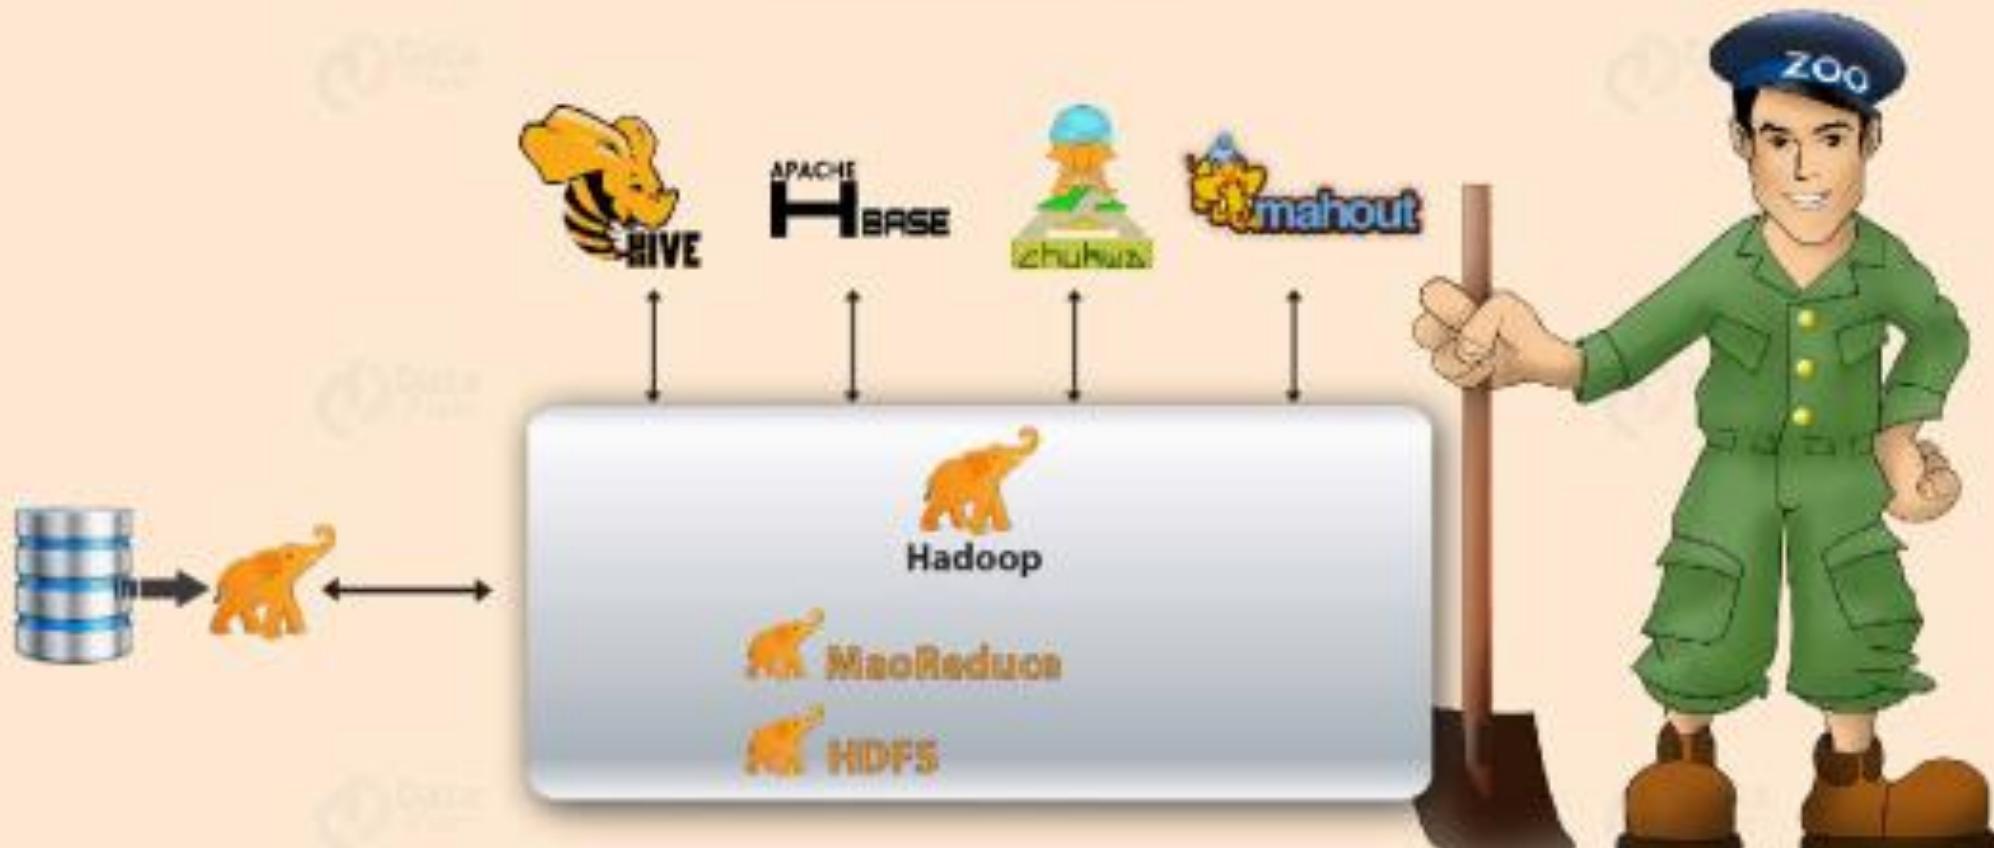


- 1** Ambari aggregates and exposes Hadoop metrics
- 2** Ambari SCOM displays cluster layout and Hadoop metric graphs
- 3** Ambari SCOM monitors Hadoop health and alerts in case of problems

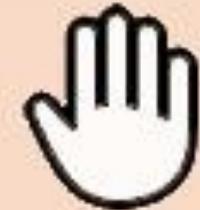
HADOOP ECOSYSTEM



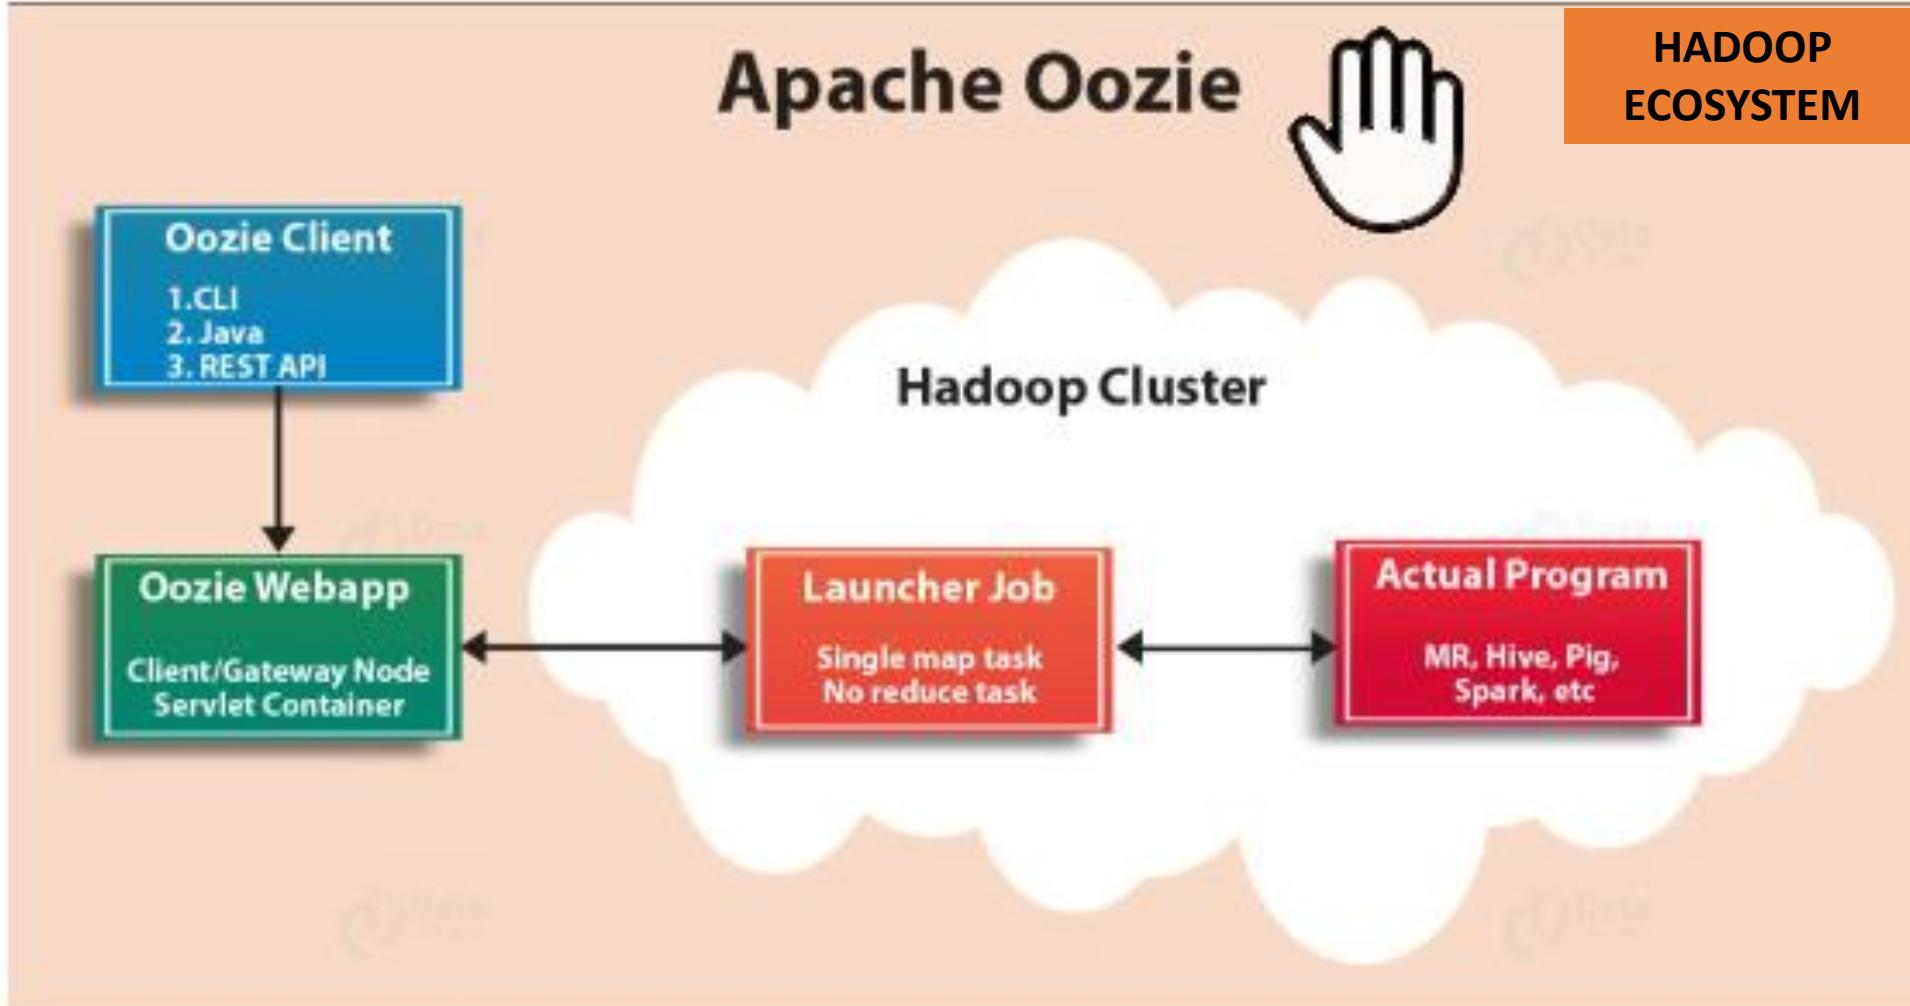
ZooKeeper



Apache Oozie

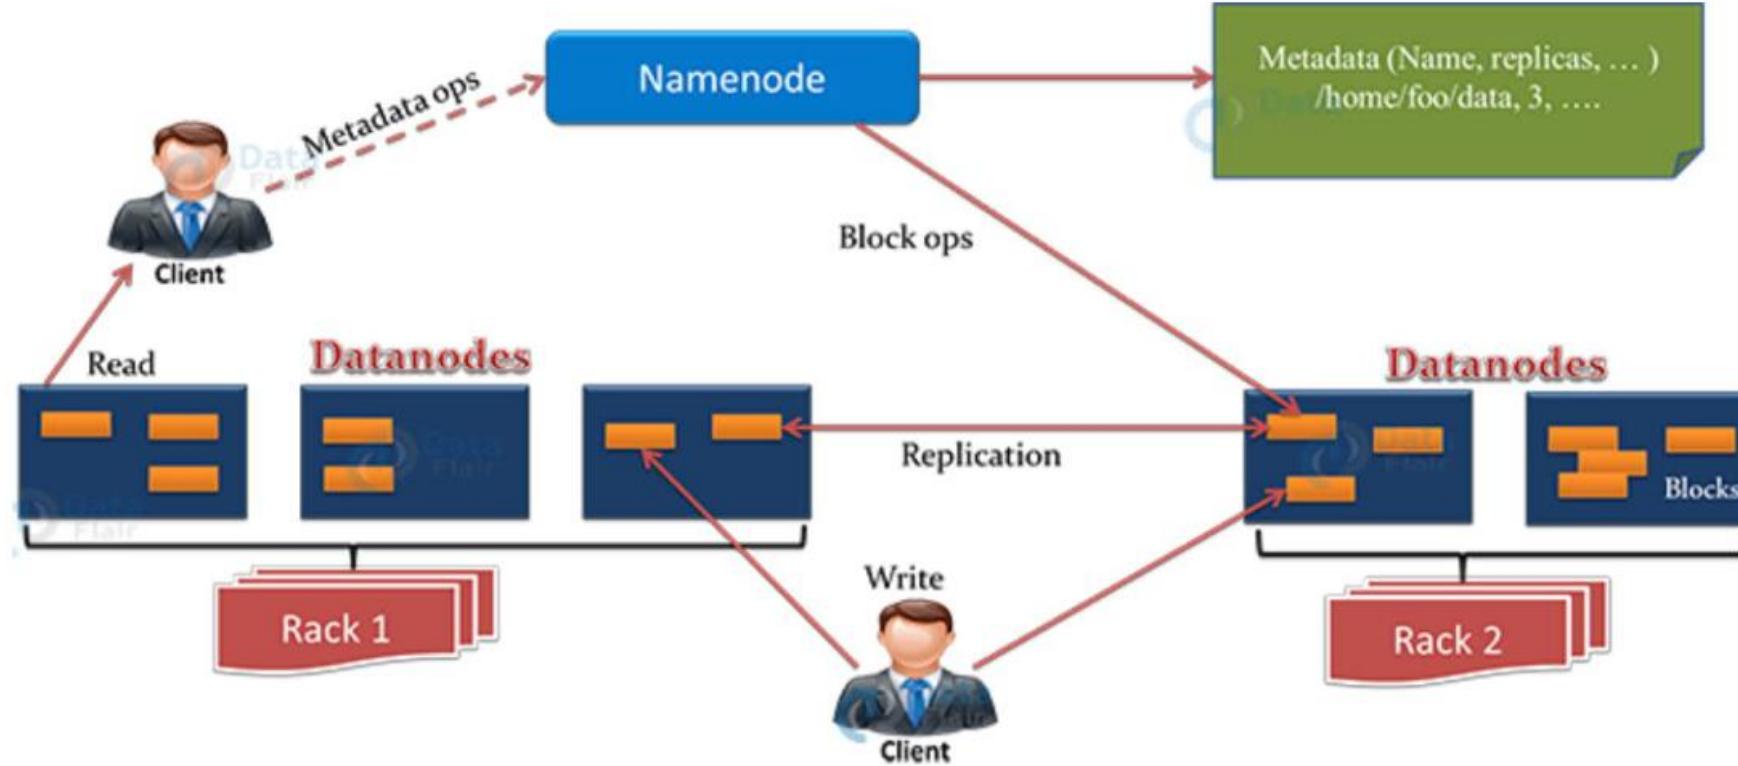


HADOOP
ECOSYSTEM

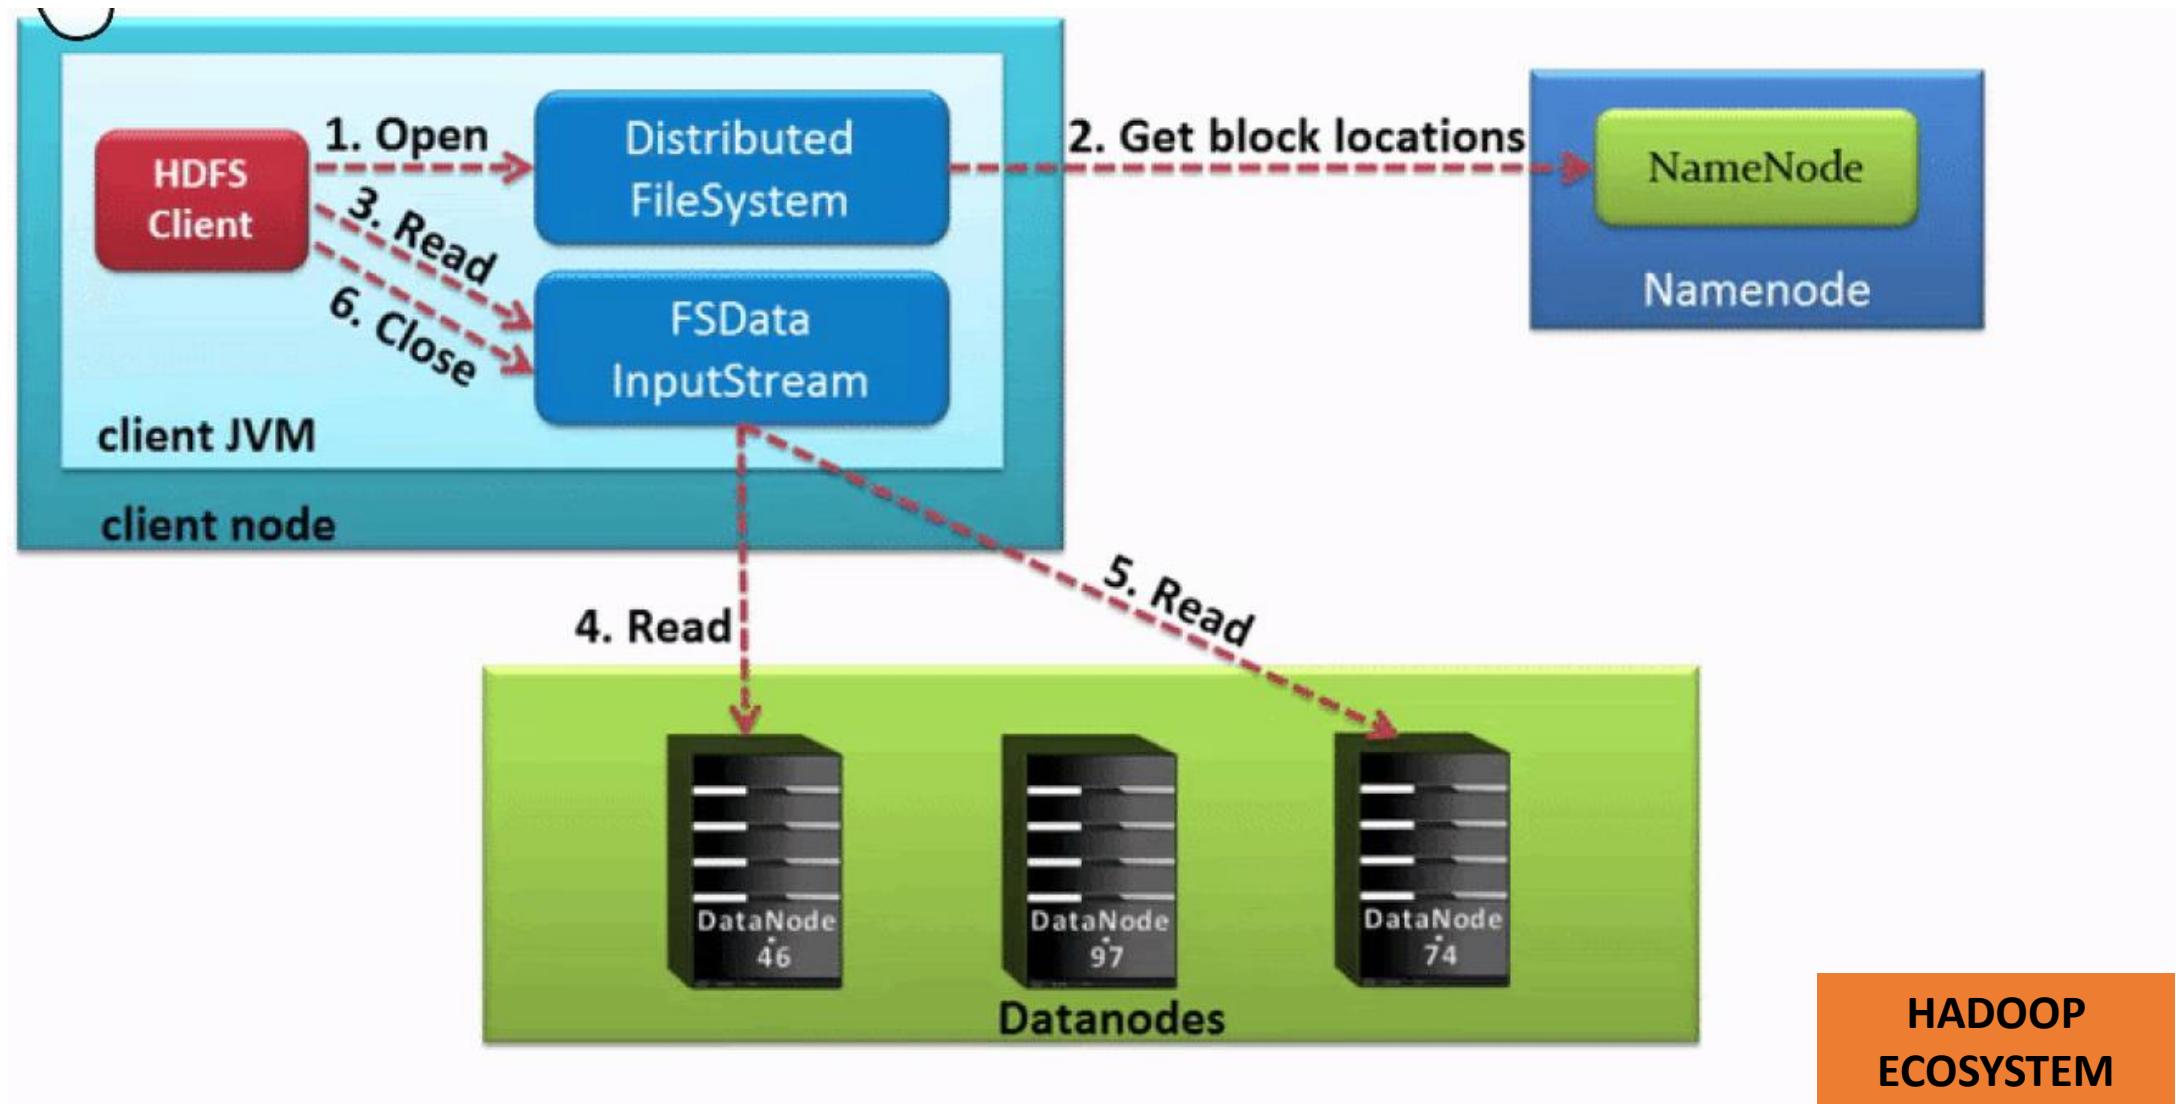


Oozie Diagram

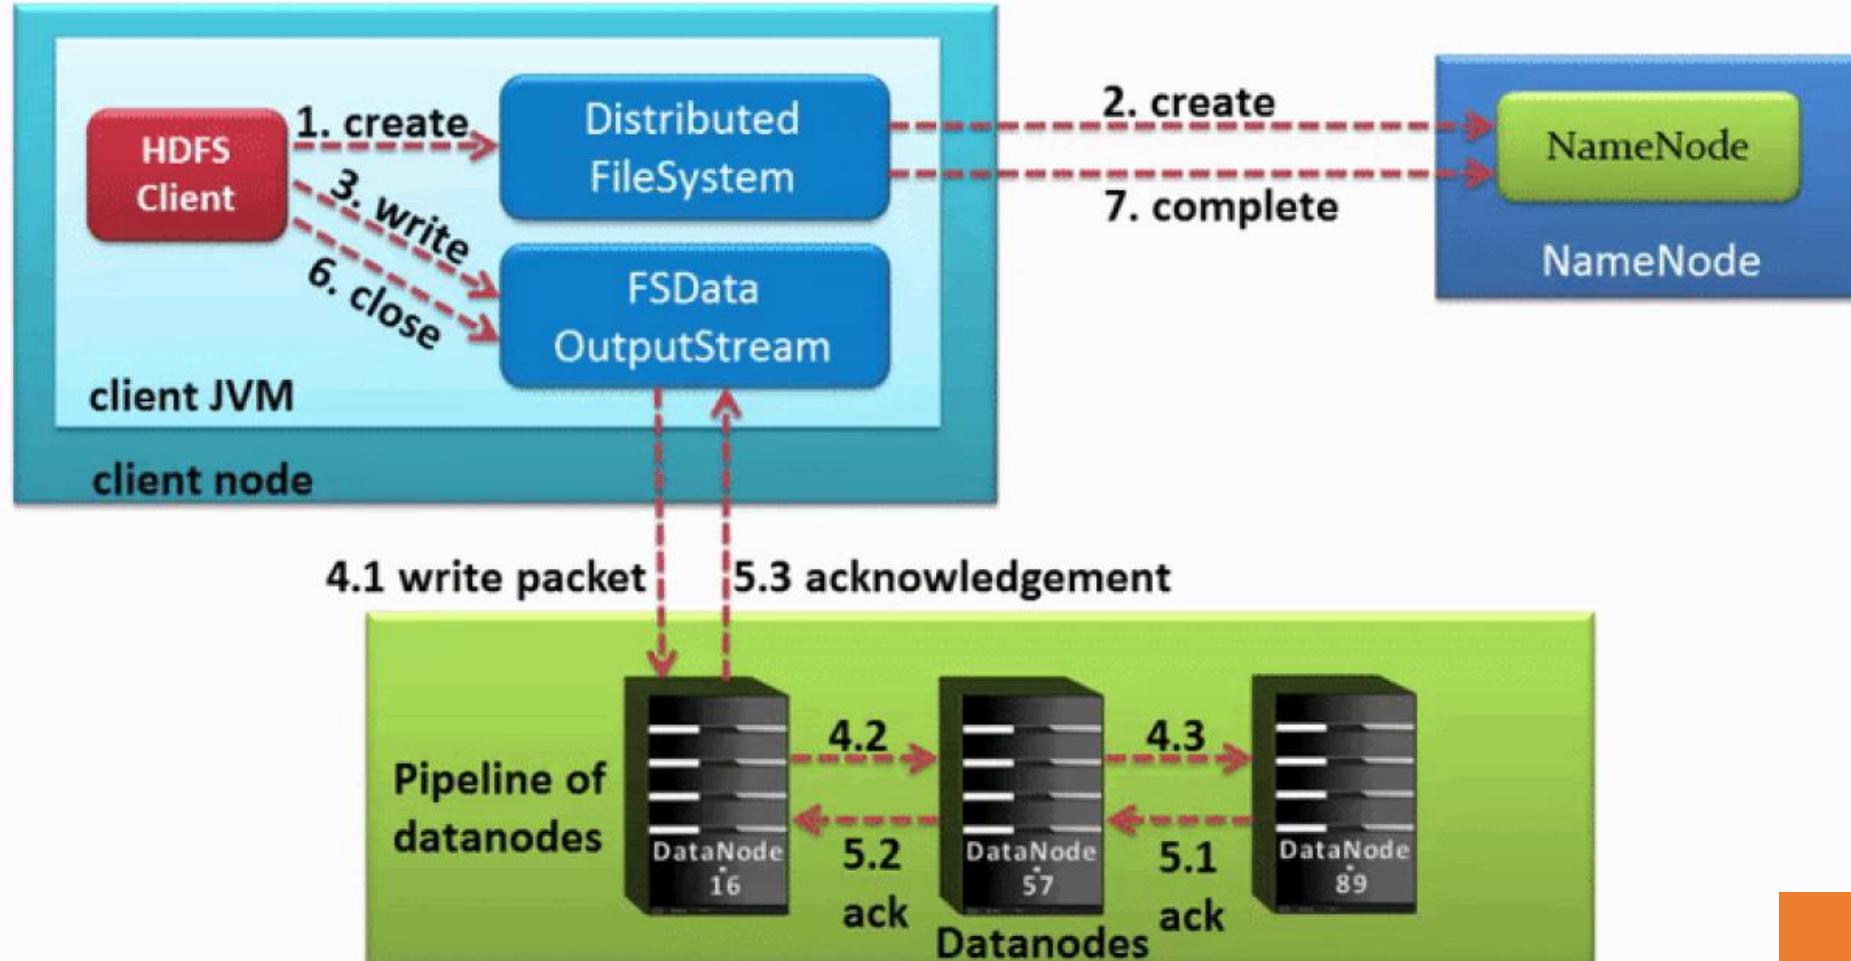
HADOOP ECOSYSTEM



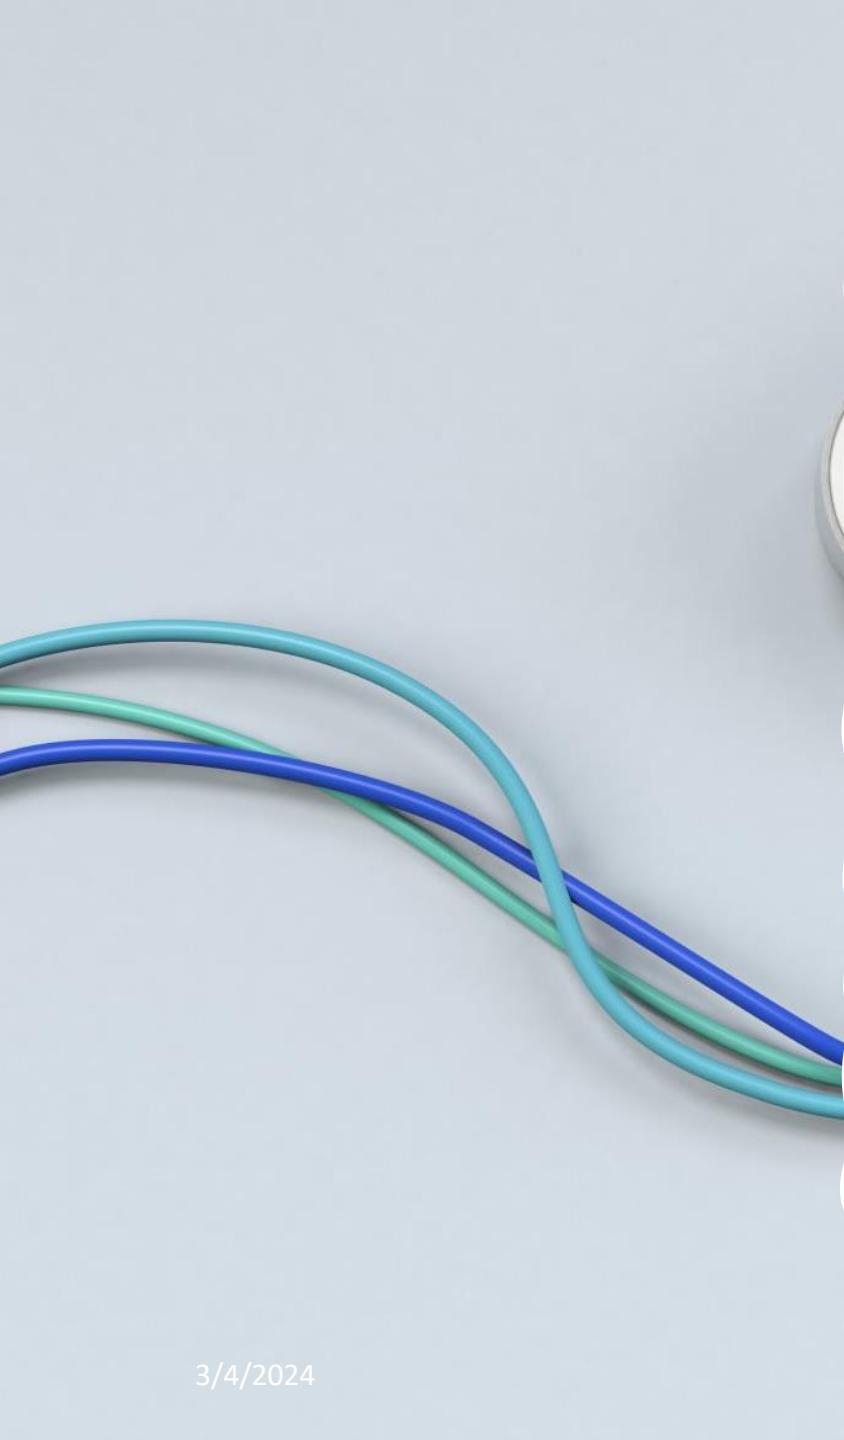
Read Operation in Hadoop HDFS



Write Operation in Hadoop HDFS with 2 replications



HADOOP
ECOSYSTEM



Summary

- Elastic Map Reduce (EMR) is the cloud-version of MR that is supported by all cloud providers in a platform as a service model (PaaS)
- Hadoop Ecosystem was developed in open-source as a cloud version for Infrastructure as a Services (IaaS) model.
- Handling a few very large data files that have more reads than writes (updates model) is what is used for large data analytics applications.